

Introduction to Machine Learning

-

How to design machine learning tools for Optics

Workshop

May 6th 2024

Louis Rossignol



COLLÈGE
DE FRANCE
1530



Reference and good ressource

- 3blue1brown on Neural Networks:
- [https://youtube.com/playlist?
list=PLZHQQObOWTQDNU6R1_67000Dx_ZCJB-3pi&si=xvbZnuQsteonsFpn](https://youtube.com/playlist?list=PLZHQQObOWTQDNU6R1_67000Dx_ZCJB-3pi&si=xvbZnuQsteonsFpn)

Outline

- Introduction
- Machine learning mechanics
 - Back-propagation of Error and Gradient Descent
 - Activation function
 - Loss function
 - Cross entropy loss and Classification
 - Mean Squared Error and Regression
 - Convolutional Neural Networks

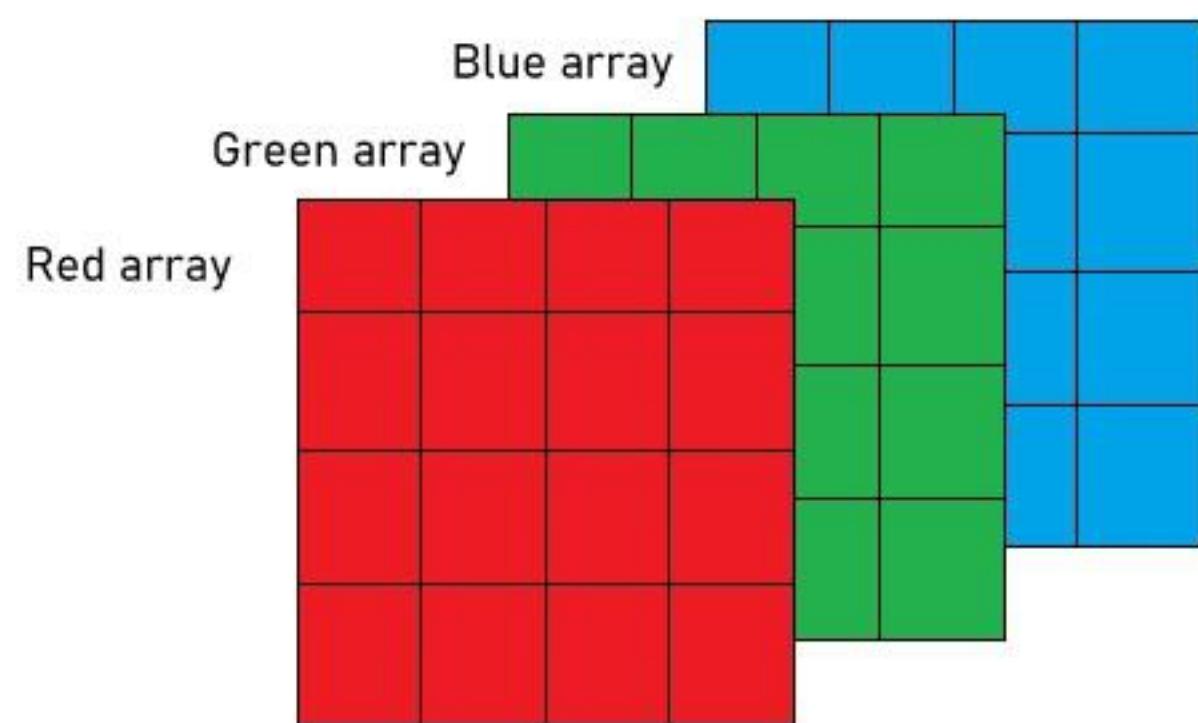
Introduction

- Data
- Model
- Loss function

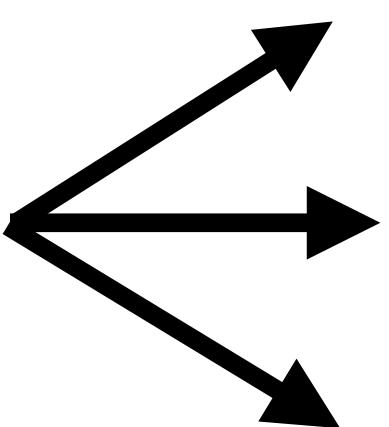
Introduction

Data

Image [1]



Arrays stacked over each other
to form a Digital Image.



Pig
Dog
Cat

Sound [2]



Classification

[1] - <https://www.google.com/url?sa=i&url=https%3A%2F%2Fwww.geeksforgeeks.org%2Fmatlab-rgb-image-representation%2F&psig=AOvVaw3nGrGp2TvHzoNbSHI8GIK8&ust=1714476975492000&source=images&cd=vfe&opi=89978449&ved=oCBIQjRxqFwoTCNDEicqq54UDFQAAAAAdAAAAABAE>

[2] - <https://www.google.com/url?sa=i&url=https%3A%2F%2Fpngtree.com%2Fso%2Fblack-and-white-sound-waves&psig=AOvVaw3ZKMWpx1T6W2Yw4UVqrS33&ust=1714476738077000&source=images&cd=vfe&opi=89978449&ved=oCBIQjRxqFwoTCID-oNmp54UDFQAAAAAdAAAAABAK>

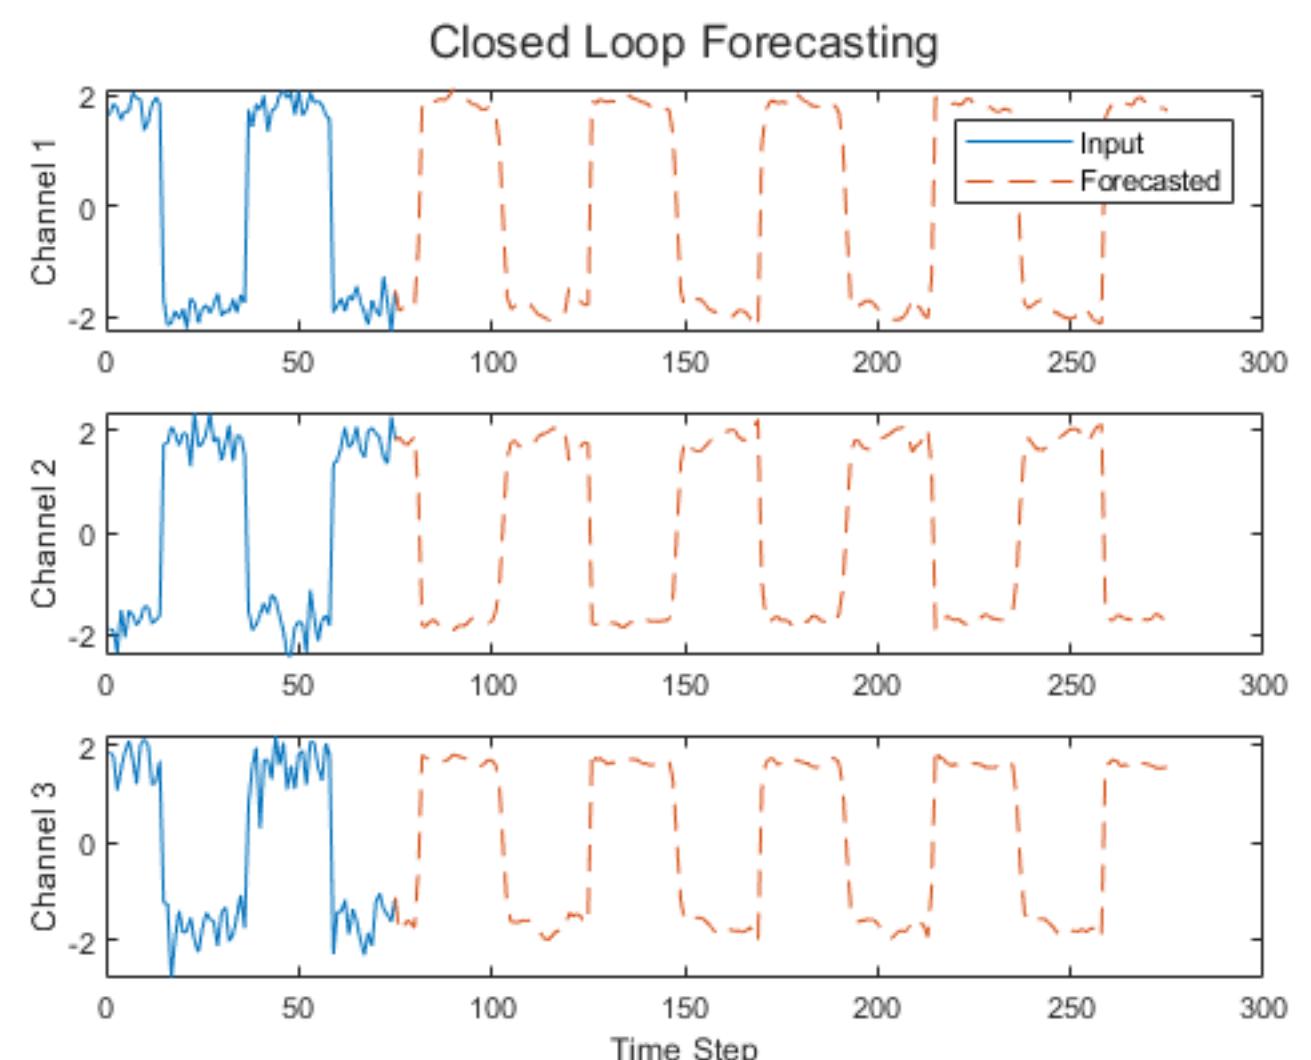
Introduction

Data

Image [1]



Sound [2]



Forecasting

[1] - <https://www.google.com/url?sa=i&url=https%3A%2F%2Fstock.adobe.com%2Ffr%2Fsearch%2Fimages%3Fk%3Dsequence&psig=AOvVaw3FXo8WKCjXAv0FTo8OsmGb&ust=1714477066240000&source=images&cd=vfe&opi=89978449&ved=oCBQQjhxqFwoTCMDZ4fWq54UDFQAAAAAdAAAABAE>

[2] - <https://www.google.com/url?sa=i&url=https%3A%2F%2Fwww.mathworks.com%2Fhelp%2Fdeeplearning%2Fug%2Ftime-series-forecasting-using-deep-learning.html&psig=AOvVaw35v6rvPsQNiSFWjHxQY4-i&ust=1714477261772000&source=images&cd=vfe&opi=89978449&ved=oCBIQjRxqFwoTCNCR9tKr54UDFQAAAAAdAAAABAJ>

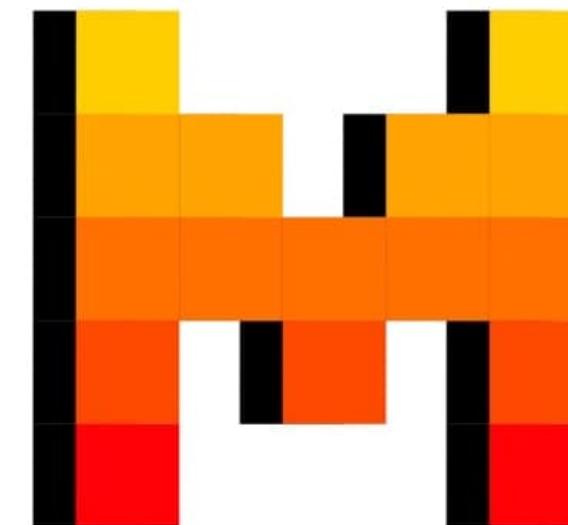
Introduction

Data

Text



Text prediction



[1] - <https://www.google.com/url?sa=i&url=https%3A%2F%2Fwww.emoji.com%2Fview%2Femoji%2F1068%2Fproposed%2Fpages&psig=AOvVaw2xFcqlKI-W9YG-DS6hZKb5&ust=1714479822672000&source=images&cd=vfe&opi=89978449&ved=oCBIQjRxqFwoTCKiygZi54UDFQAAAAAdAAAAABAE>

[2] - <https://www.google.com/url?sa=i&url=https%3A%2F%2Fwww.geeky-gadgets.com%2Fchat-gpt-login%2F&psig=AOvVawozJUSVrrfJX2mgnomfiYvv&ust=1714479965542000&source=images&cd=vfe&opi=89978449&ved=oCBIQjRxqFwoTCPDb5aS254UDFQAAAAAdAAAAABAE>
<https://www.google.com/url?sa=i&url=https%3A%2F%2Fwww.jaimelesstartups.fr%2Fmistral-ai-ia-generative-lm%2F&psig=AOvVaw3HoeZ6G6buYd5ZIUq6MDgK&ust=1714479988170000&source=images&cd=vfe&opi=89978449&ved=oCBIQjRxqFwoTCNCNzqO254UDFQAAAAAdAAAAABAK>
<https://plugins.jetbrains.com/plugin/17718-github-copilot>
<https://www.google.com/url?sa=i&url=https%3A%2F%2Fca.linkedin.com%2Fcompany%2Fblackboxtech&psig=AOvVaw363p8kpYRiUvRuvnxwoGb9&ust=1714480207899000&source=images&cd=vfe&opi=89978449&ved=oCBQQjhxqFwoTCNiZoM-254UDFQAAAAAdAAAAABAE>

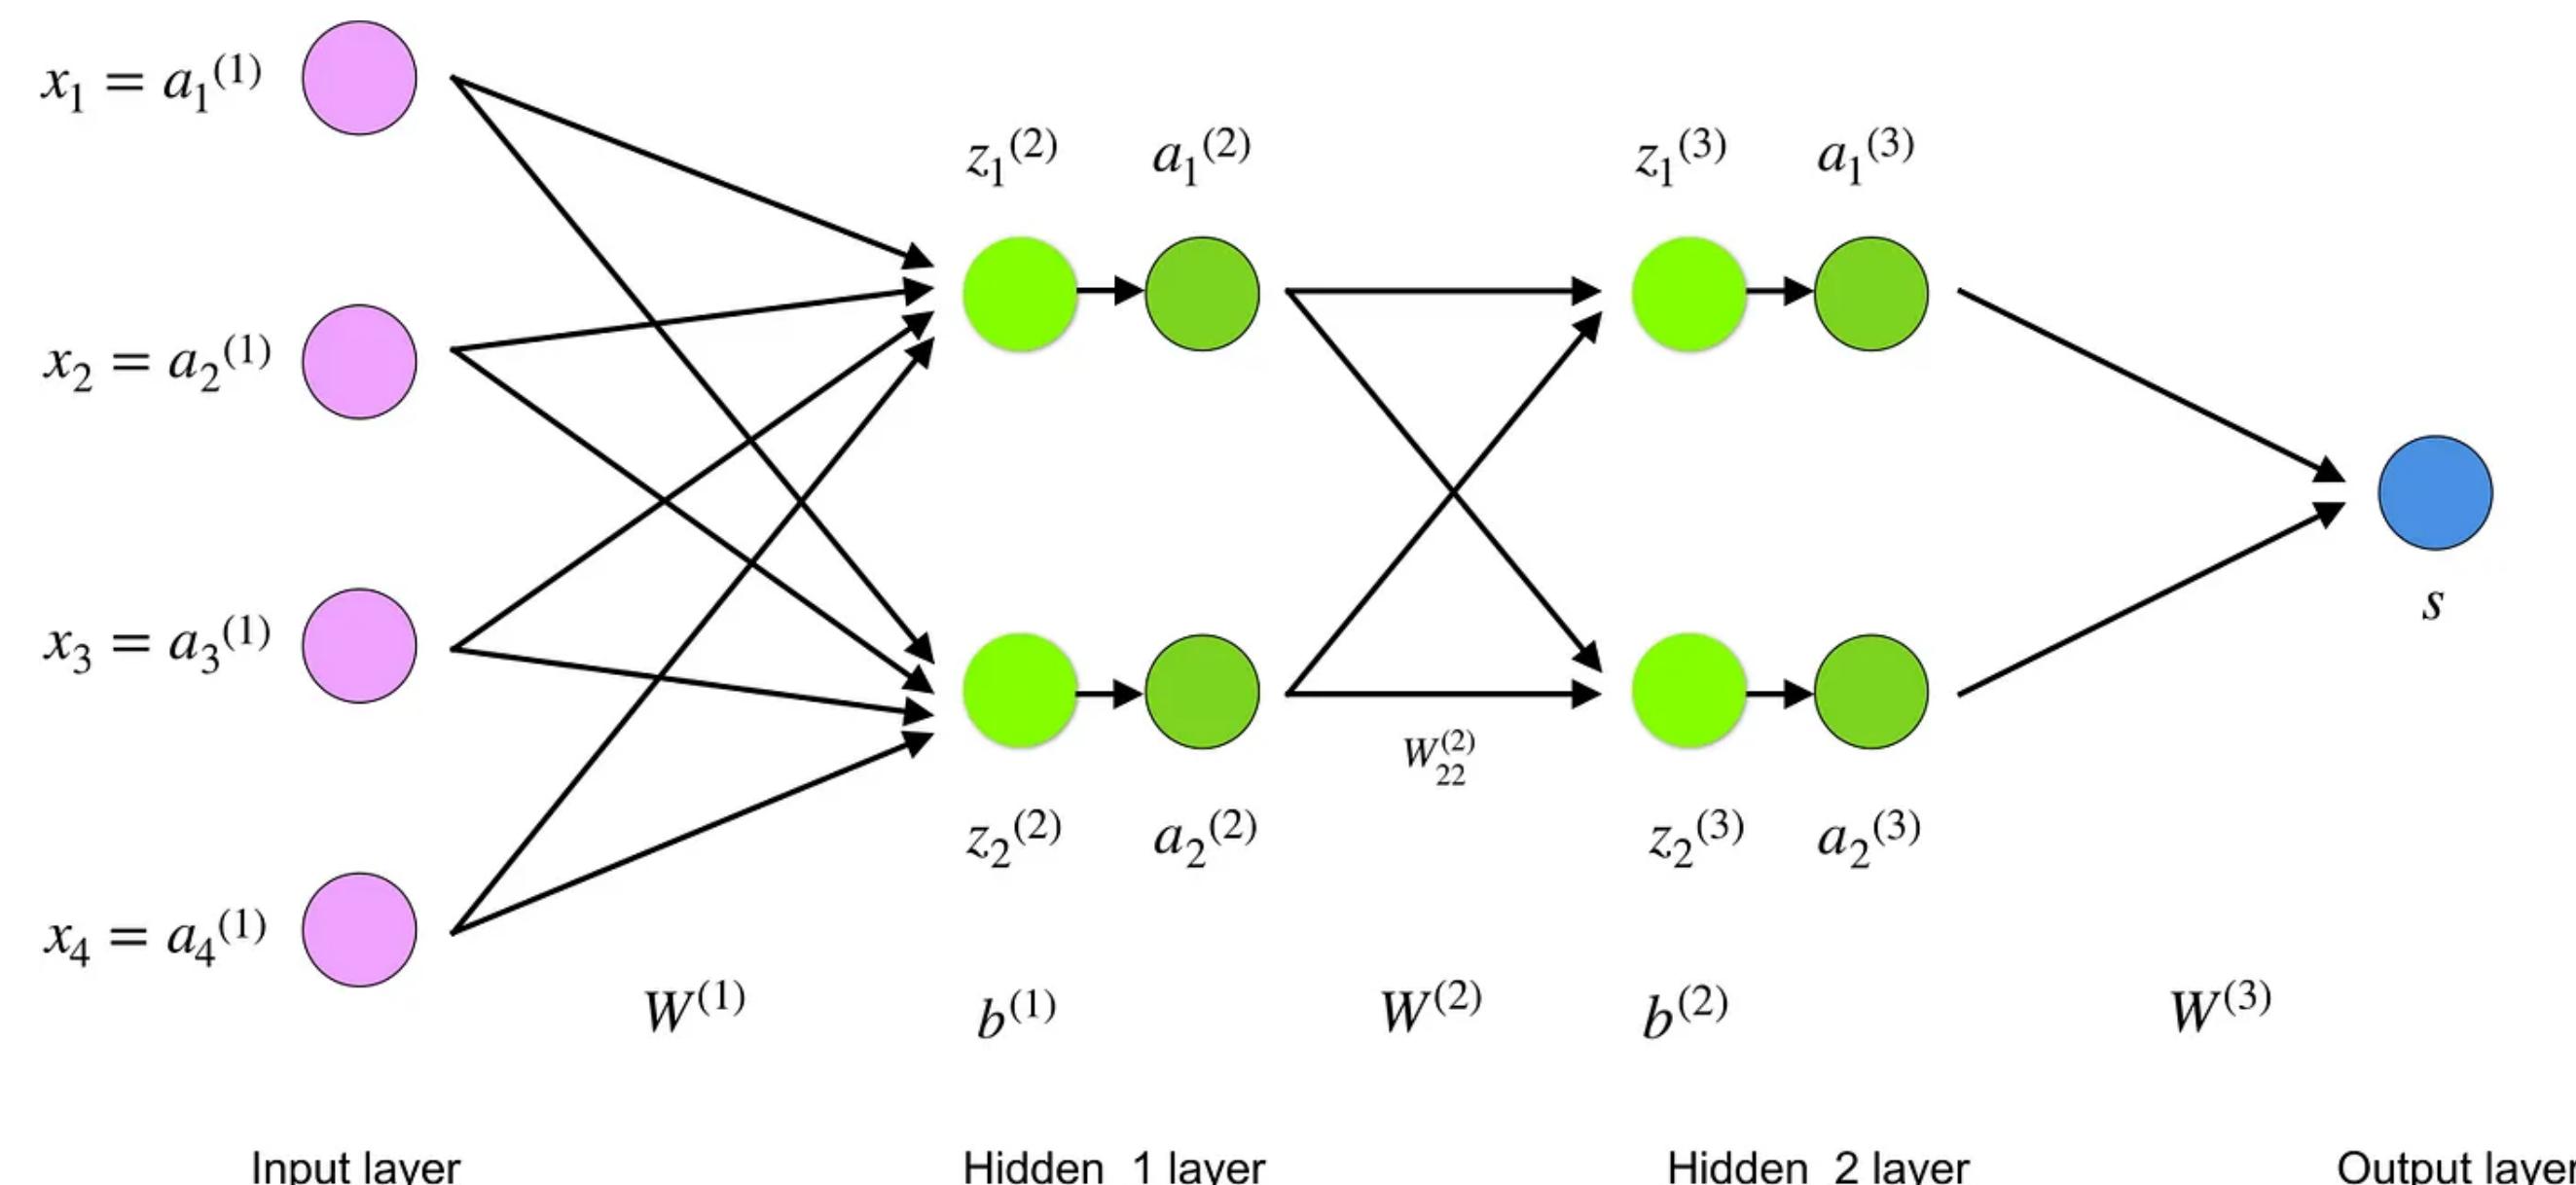
Important

- Generalization
- Data formatting
- **Model**

Introduction

Model

Fully connected layer [1]



Input layer

Hidden_1 layer

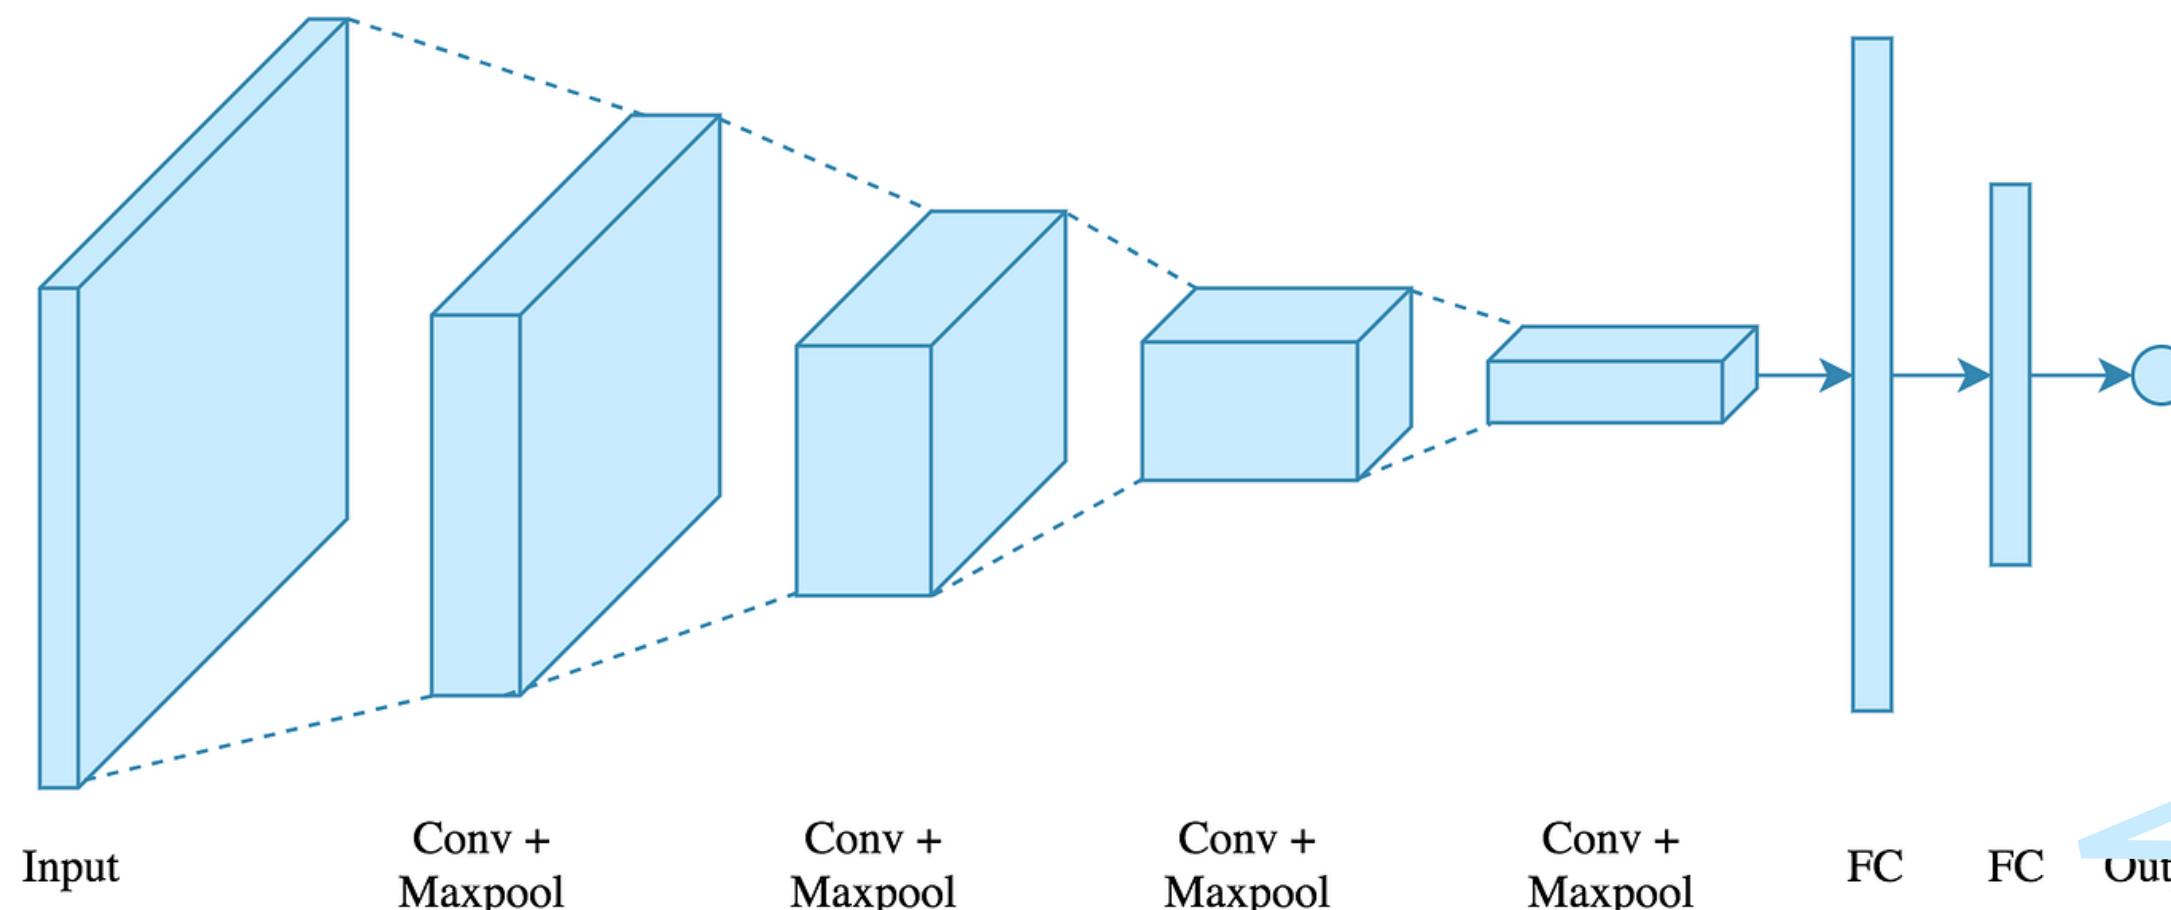
Hidden_2 layer

Output layer

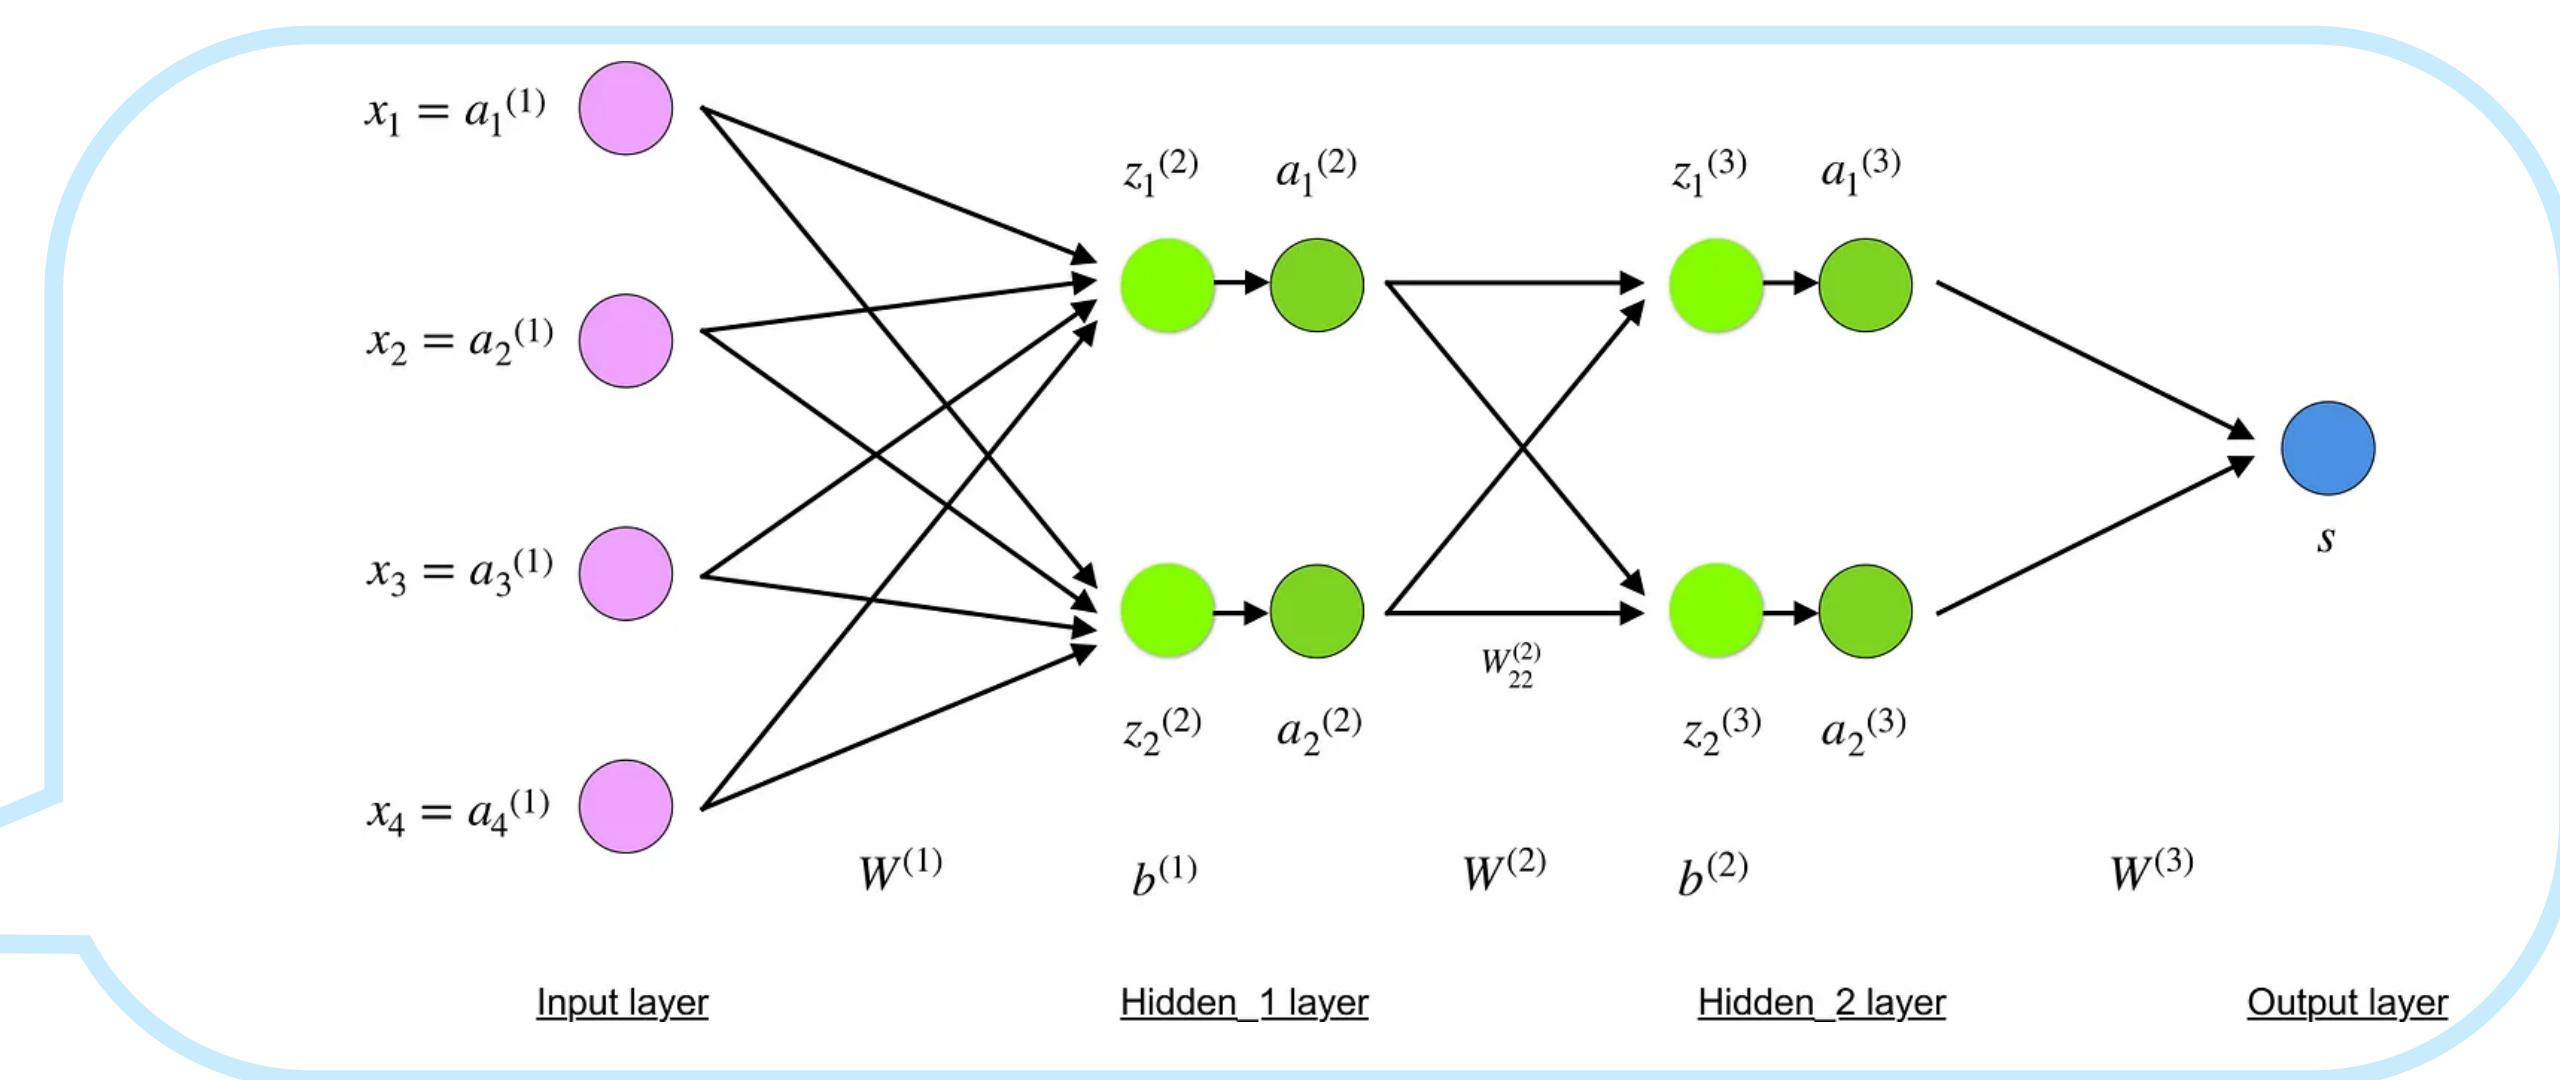
Introduction

Model

Convolutional Neural Network [1]



Fully connected layer [2]



[1] -<https://www.google.com/url?sa=i&url=https%3A%2F%2Fmedium.com%2F40mayankverma05032001%2Fbinary-classification-using-convolution-neural-network-cnn-model-6e35cdf5bdb&psig=AOvVaw2y5GRO7aT-Wo1WYOiBYh3j&ust=1714480373134000&source=images&cd=vfe&opi=89978449&ved=oCBQQjhxqFwoTCJCdup6354UDFQAAAAAdAAAAABAE>

[2] -<https://towardsdatascience.com/understanding-backpropagation-algorithm-7bb3aa2f95fd>

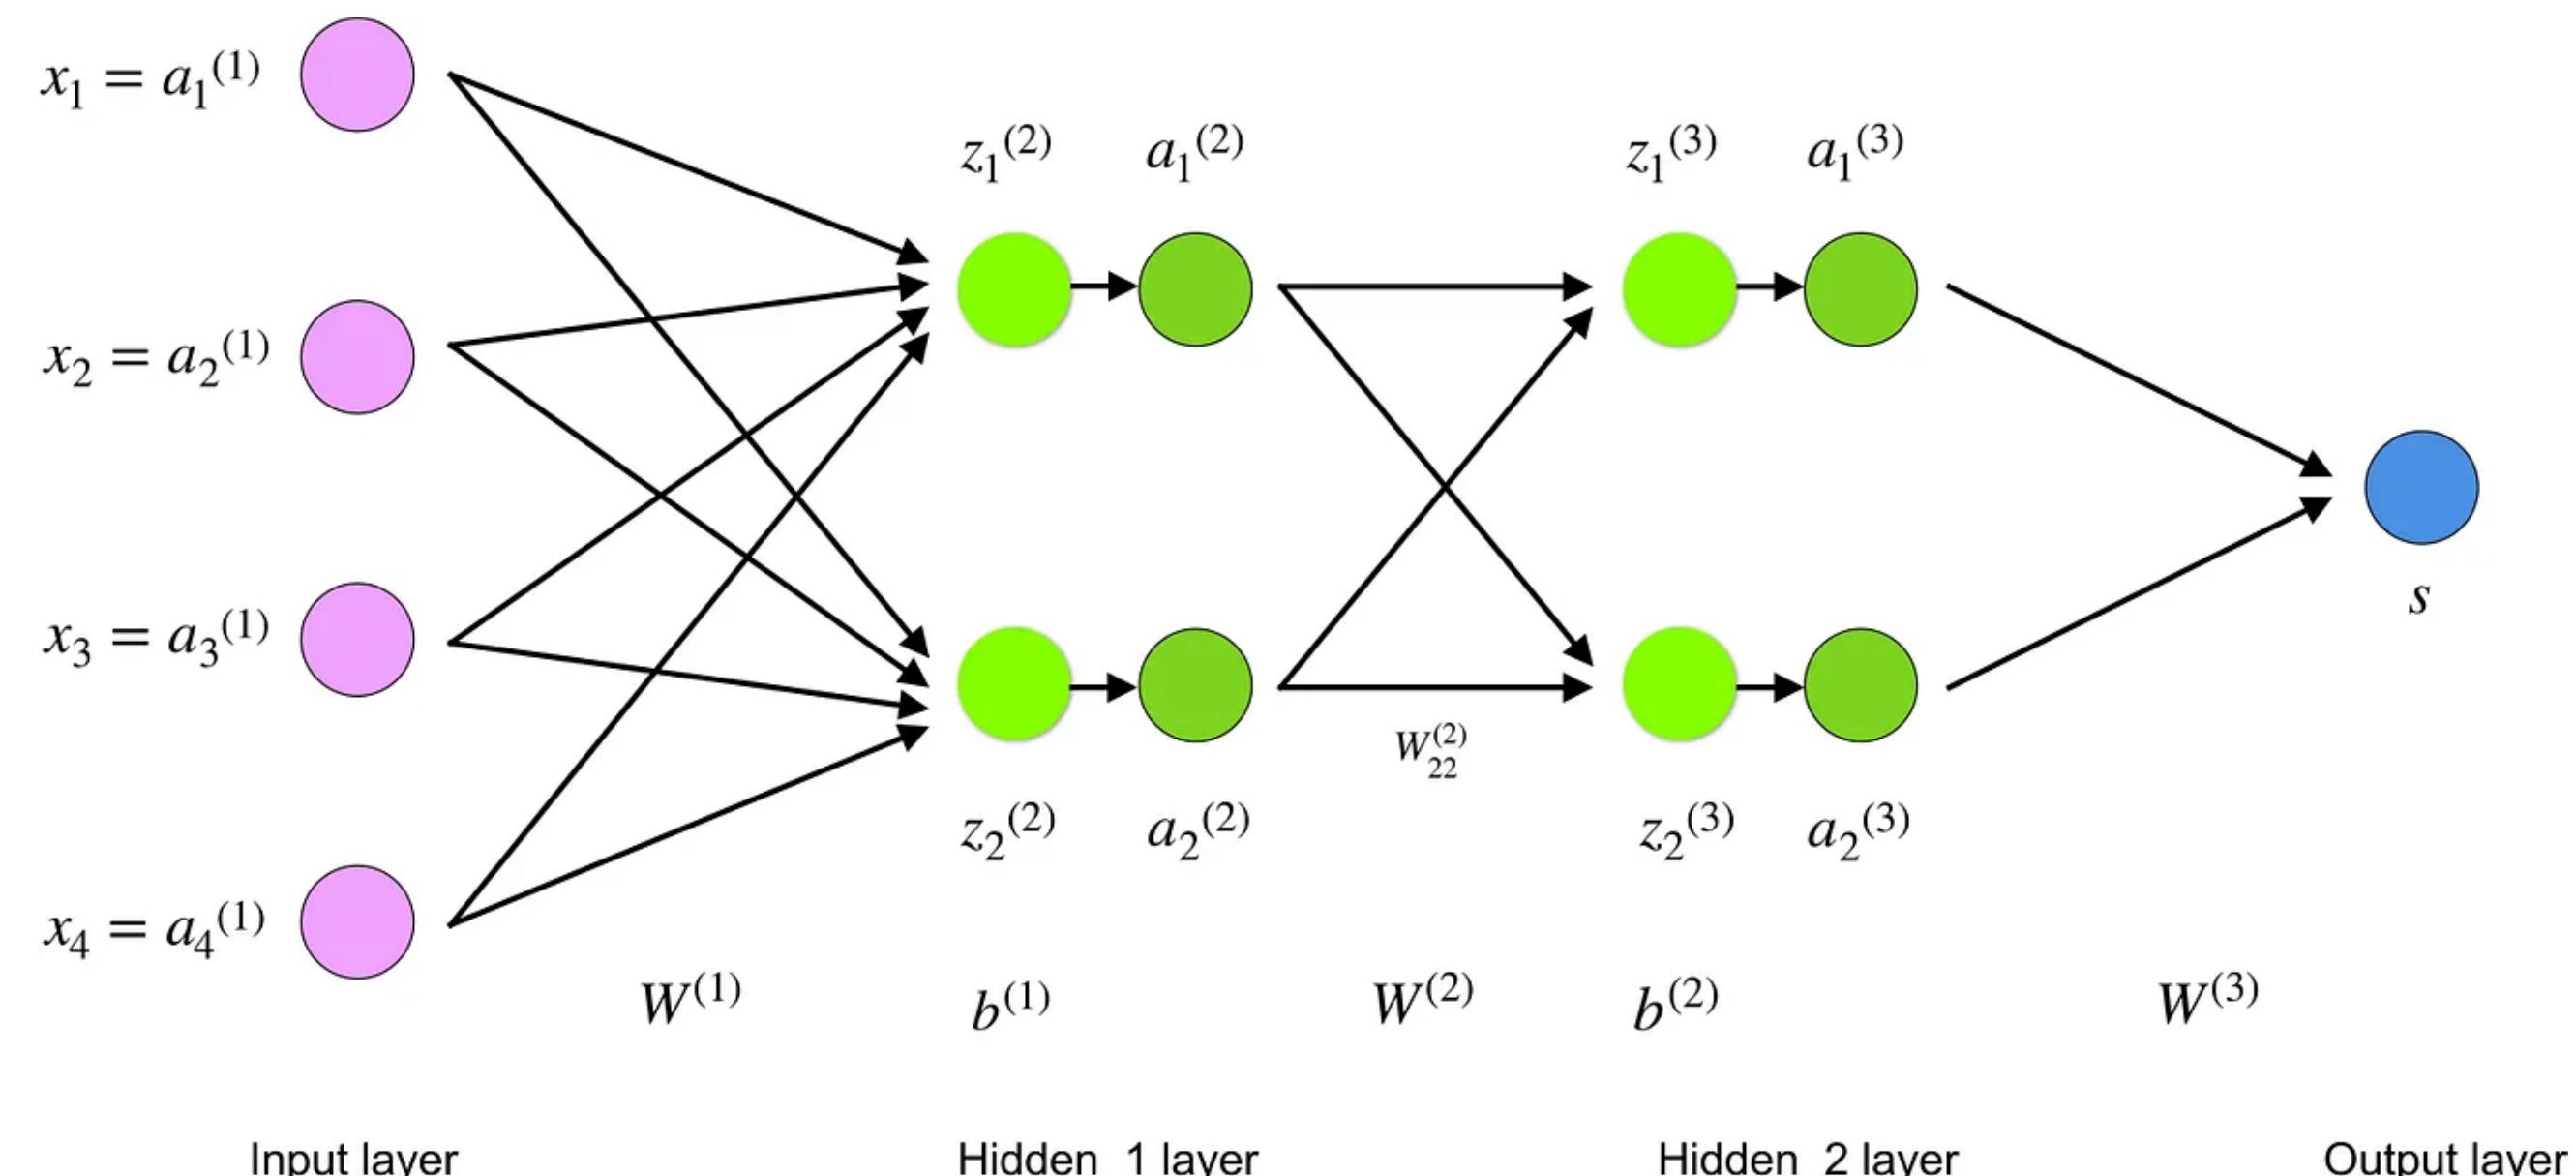
Important

- The proper model for data
- Mindful of the memory and computing power
- **Loss function**

Introduction

Model

Fully connected layer [1]



Input layer

Hidden_1 layer

Hidden_2 layer

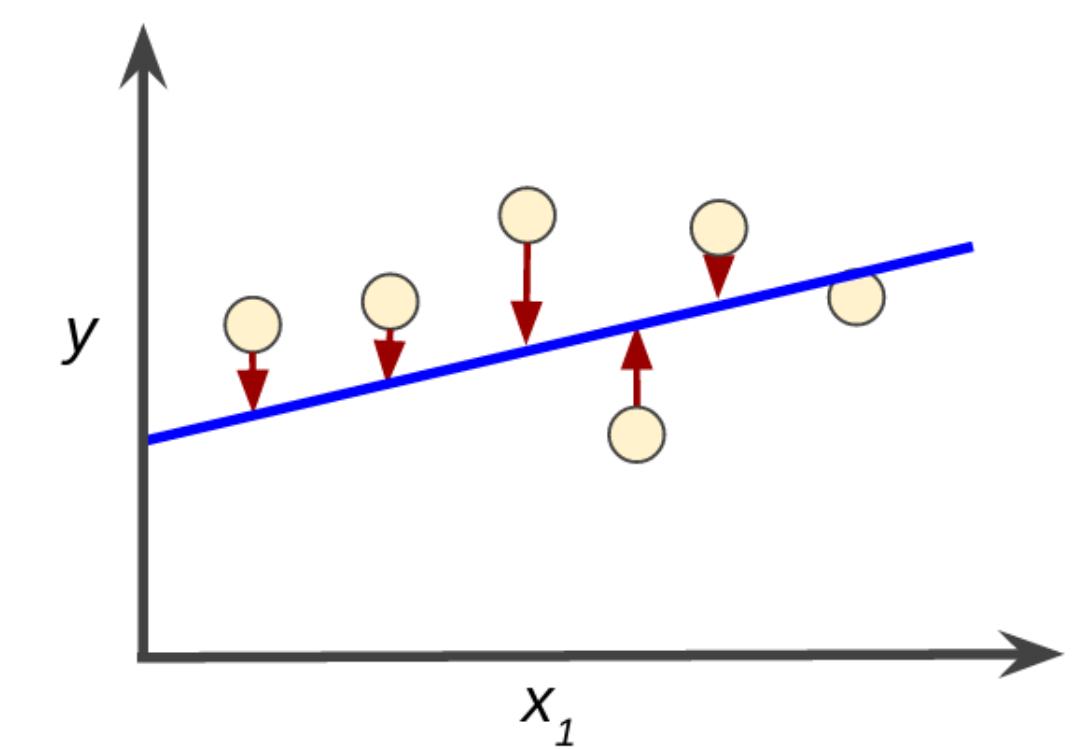
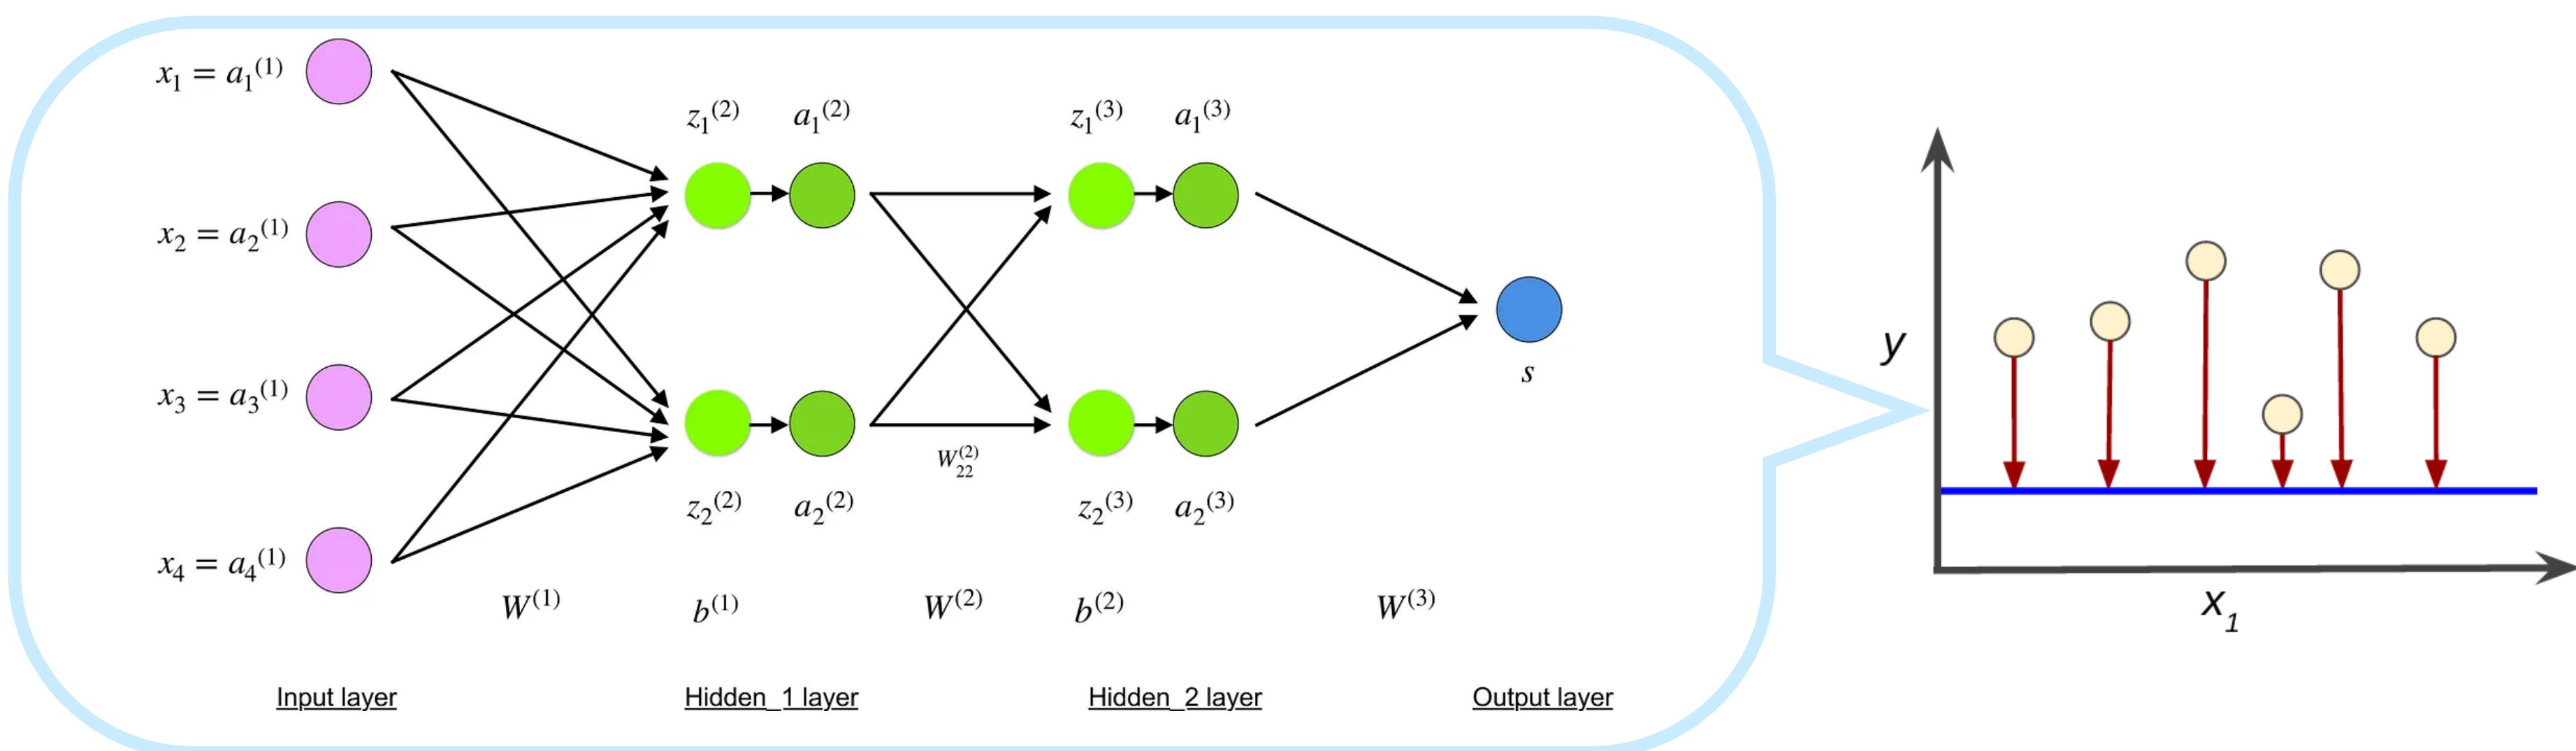
Output layer

Introduction

- Loss function

Fully connected layer [1]

Loss Function [1]



[1] -<https://www.google.com/url?sa=i&url=https%3A%2F%2Fmedium.com%2F40mayankverma05032001%2Fbinary-classification-using-convolution-neural-network-cnn-model-6e35cdf5bdbb&psig=AOvVaw2y5GRO7aT-Wo1WYOiBYh3j&ust=1714480373134000&source=images&cd=vfe&opi=89978449&ved=0CBQQjhxqFwoTCJCdup6354UDFQAAAAAdAAAAABAE>

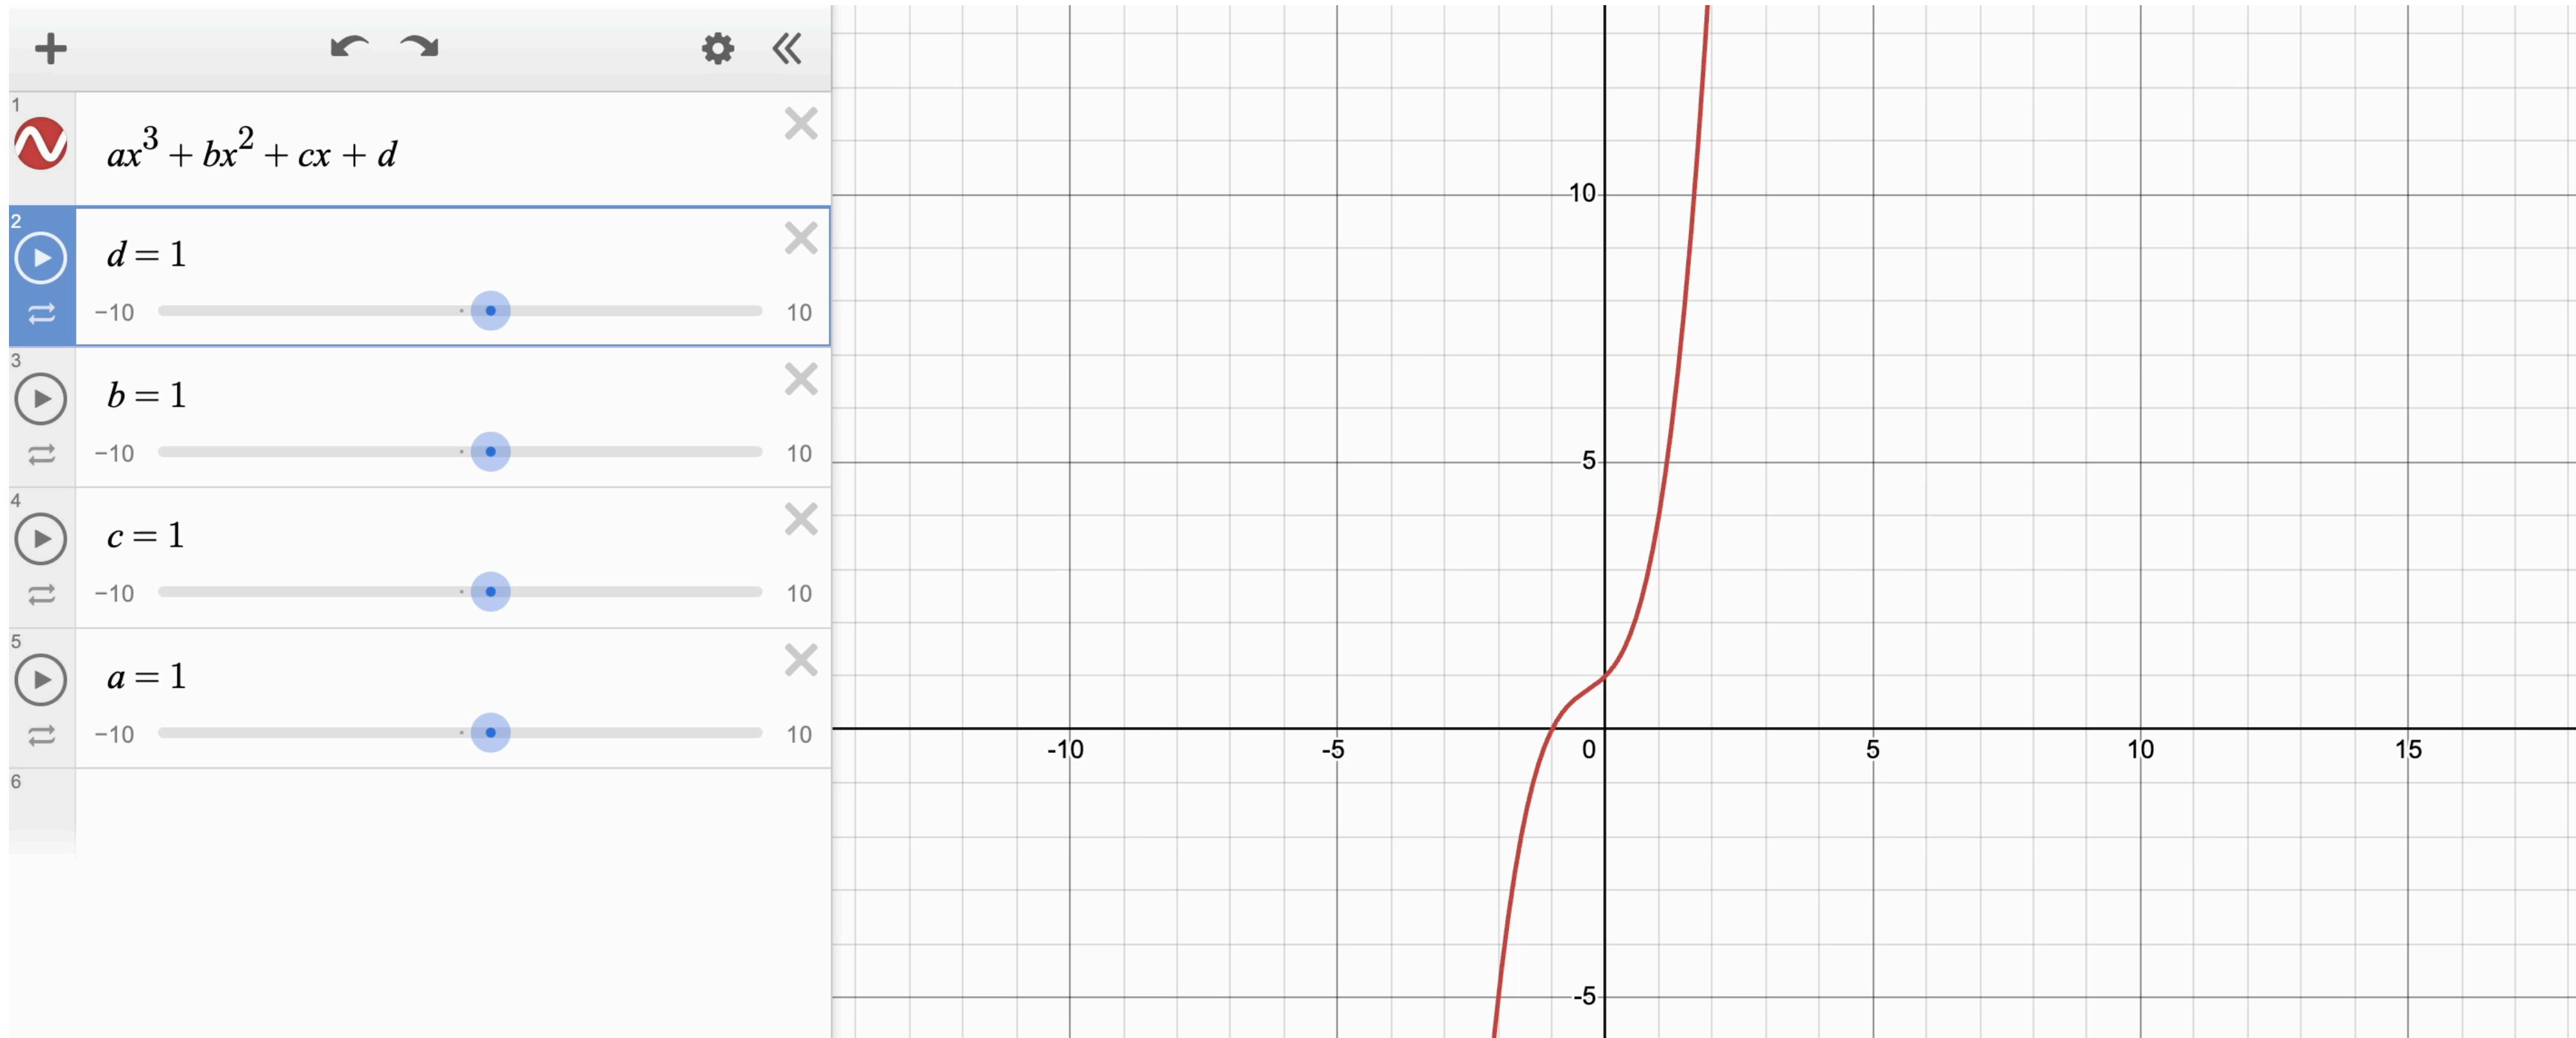
[2] -<https://developers.google.com/machine-learning/crash-course/descending-into-ml/training-and-loss>

Machine learning mechanics

- Back-propagation of Error and Gradient Descent
- Activation function
- Convolutional Neural Networks
- Loss function

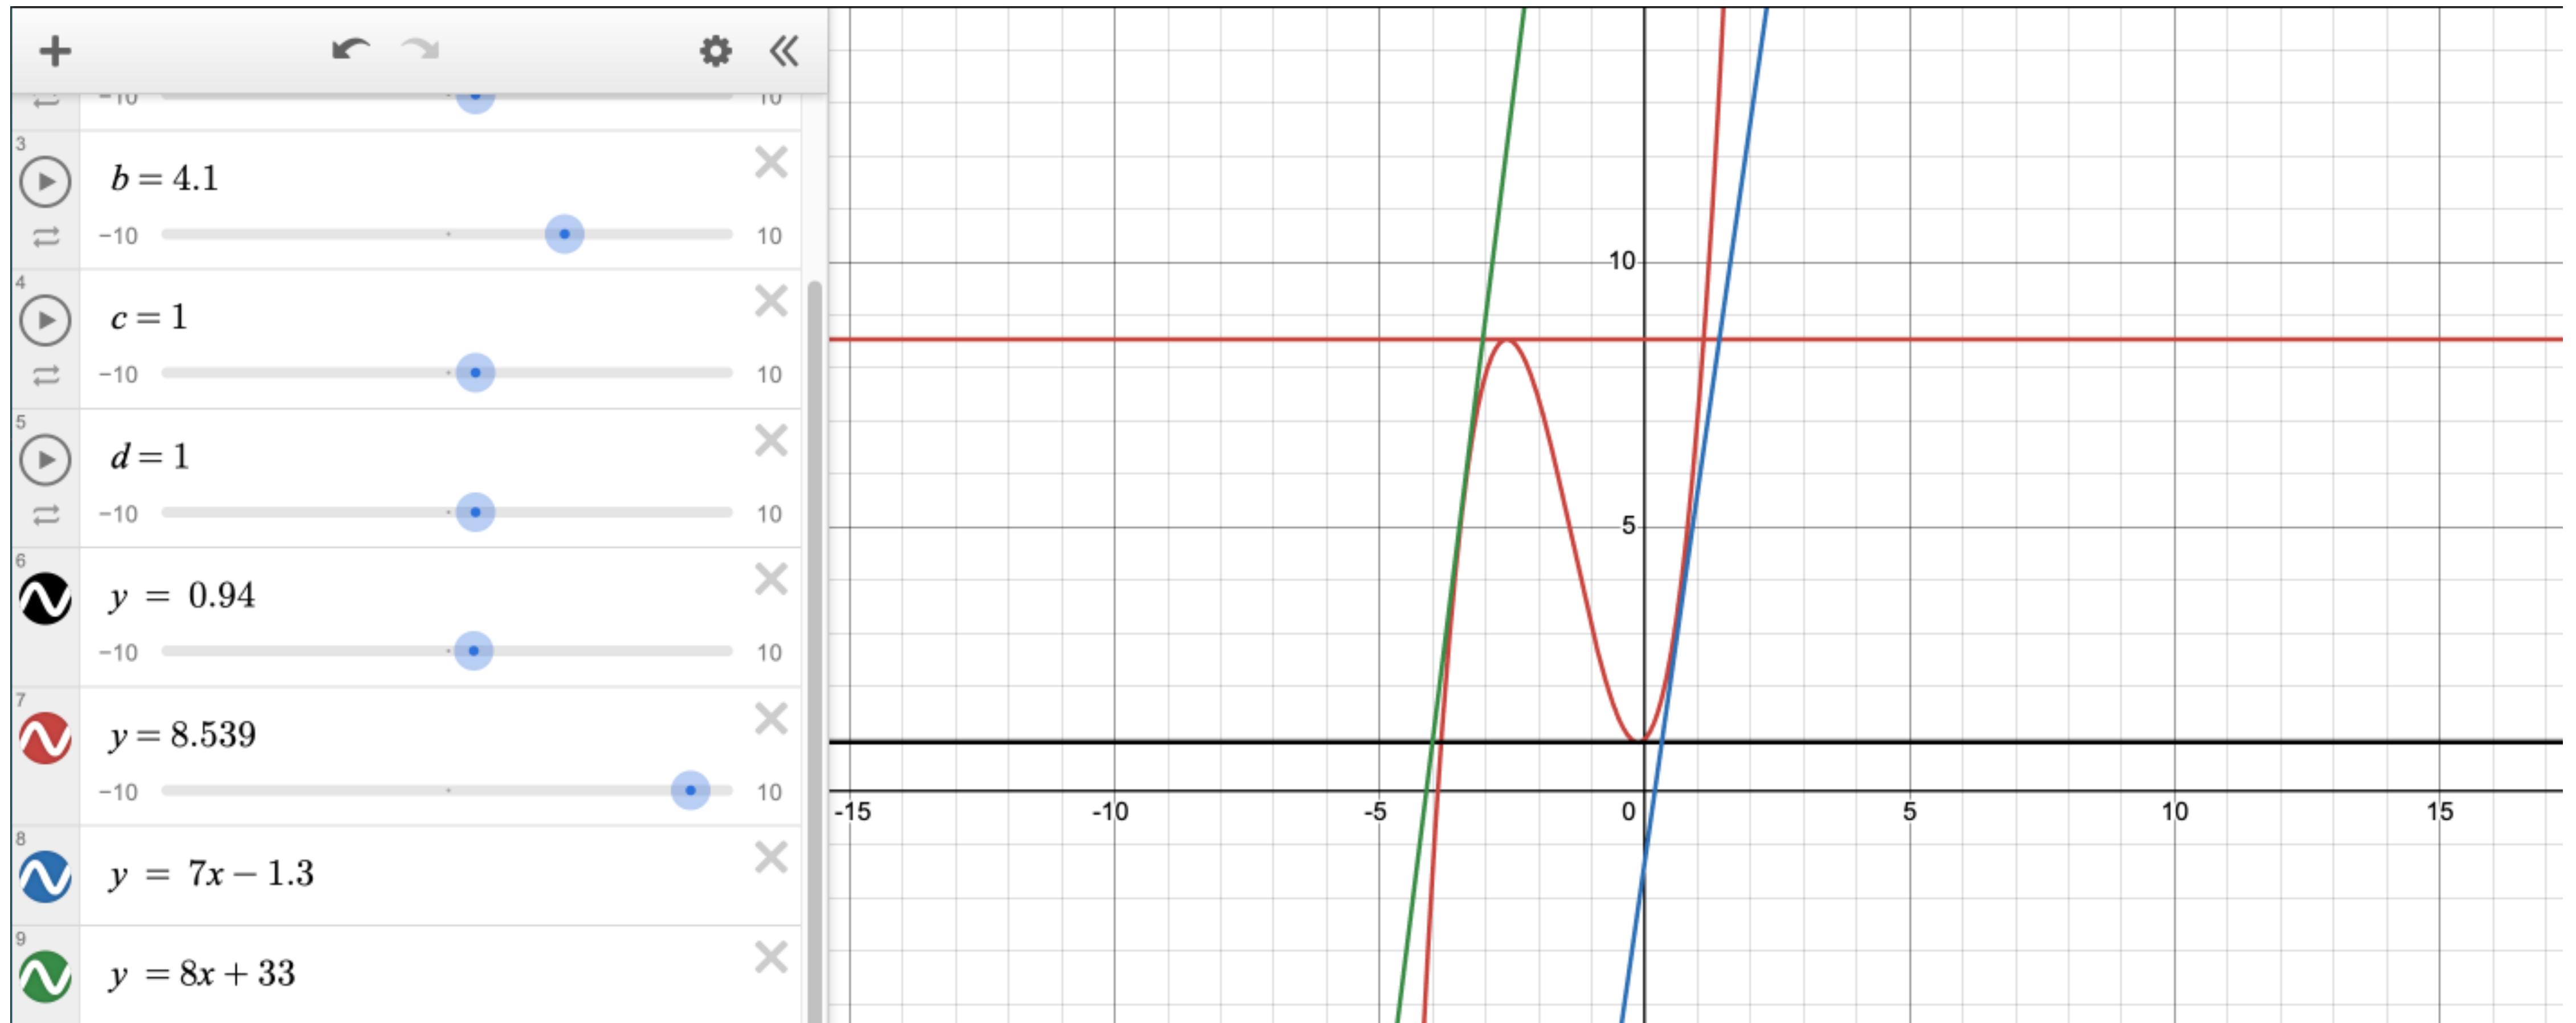
Machine learning mechanics

Back-propagation of Error and Gradient Descent



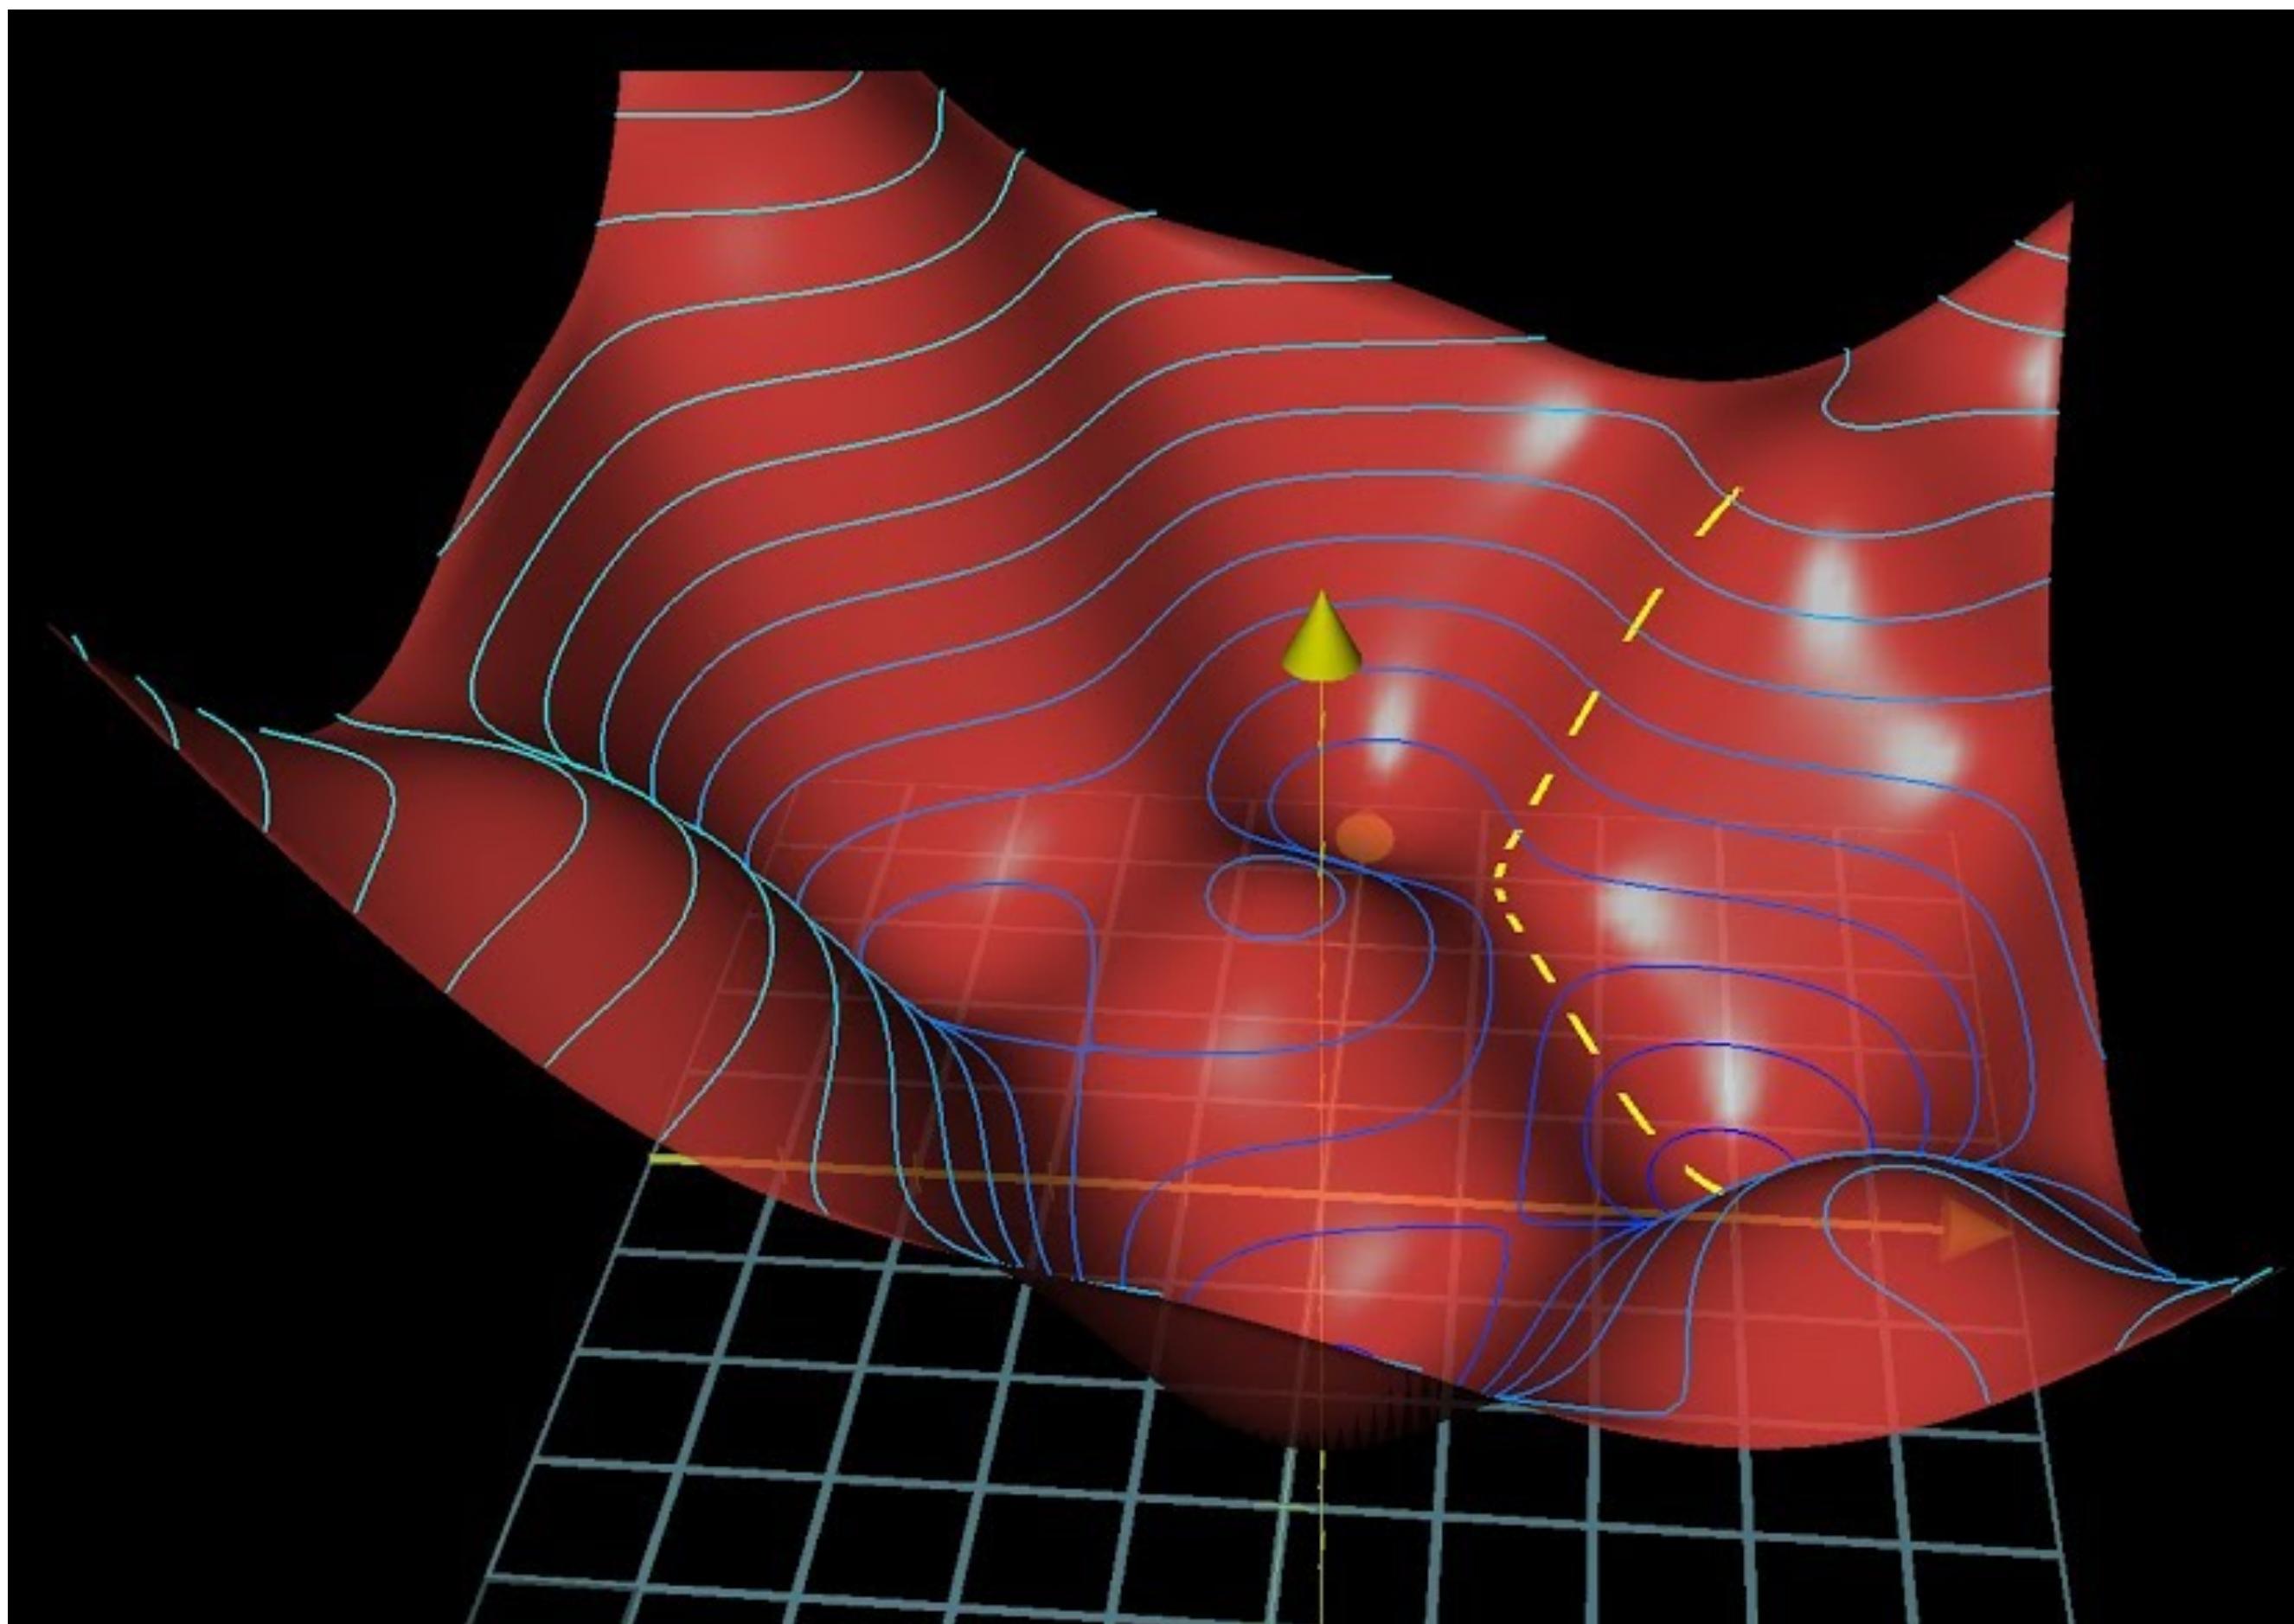
Machine learning mechanics

Back-propagation of Error and Gradient Descent



Machine learning mechanics

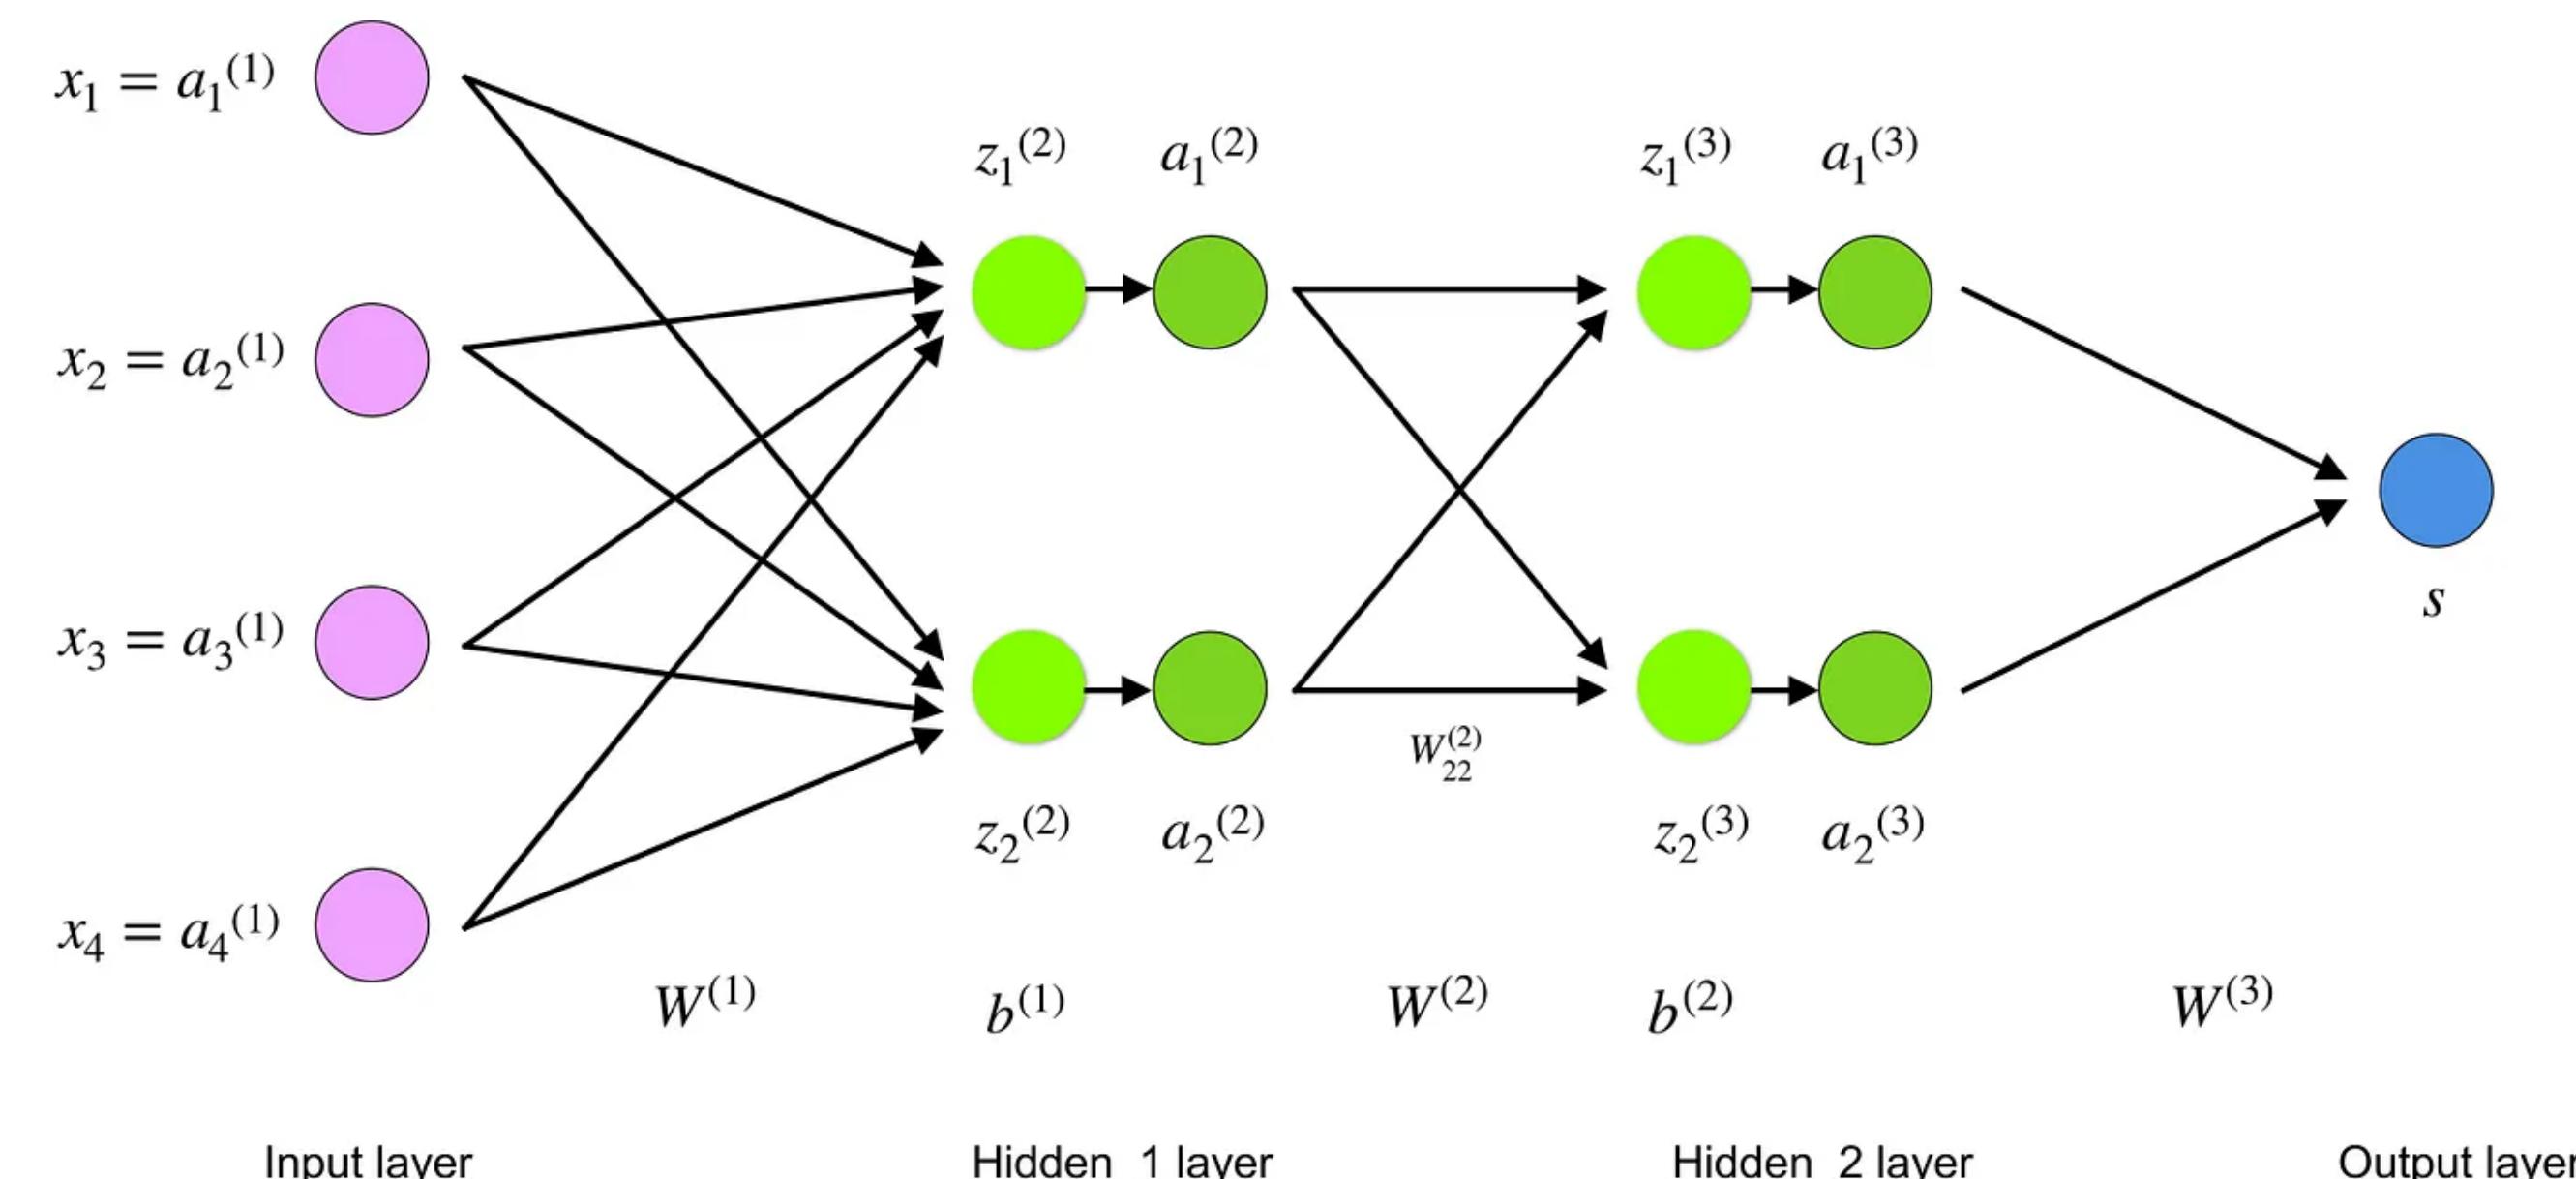
- Gradient Descent



Machine learning mechanics

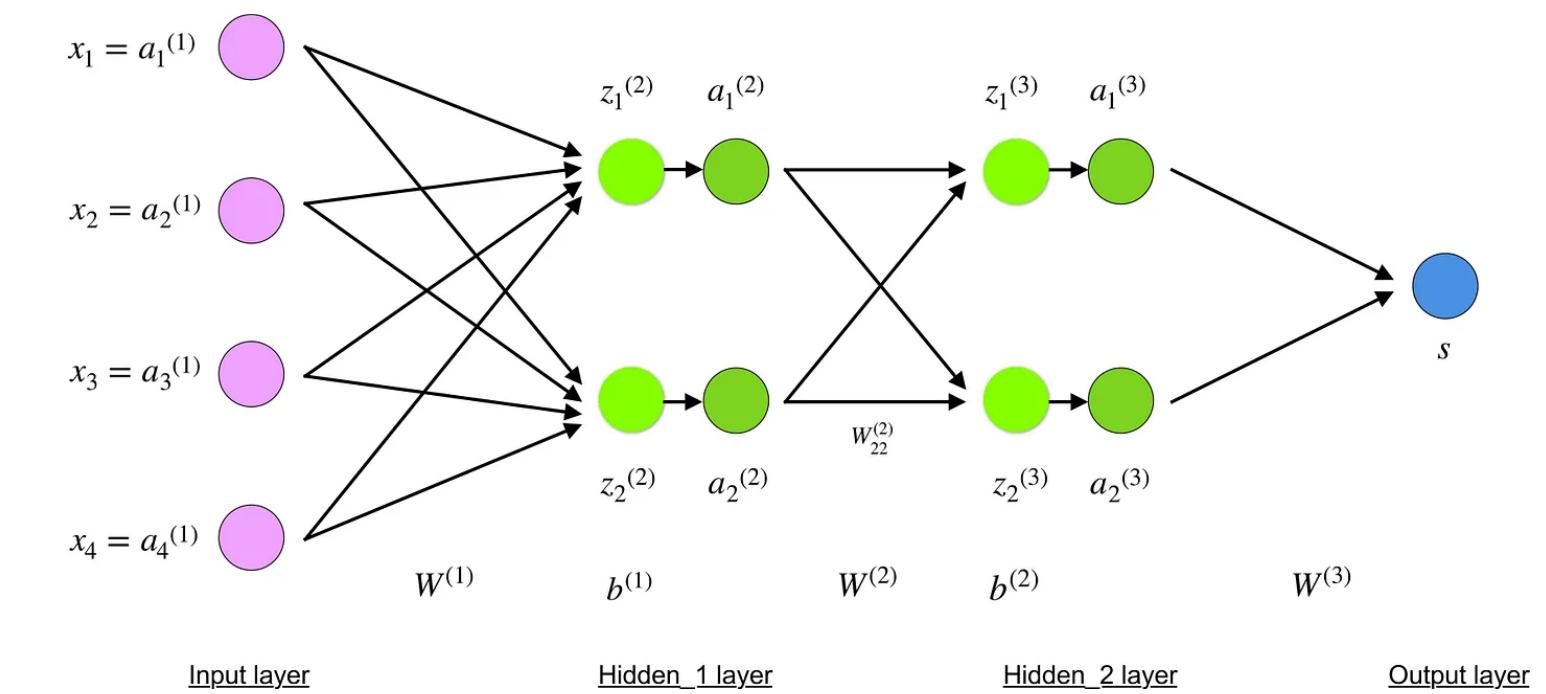
Back-propagation of Error and Gradient Descent

Fully connected layer [1]



Machine learning mechanics

Back-propagation of Error and Gradient Descent



$$W^{(1)} = \begin{bmatrix} W_{11}^{(1)} & W_{12}^{(1)} & W_{13}^{(1)} & W_{14}^{(1)} \\ W_{21}^{(1)} & W_{22}^{(1)} & W_{23}^{(1)} & W_{24}^{(1)} \end{bmatrix} \quad x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} \quad b^{(1)} = \begin{bmatrix} b_1^{(1)} \\ b_2^{(1)} \end{bmatrix}$$

$$z^{(2)} = \begin{bmatrix} W_{11}^{(1)}x_1 + W_{12}^{(1)}x_2 + W_{13}^{(1)}x_3 + W_{14}^{(1)}x_4 \\ W_{21}^{(1)}x_1 + W_{22}^{(1)}x_2 + W_{23}^{(1)}x_3 + W_{24}^{(1)}x_4 \end{bmatrix} + \begin{bmatrix} b_1^{(1)} \\ b_2^{(1)} \end{bmatrix}$$

Machine learning mechanics

Forward Propagation

$x = a^{(1)}$ *Input layer*

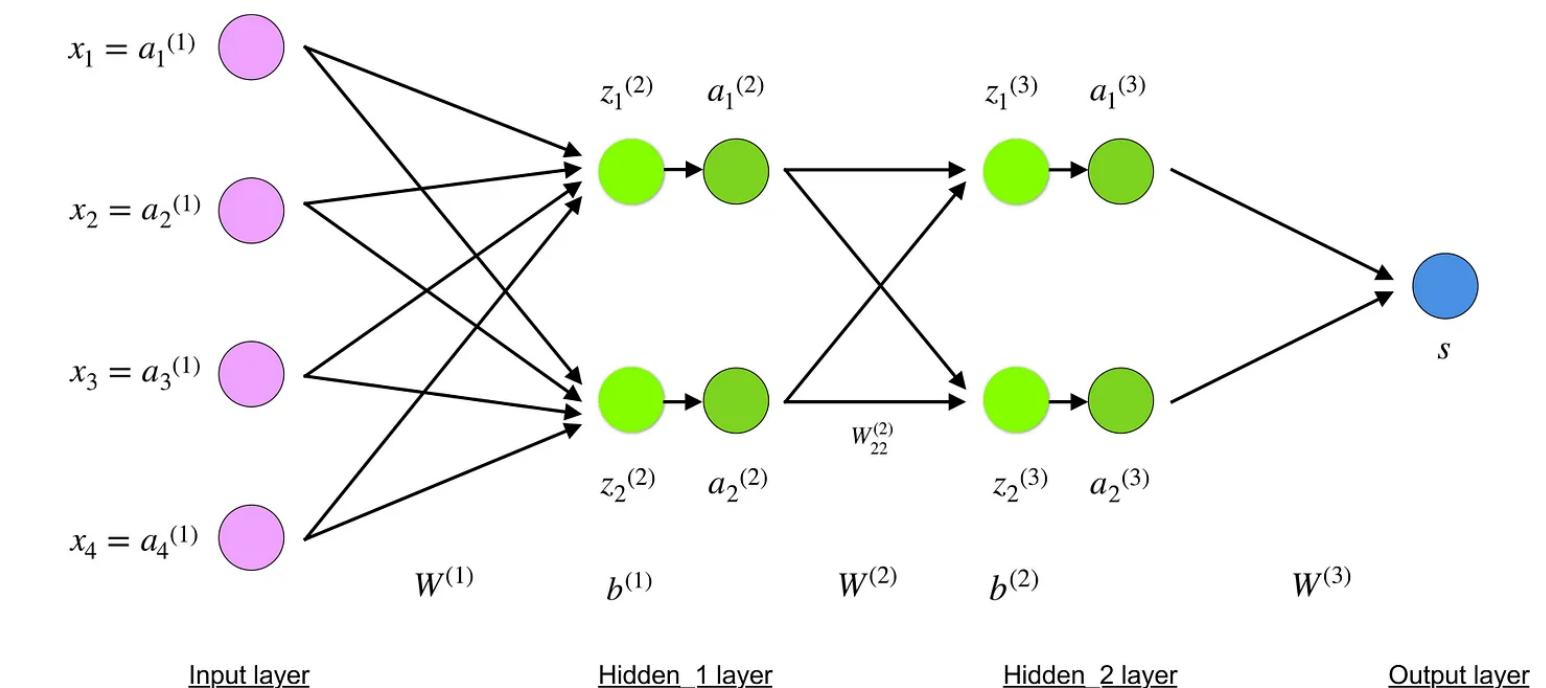
$z^{(2)} = W^{(1)}x + b^{(1)}$ *neuron value at Hidden₁ layer*

$a^{(2)} = f(z^{(2)})$ *activation value at Hidden₁ layer*

$z^{(3)} = W^{(2)}a^{(2)} + b^{(2)}$ *neuron value at Hidden₂ layer*

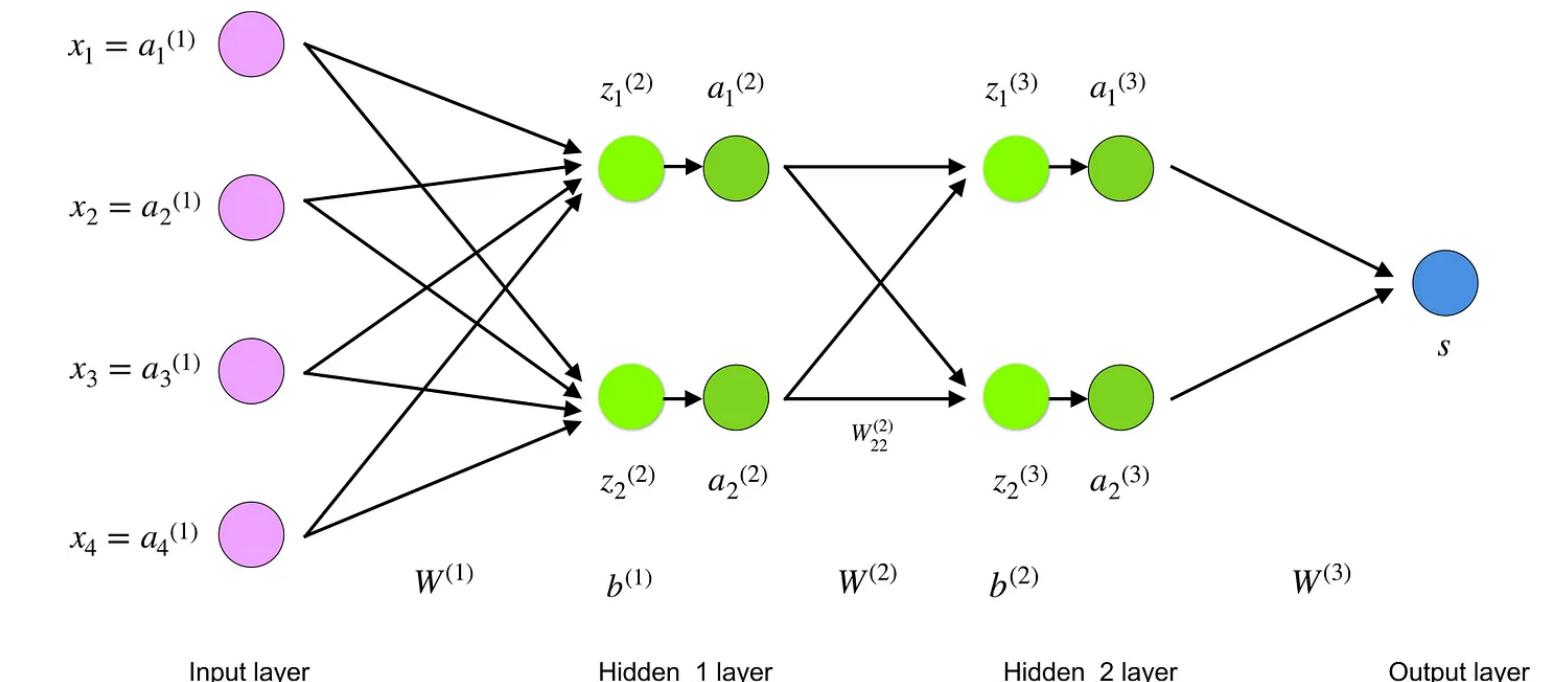
$a^{(3)} = f(z^{(3)})$ *activation value at Hidden₂ layer*

$s = W^{(3)}a^{(3)}$ *Output layer* $\implies C = cost(s, y)$



Machine learning mechanics

Back Propagation



$$z_j^l = \sum_{k=1}^m w_{jk}^l a_k^{l-1} + b_j^l \quad \text{by definition}$$

m – number of neurons in l – 1 layer

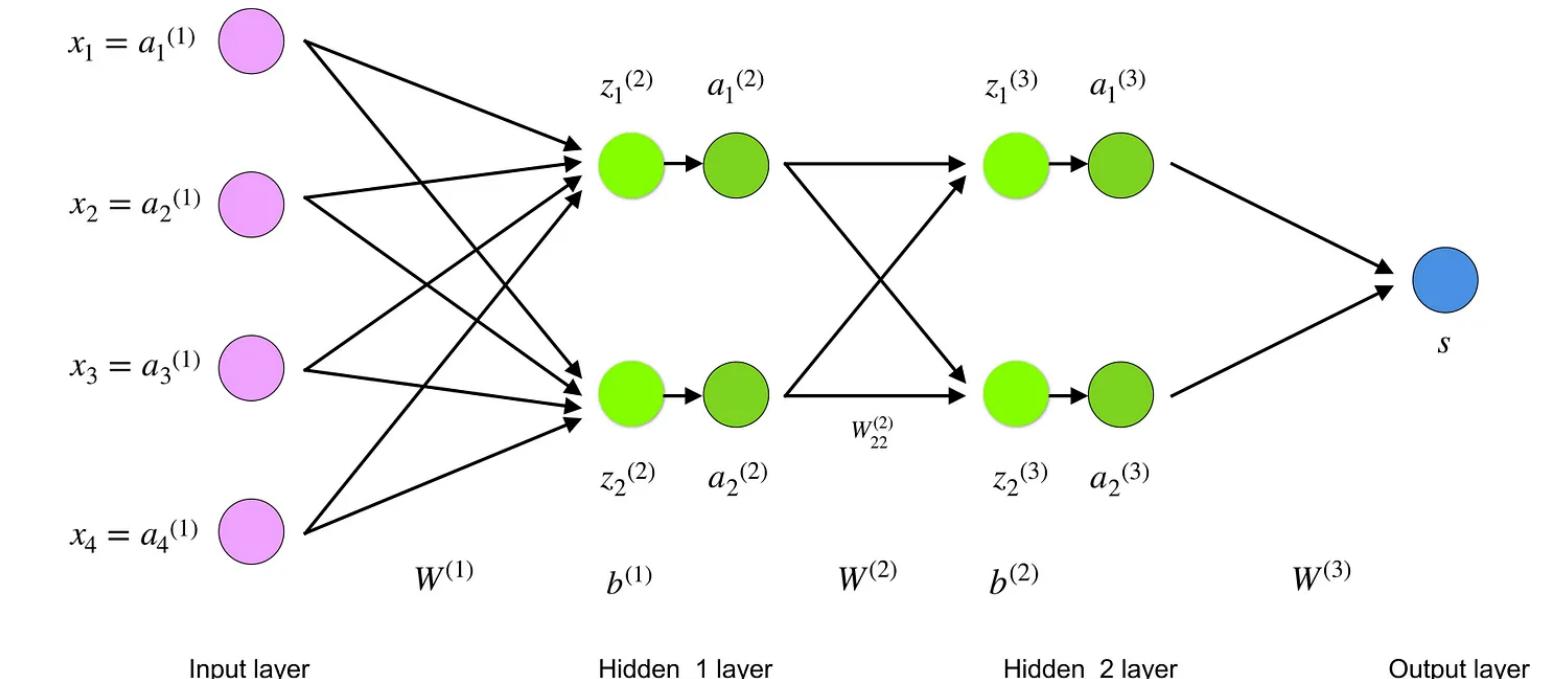
$$\frac{\partial C}{\partial w_{jk}^l} = \frac{\partial C}{\partial z_j^l} \frac{\partial z_j^l}{\partial w_{jk}^l}$$

chain rule $\frac{\partial z_j^l}{\partial w_{jk}^l} = a_k^{l-1}$ *by differentiation (calculating derivative)*

$$\frac{\partial C}{\partial w_{jk}^l} = \frac{\partial C}{\partial z_j^l} a_k^{l-1} \quad \text{final value}$$

Machine learning mechanics

Back Propagation



$$z_j^l = \sum_{k=1}^m w_{jk}^l a_k^{l-1} + b_j^l \quad \text{by definition}$$

m – number of neurons in $l - 1$ layer

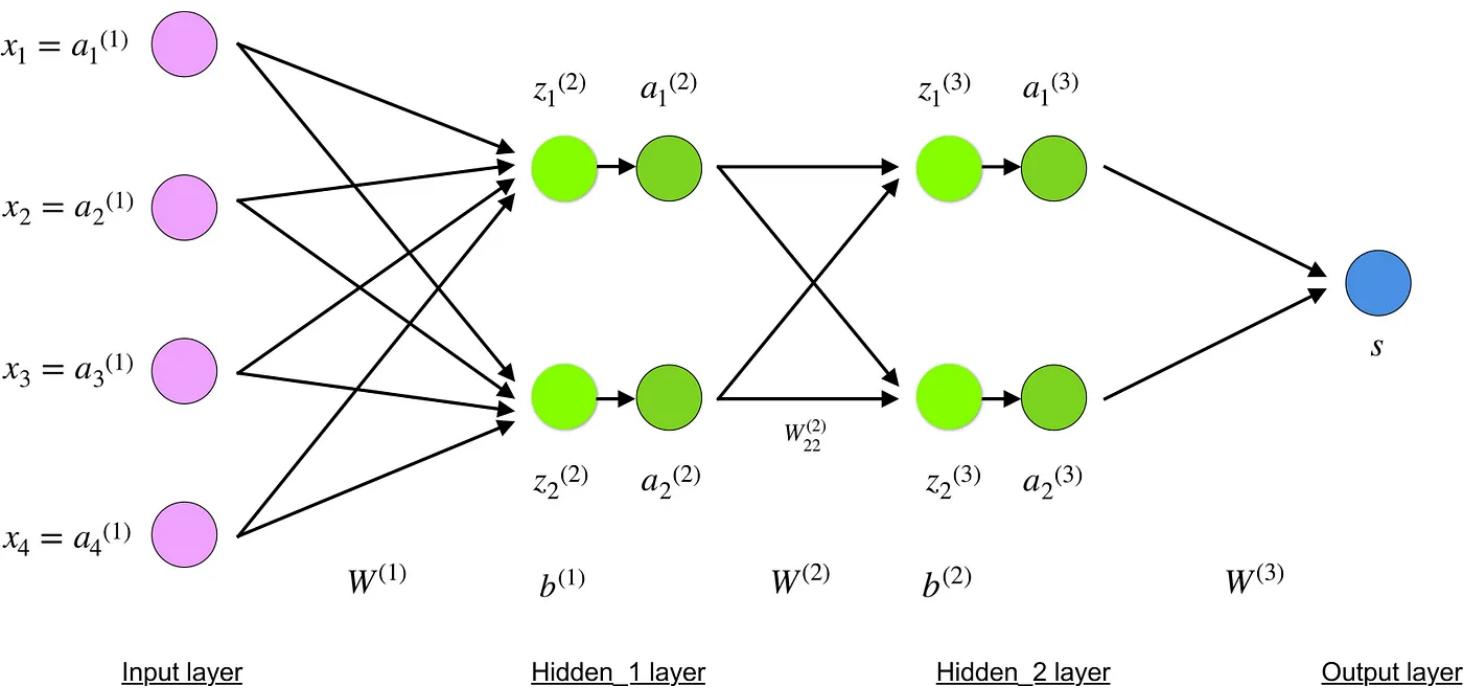
$$\frac{\partial C}{\partial b_j^l} = \frac{\partial C}{\partial z_j^l} \frac{\partial z_j^l}{\partial b_j^l}$$

chain rule $\frac{\partial z_j^l}{\partial b_j^l} = 1$ by differentiation (calculating derivative)

$$\frac{\partial C}{\partial b_j^l} = \frac{\partial C}{\partial z_j^l} 1 \quad \text{final value}$$

Machine learning mechanics

- Back Propagation



while (termination condition not met)

$$w := w - \epsilon \frac{\partial C}{\partial w}$$

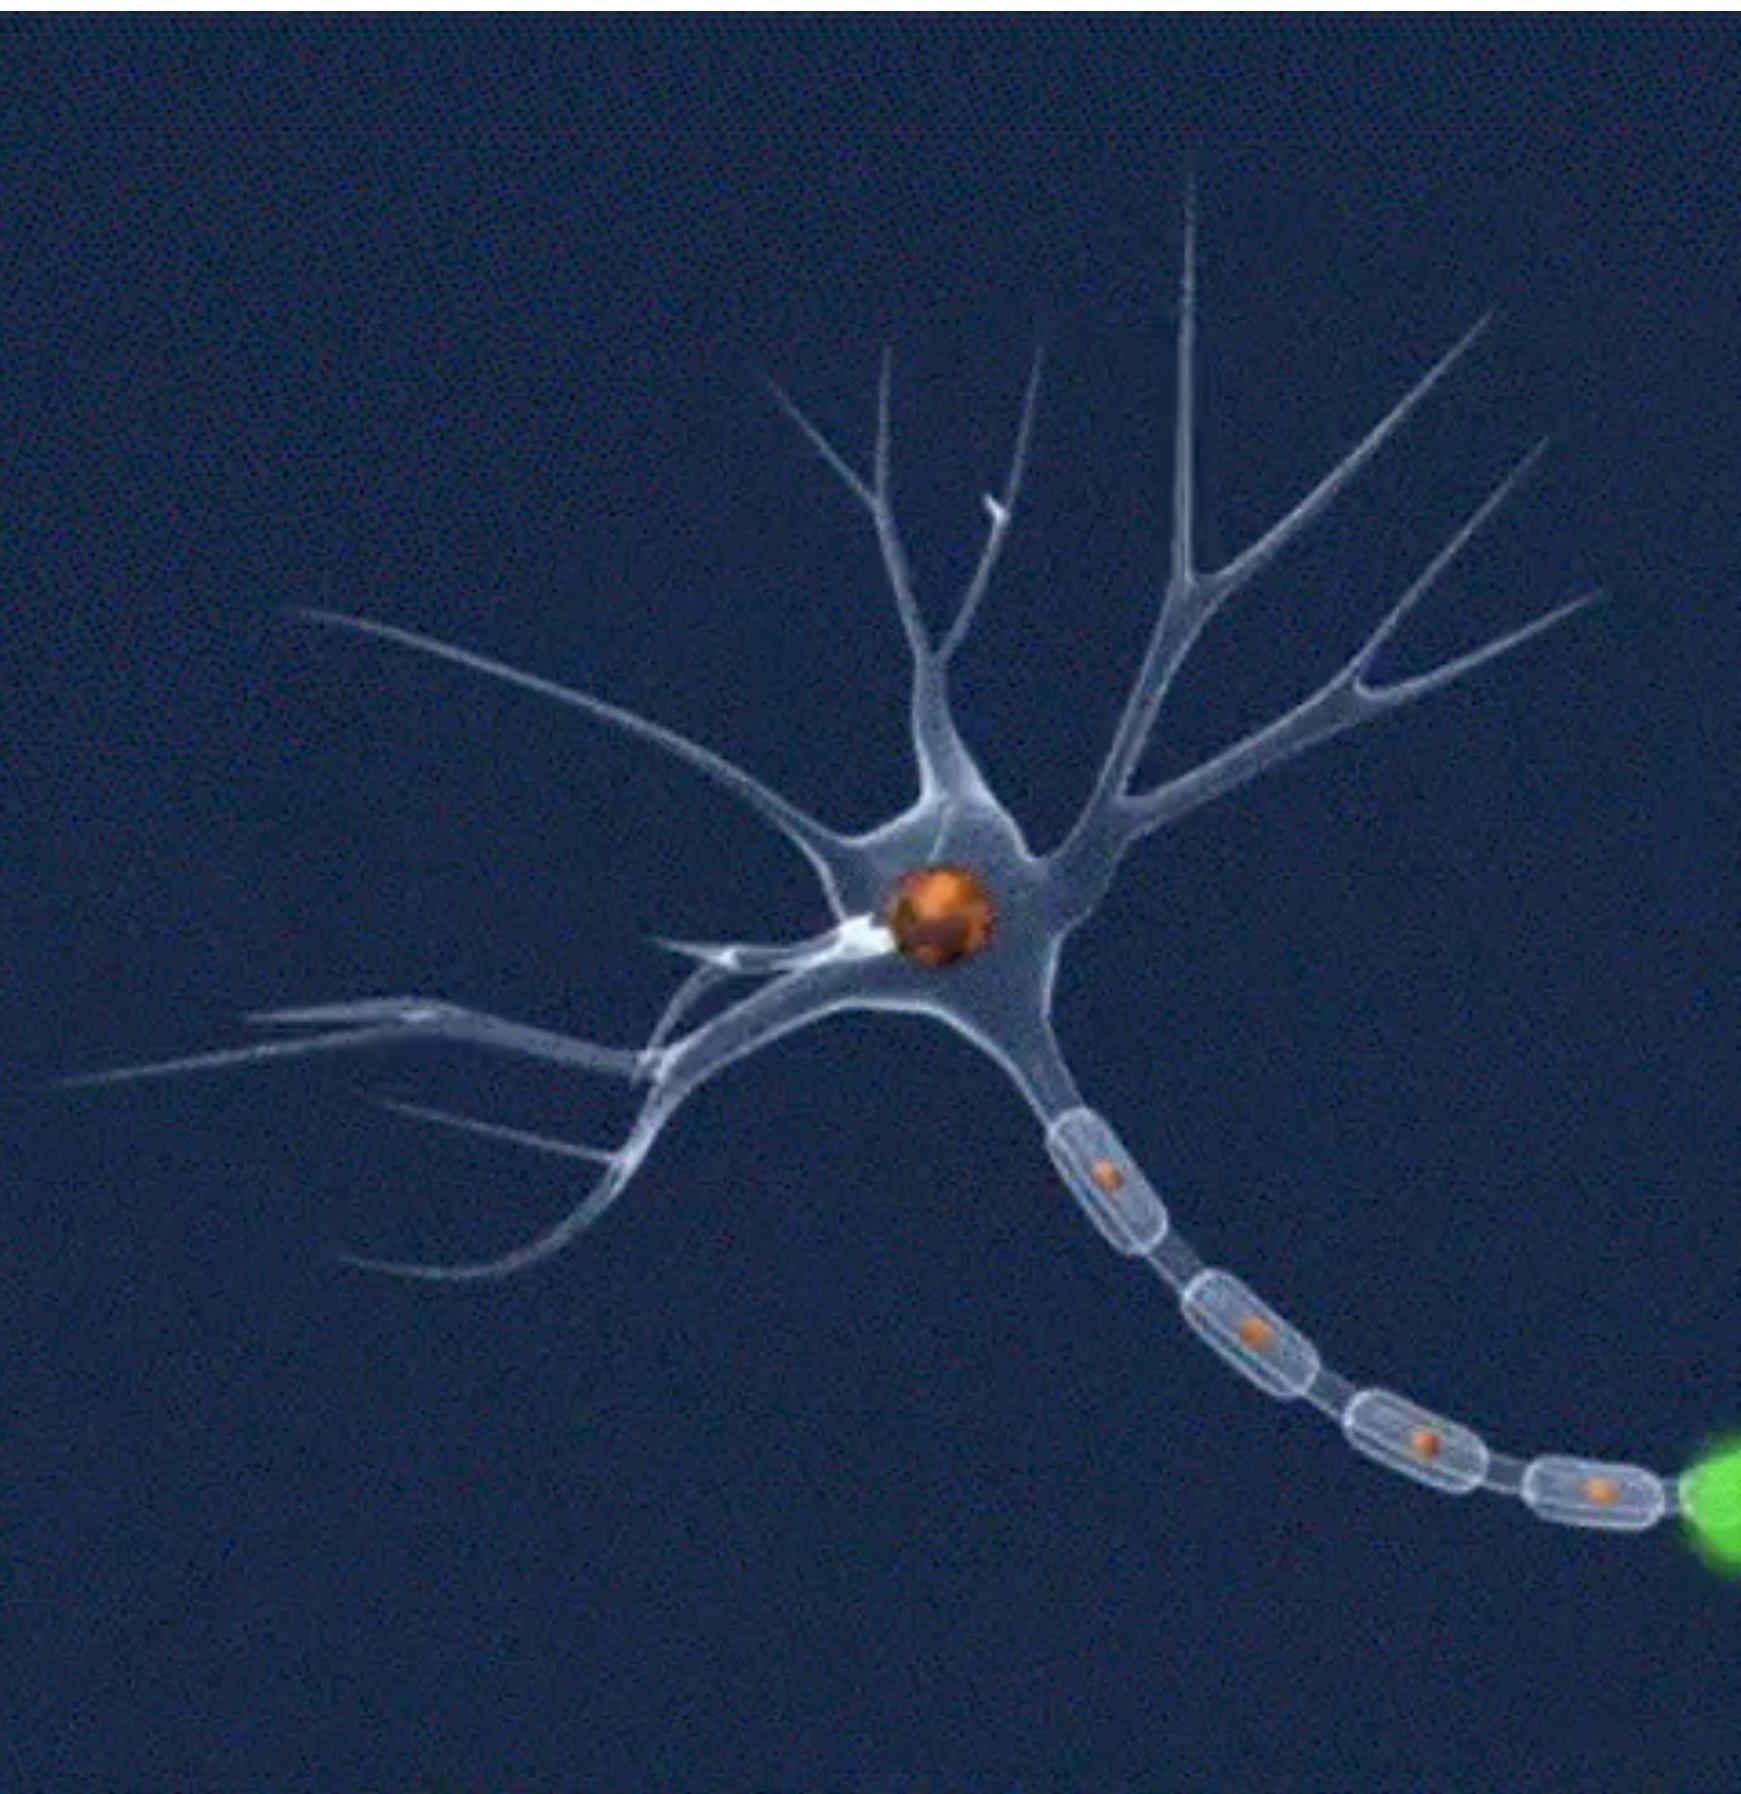
Learning rate ϵ

$$b := b - \epsilon \frac{\partial C}{\partial b}$$

end

Machine learning mechanics

Activation Functions

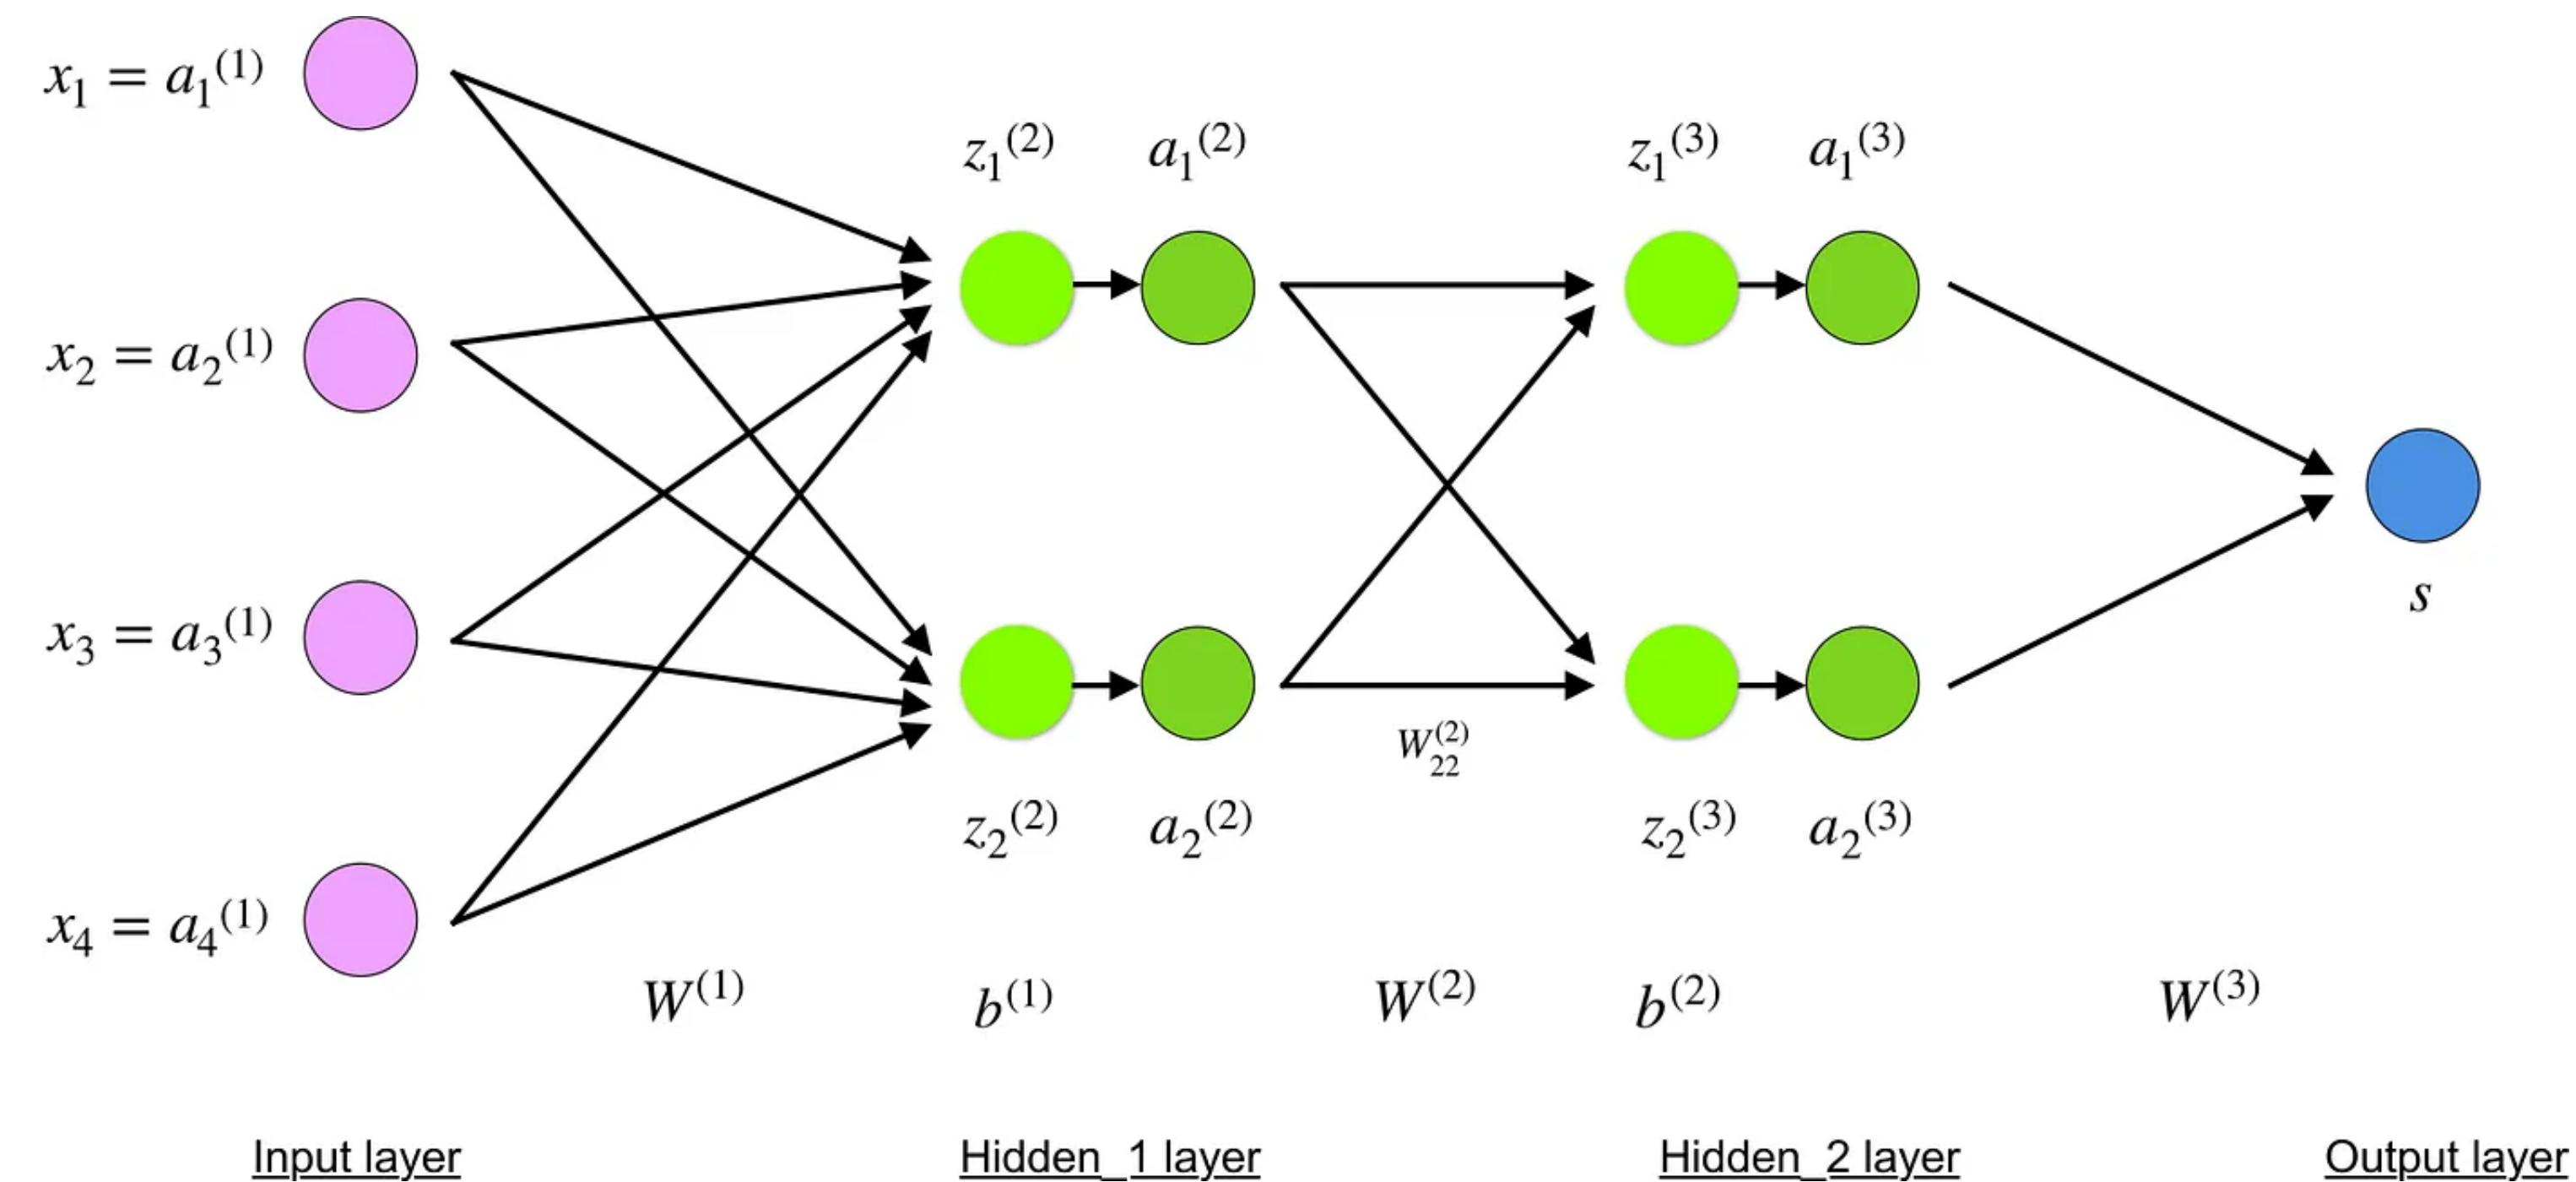


<https://medium.com/the-theory-of-everything/understanding-activation-functions-in-neural-networks-9491262884eo>

<https://www.nichd.nih.gov/newsroom/releases/031813-backwards-neurons>

Machine learning mechanics

Activation Functions

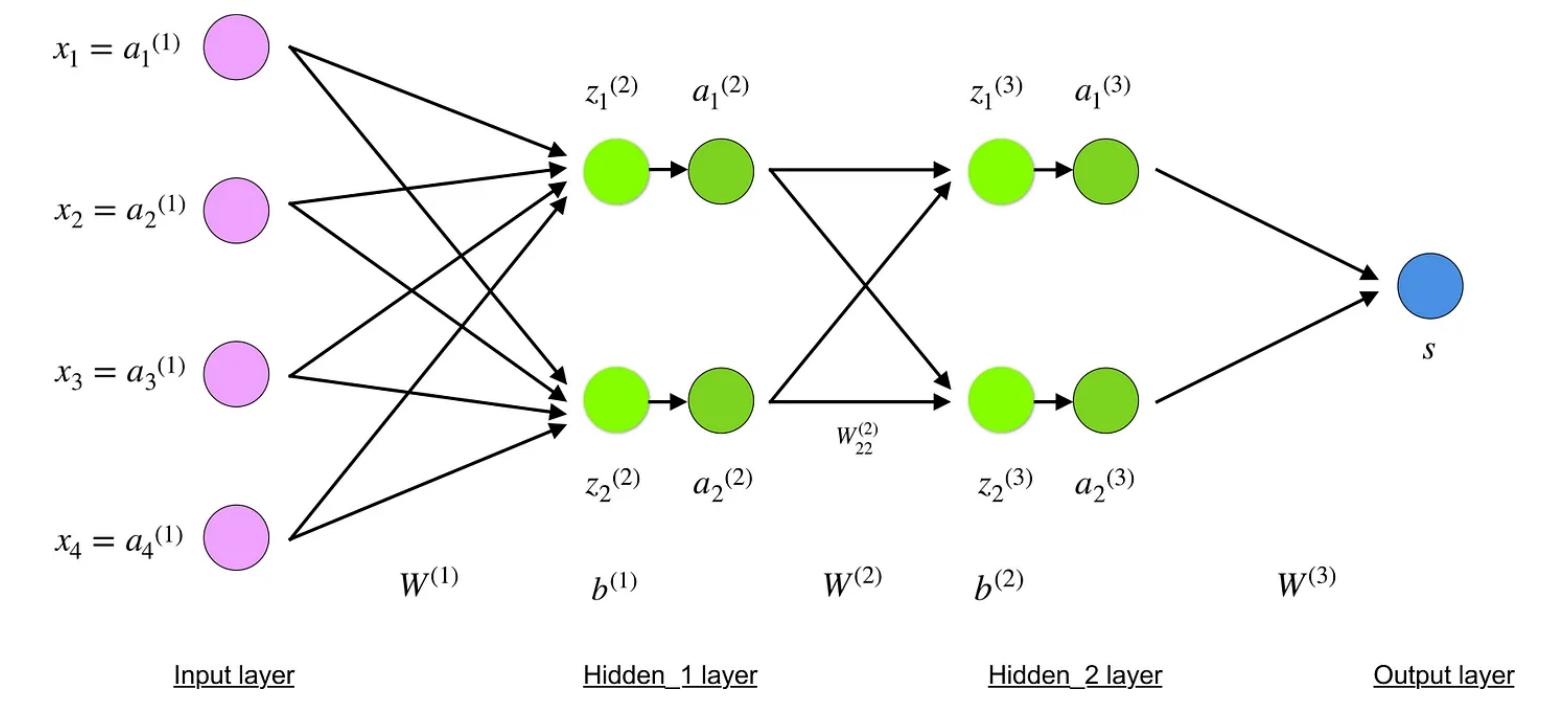


$$Y = \sum (\text{weight} * \text{input}) + \text{bias}$$

Machine learning mechanics

-

Activation Functions

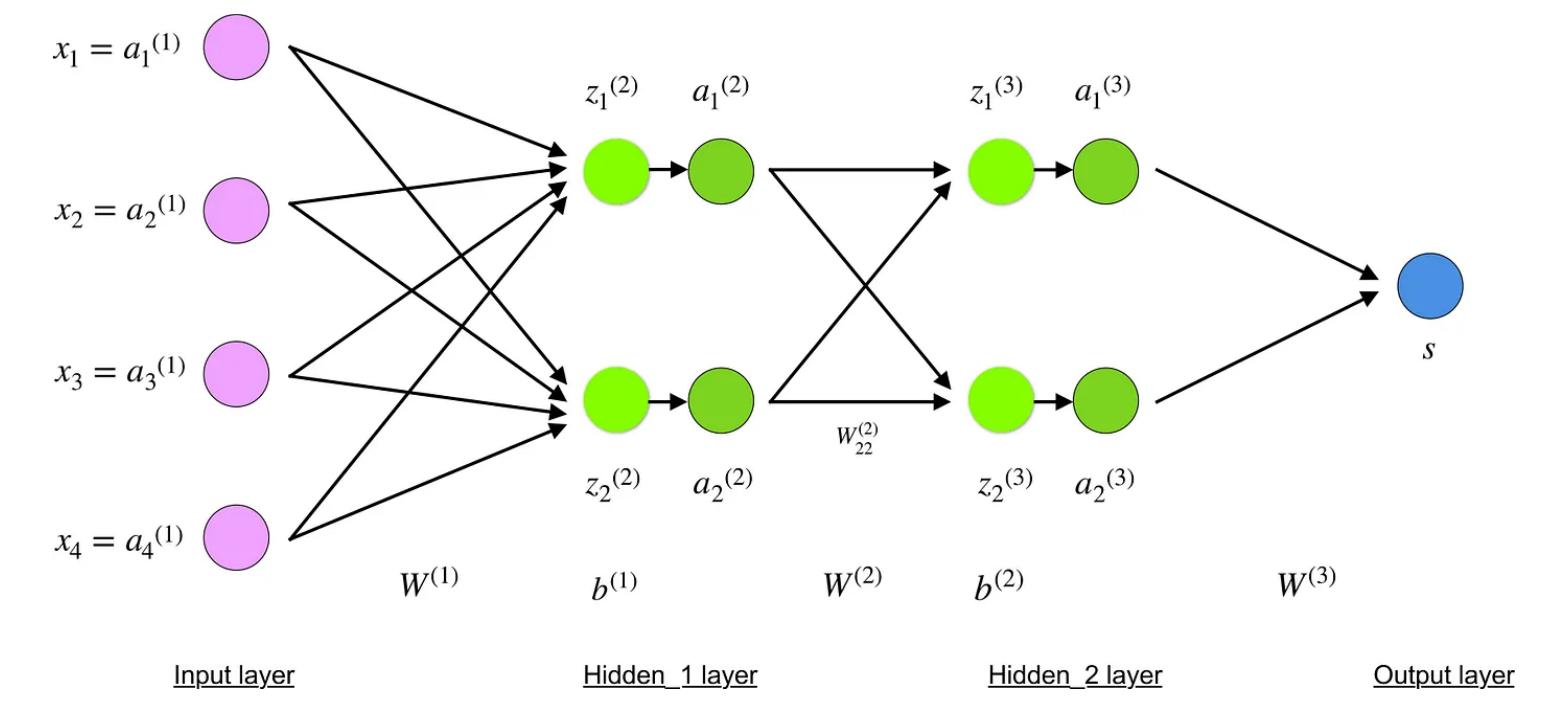


$$Y = \sum (\text{weight} * \text{input}) + \text{bias}$$

Machine learning mechanics

-

Activation Functions

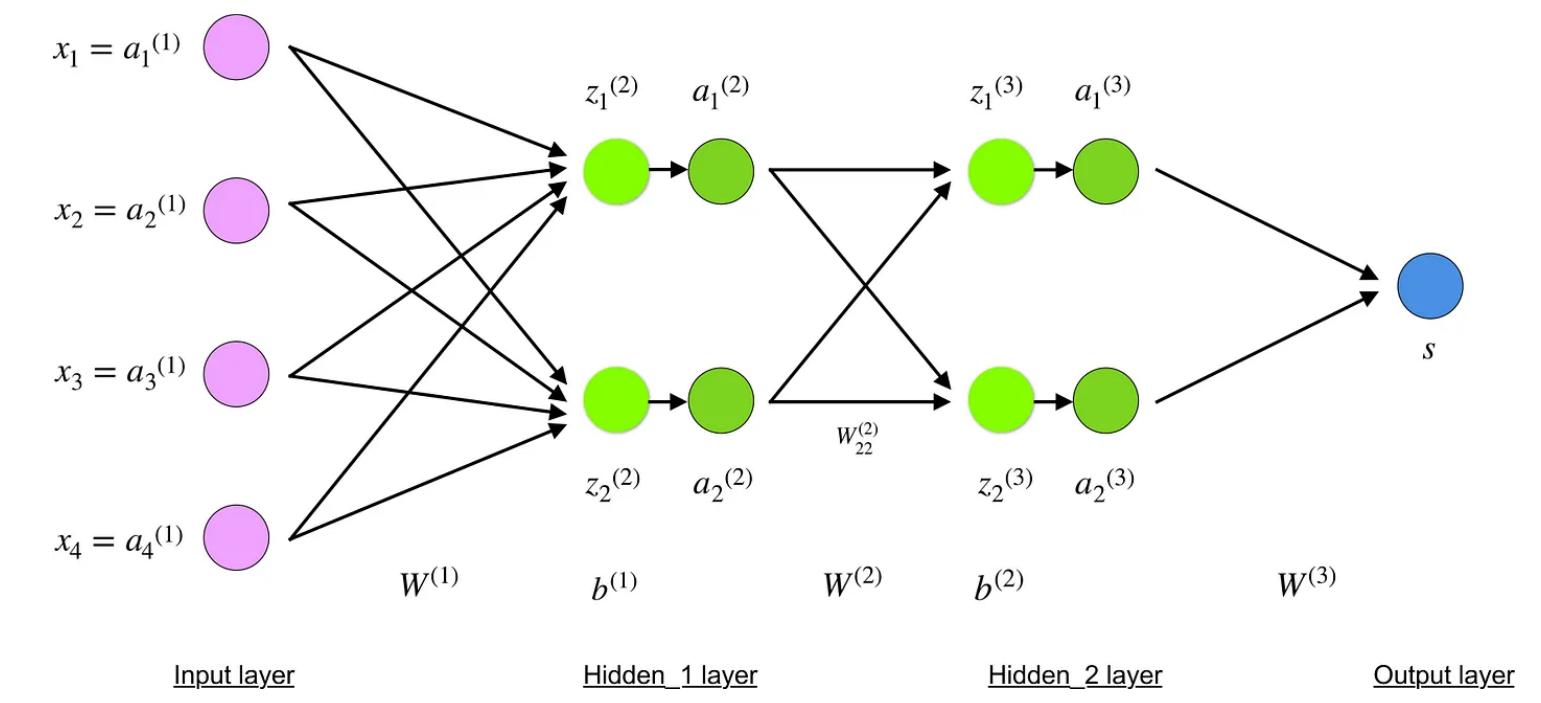


$$Y \in [-\infty; \infty]$$

Machine learning mechanics

-

Activation Functions

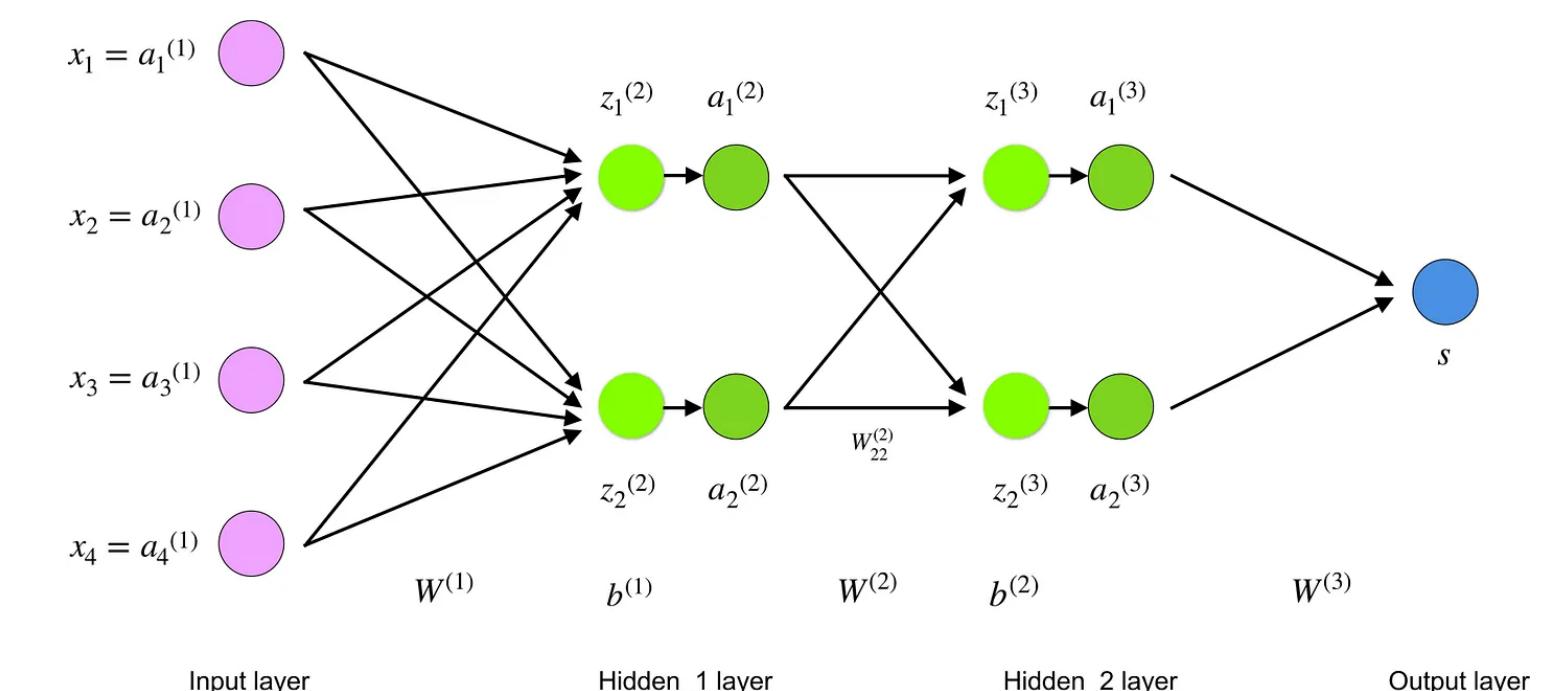
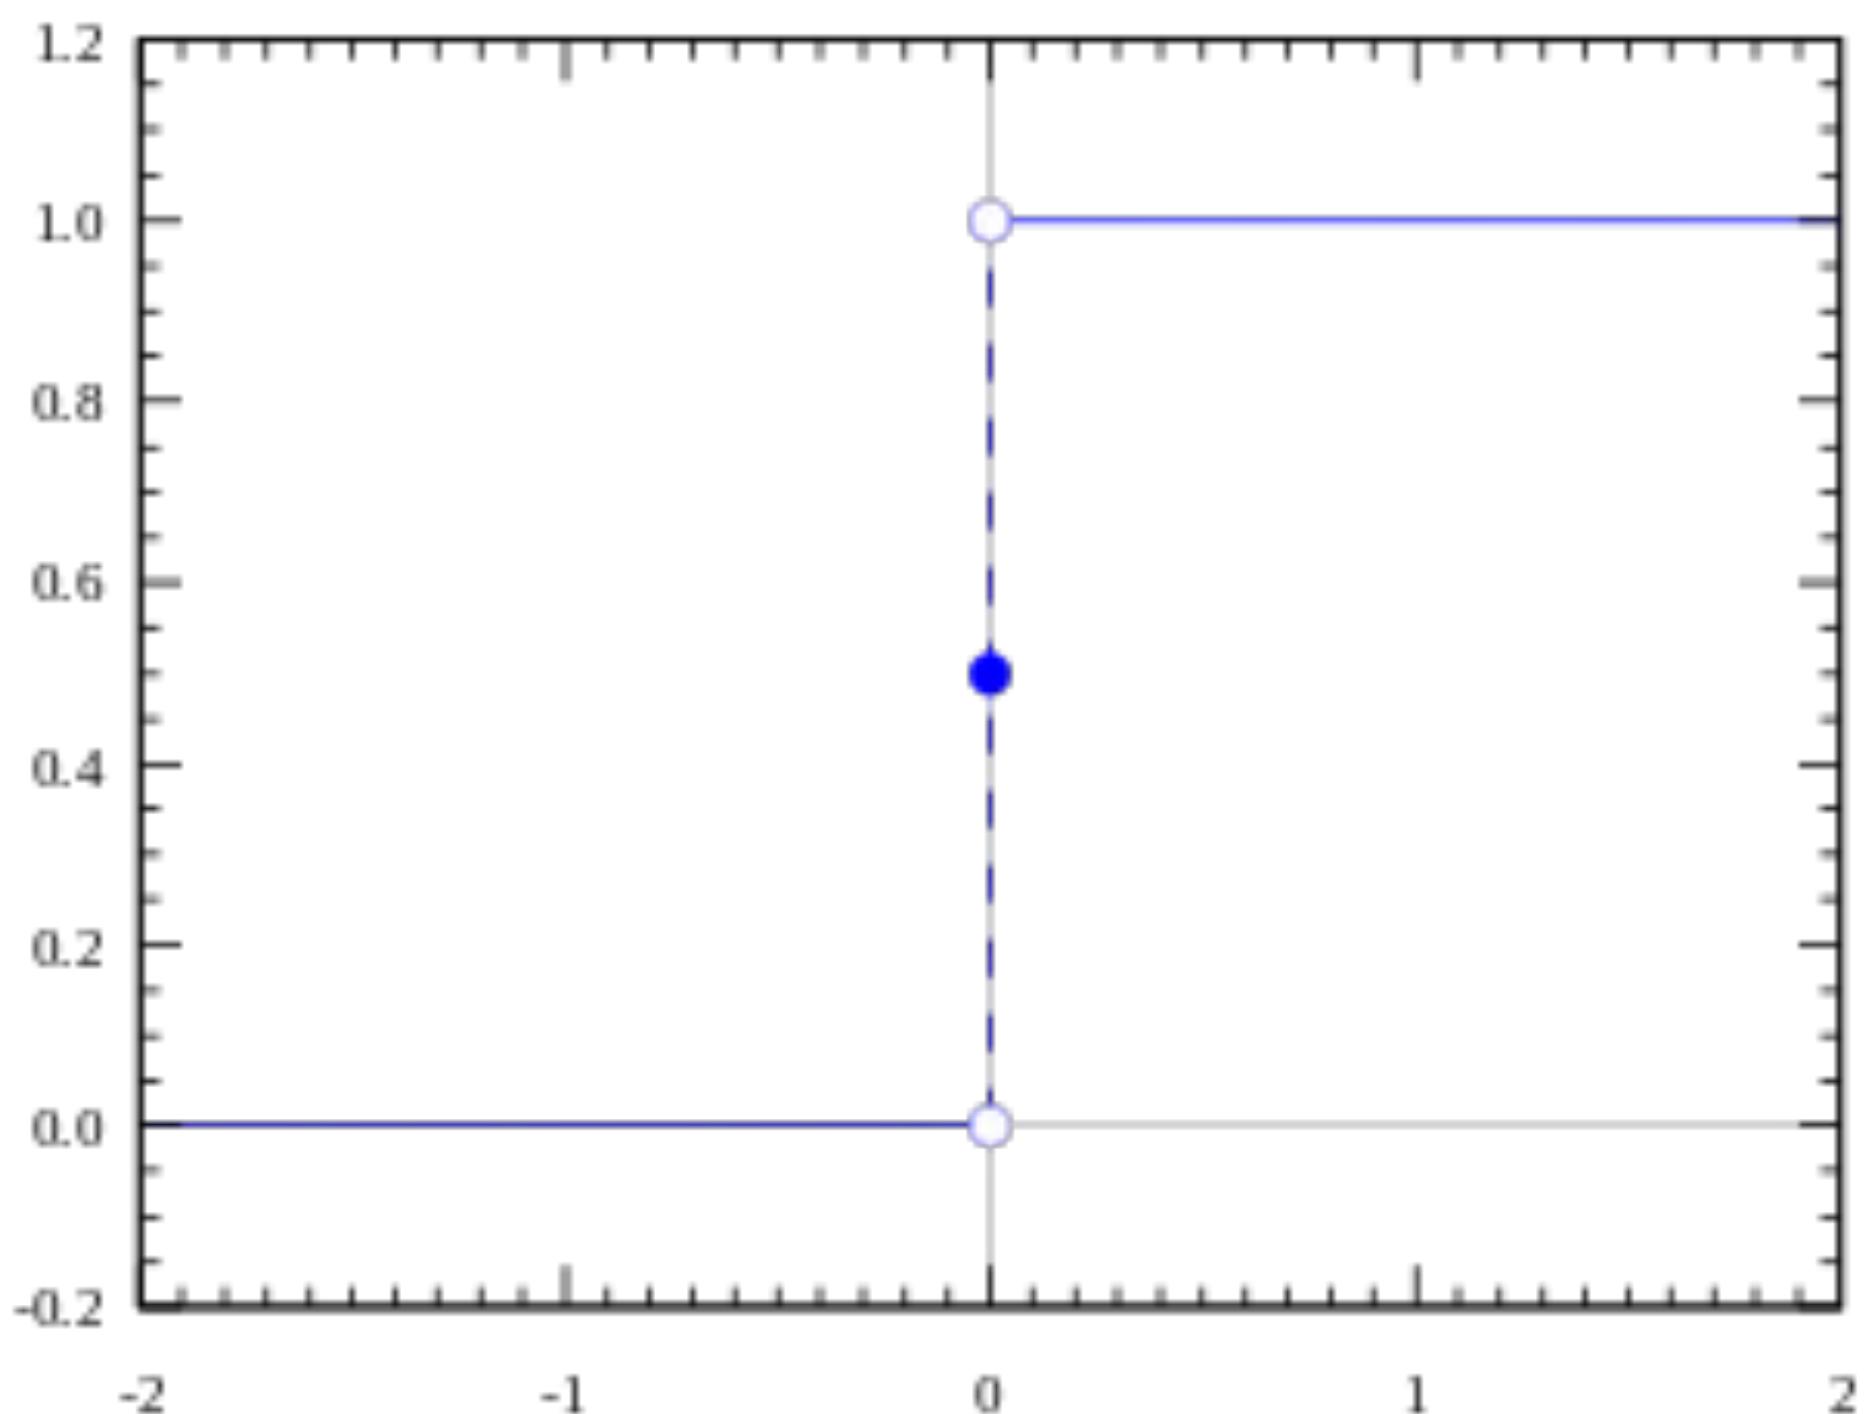


$$f(Y) \in [\min ; \max]$$

Machine learning mechanics

- Activation Functions

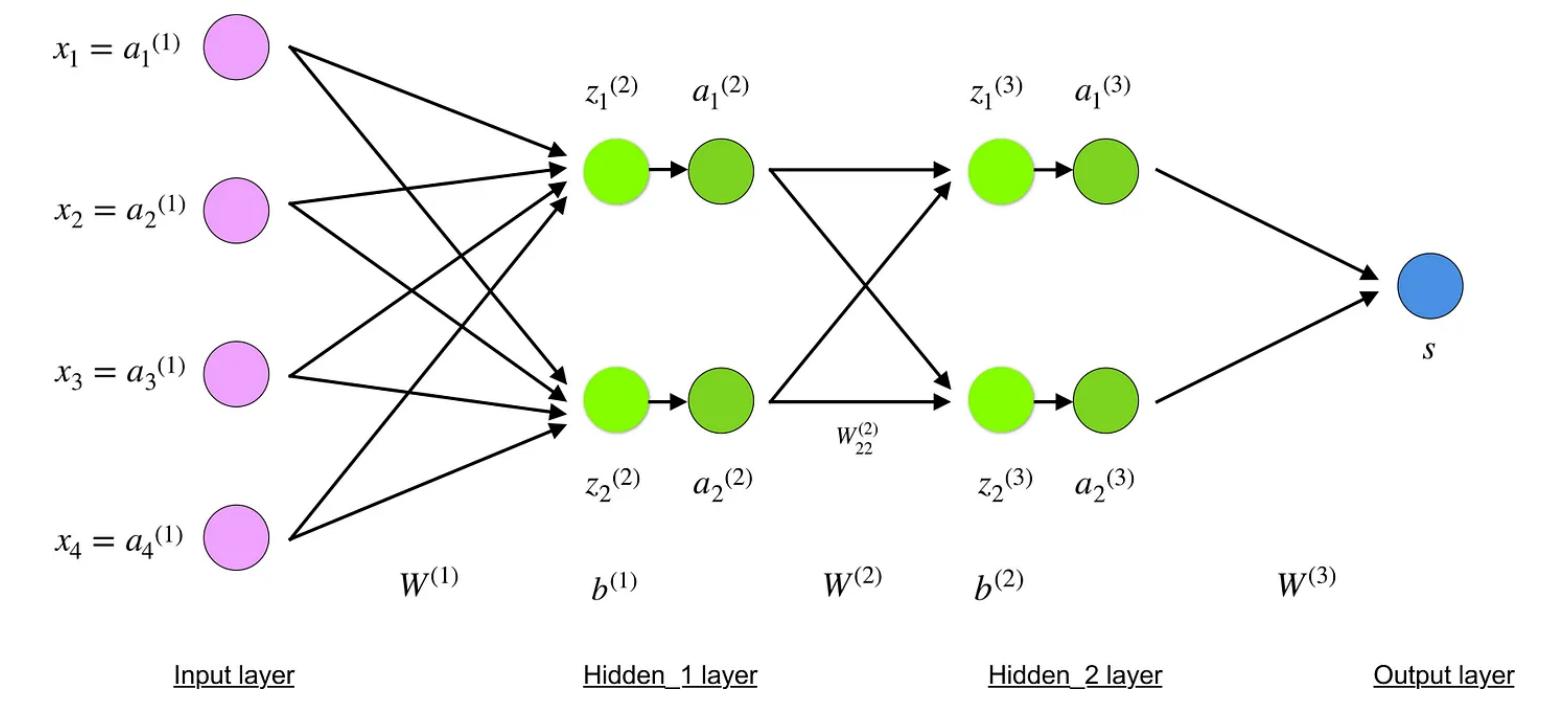
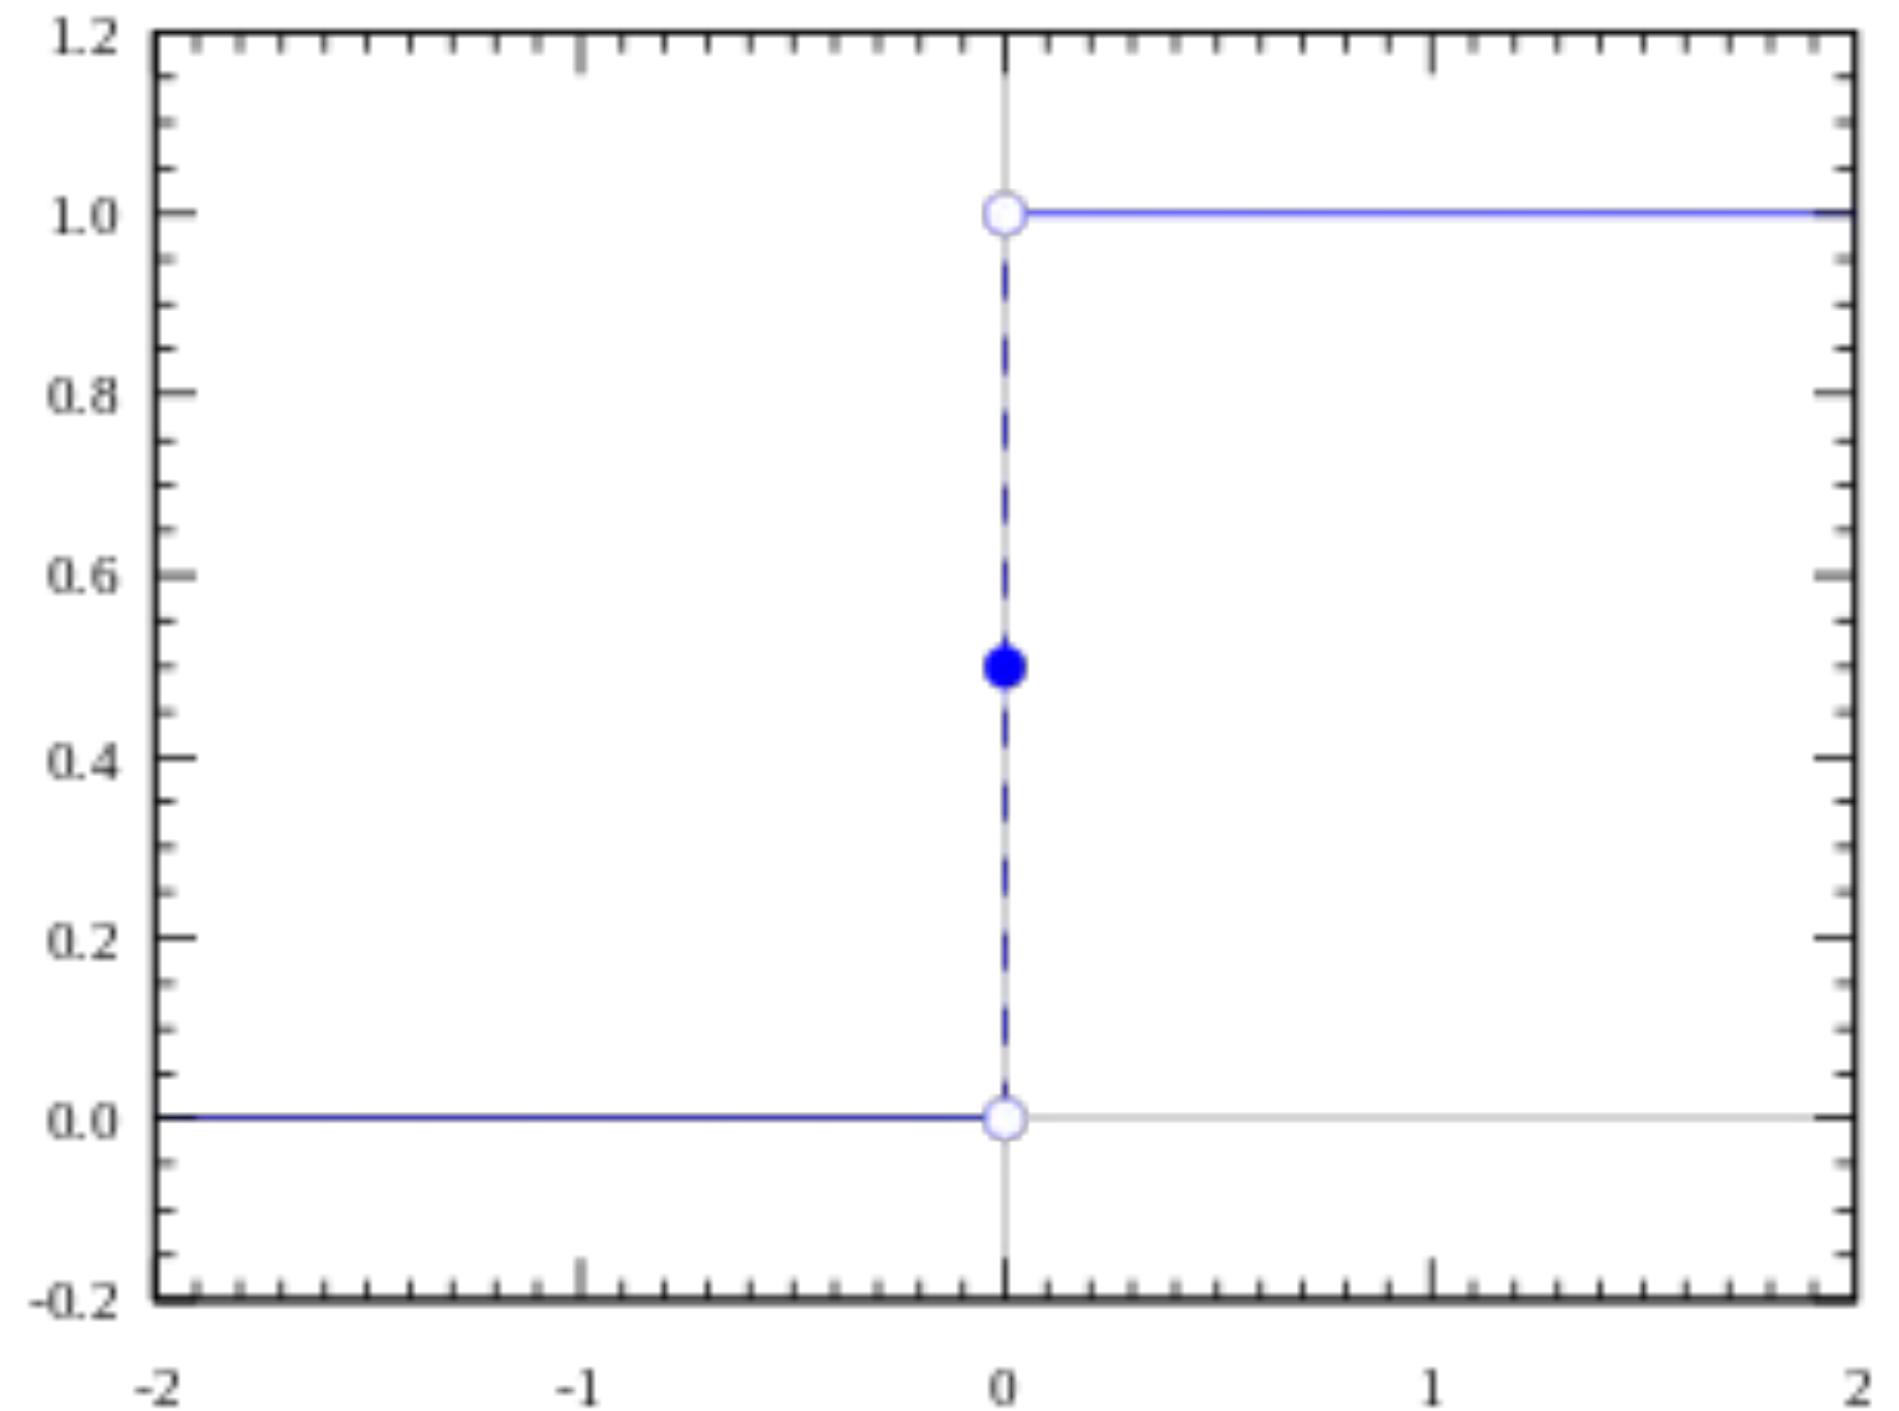
- Step function



Machine learning mechanics

- Activation Functions

Step function



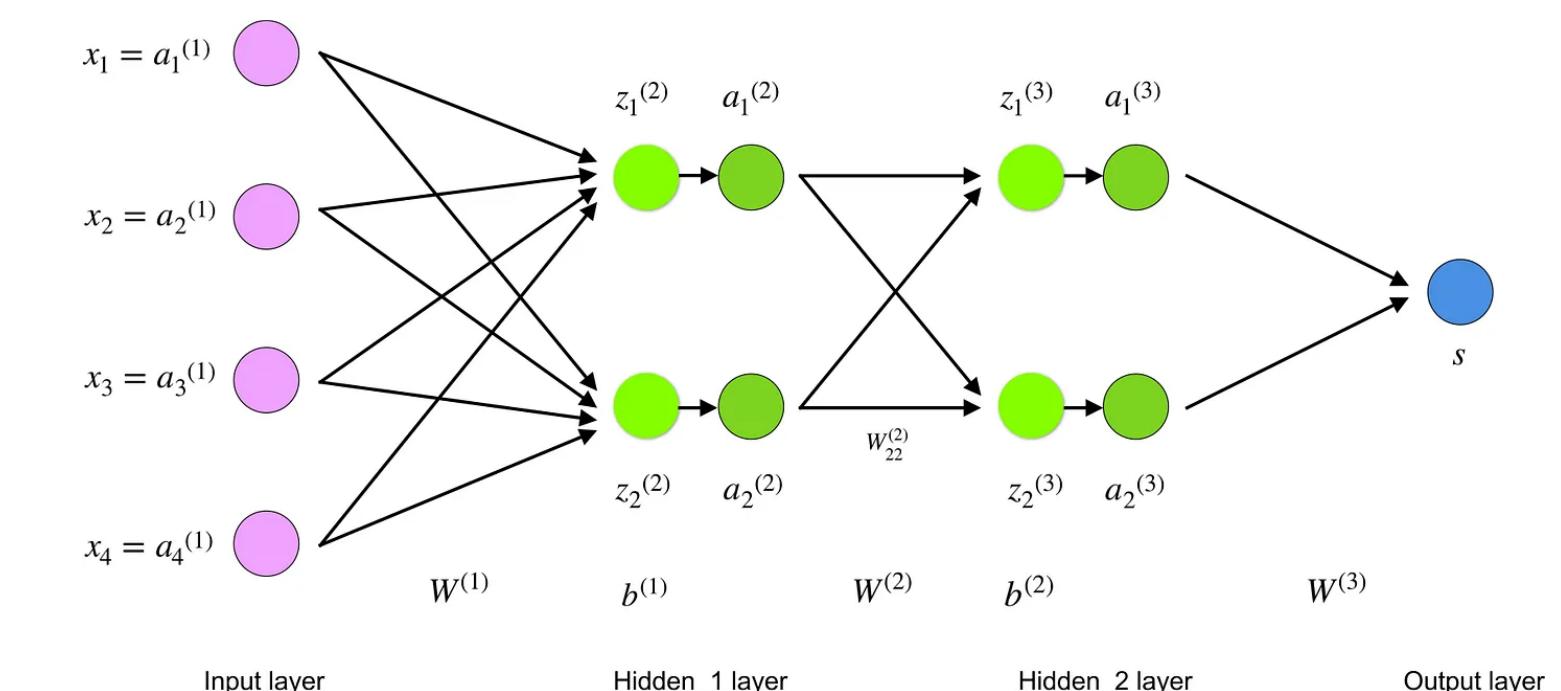
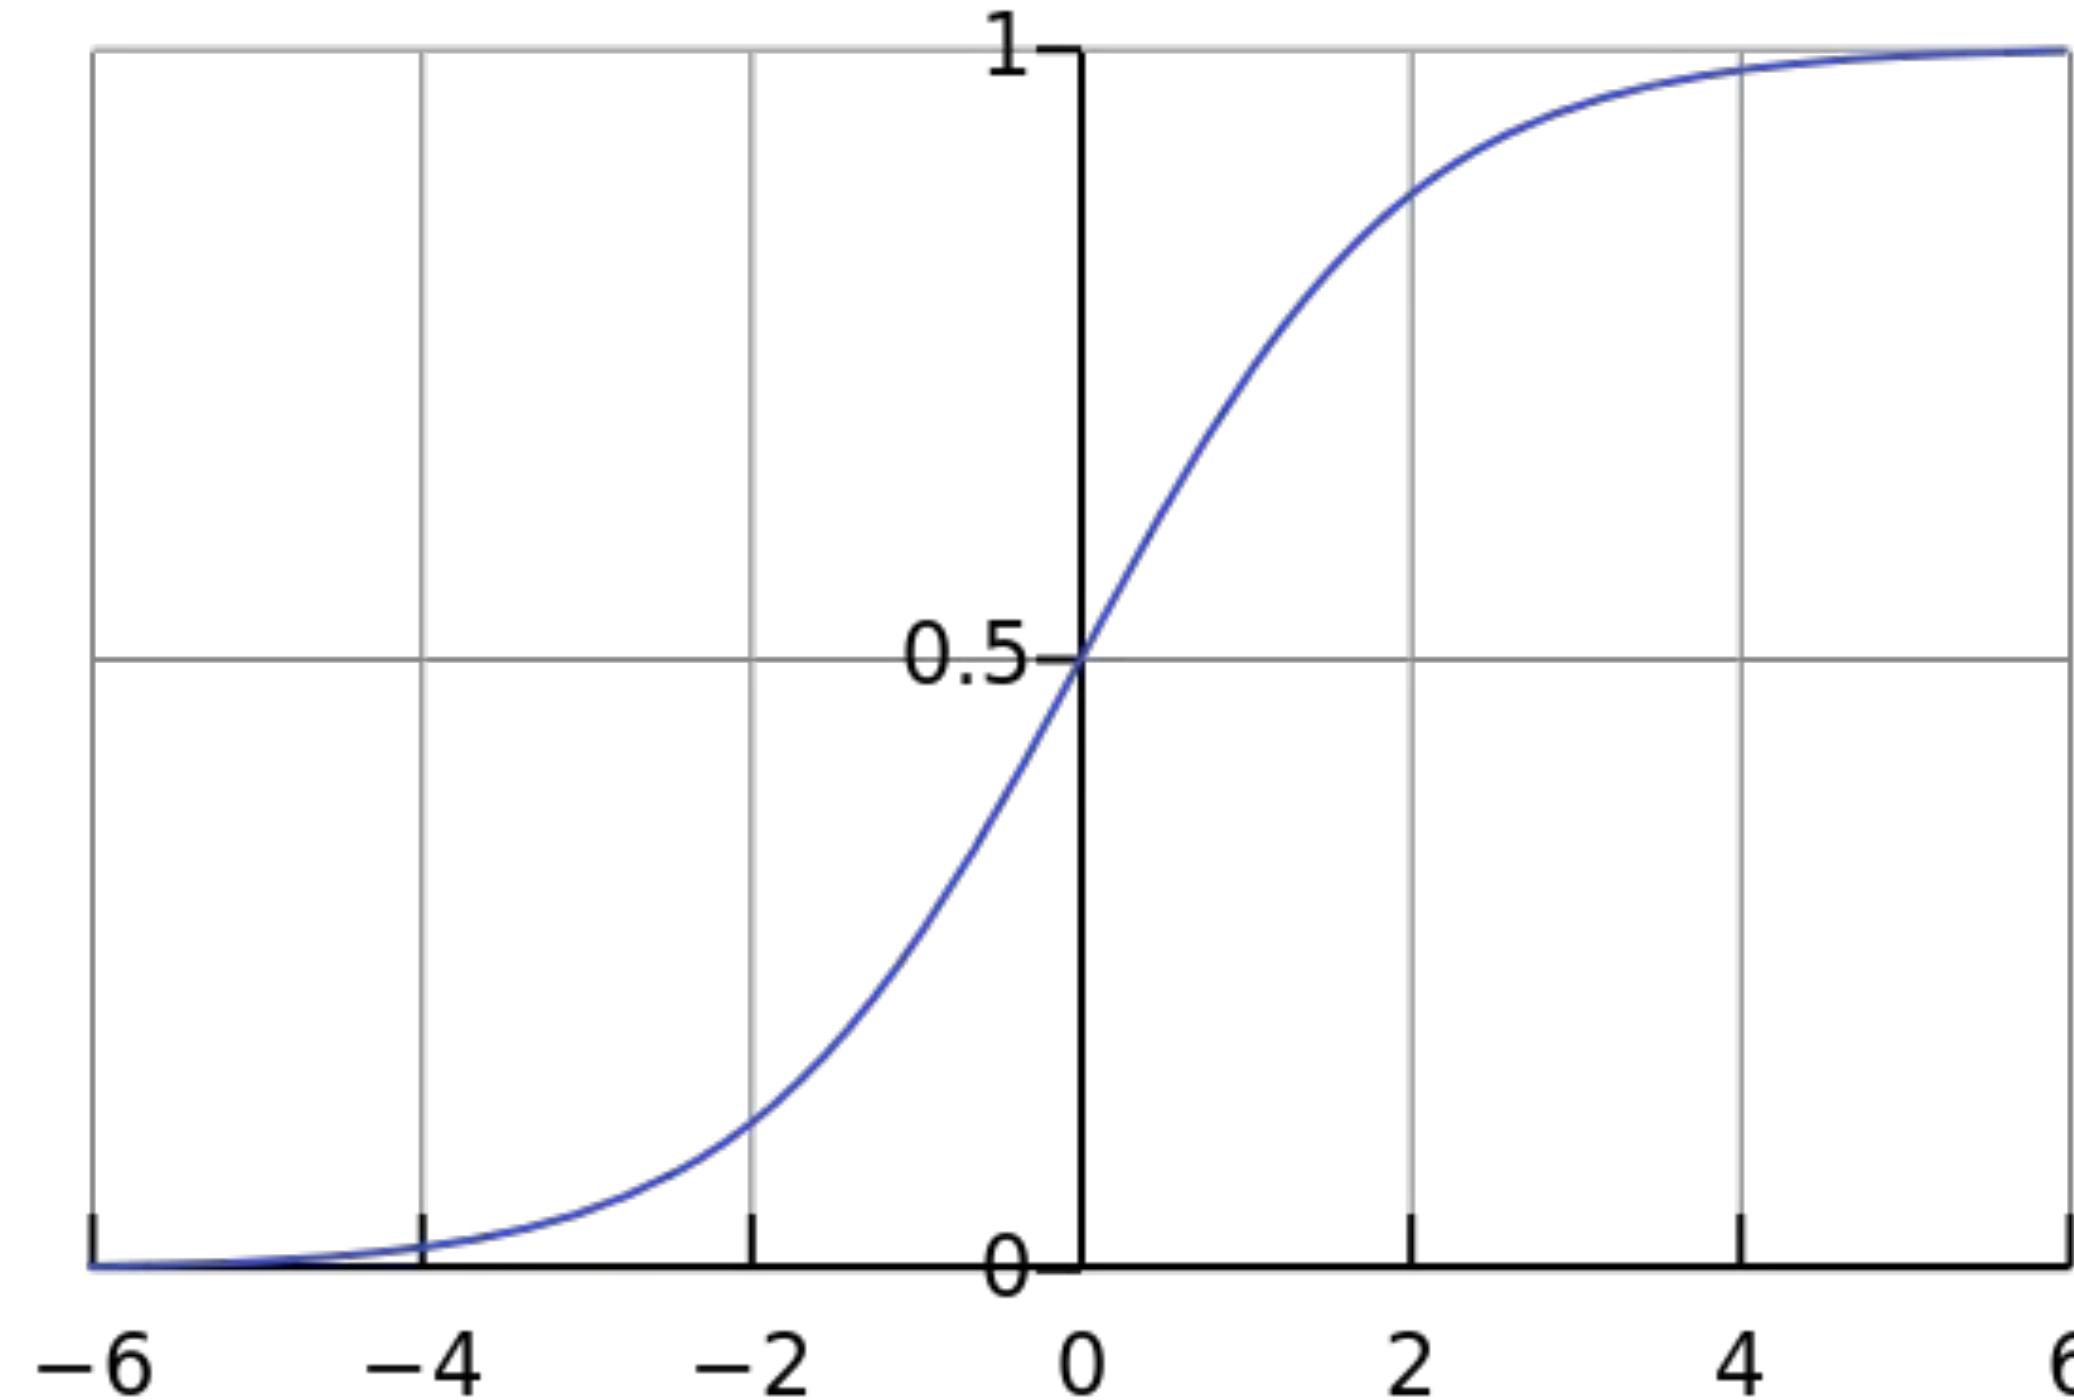
- Pros:
 - Mimics the neuron
 - $Y \in [0; 1]$
- Cons:
 - No smooth gradient
 - Binary

Machine learning mechanics

- Activation Functions

Sigmoid function

$$A = \frac{1}{1+e^{-x}}$$

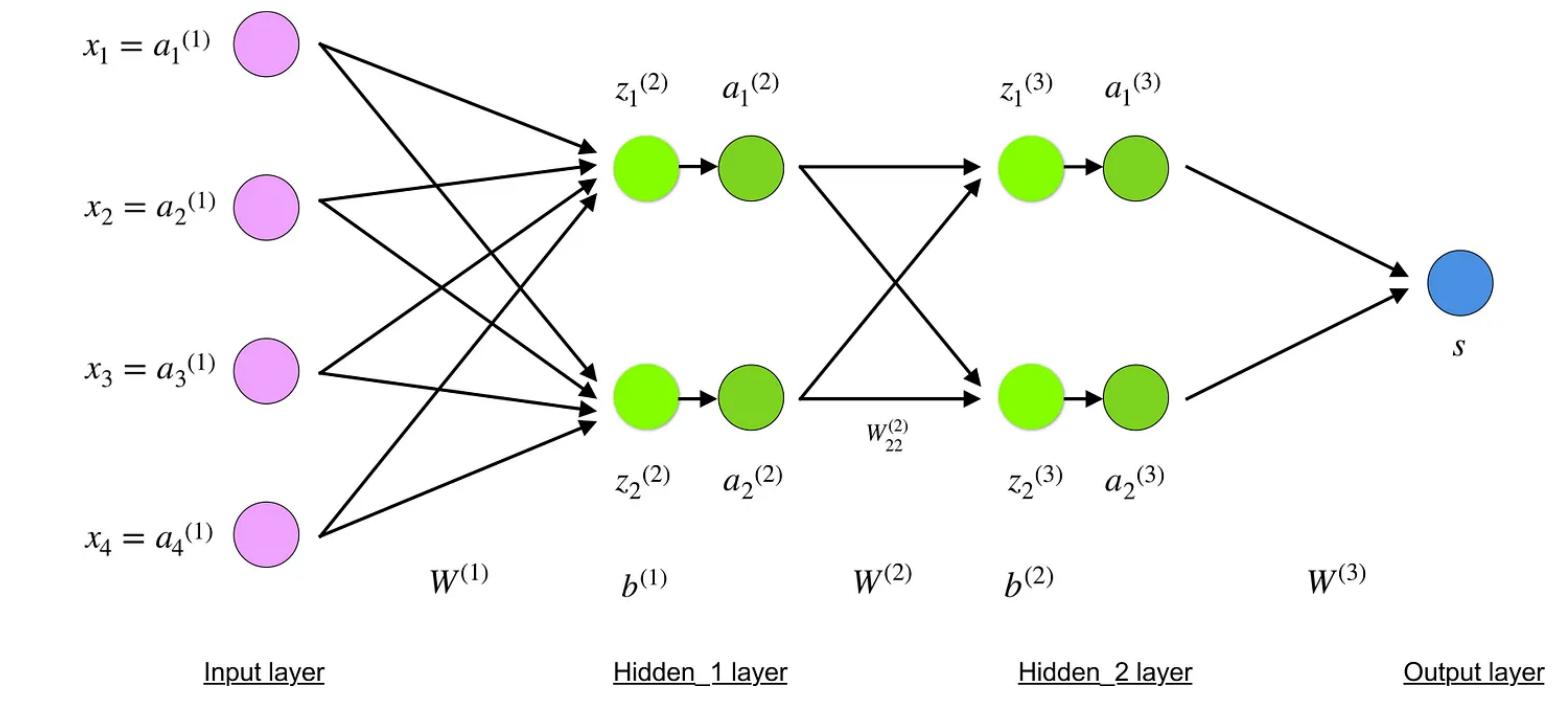
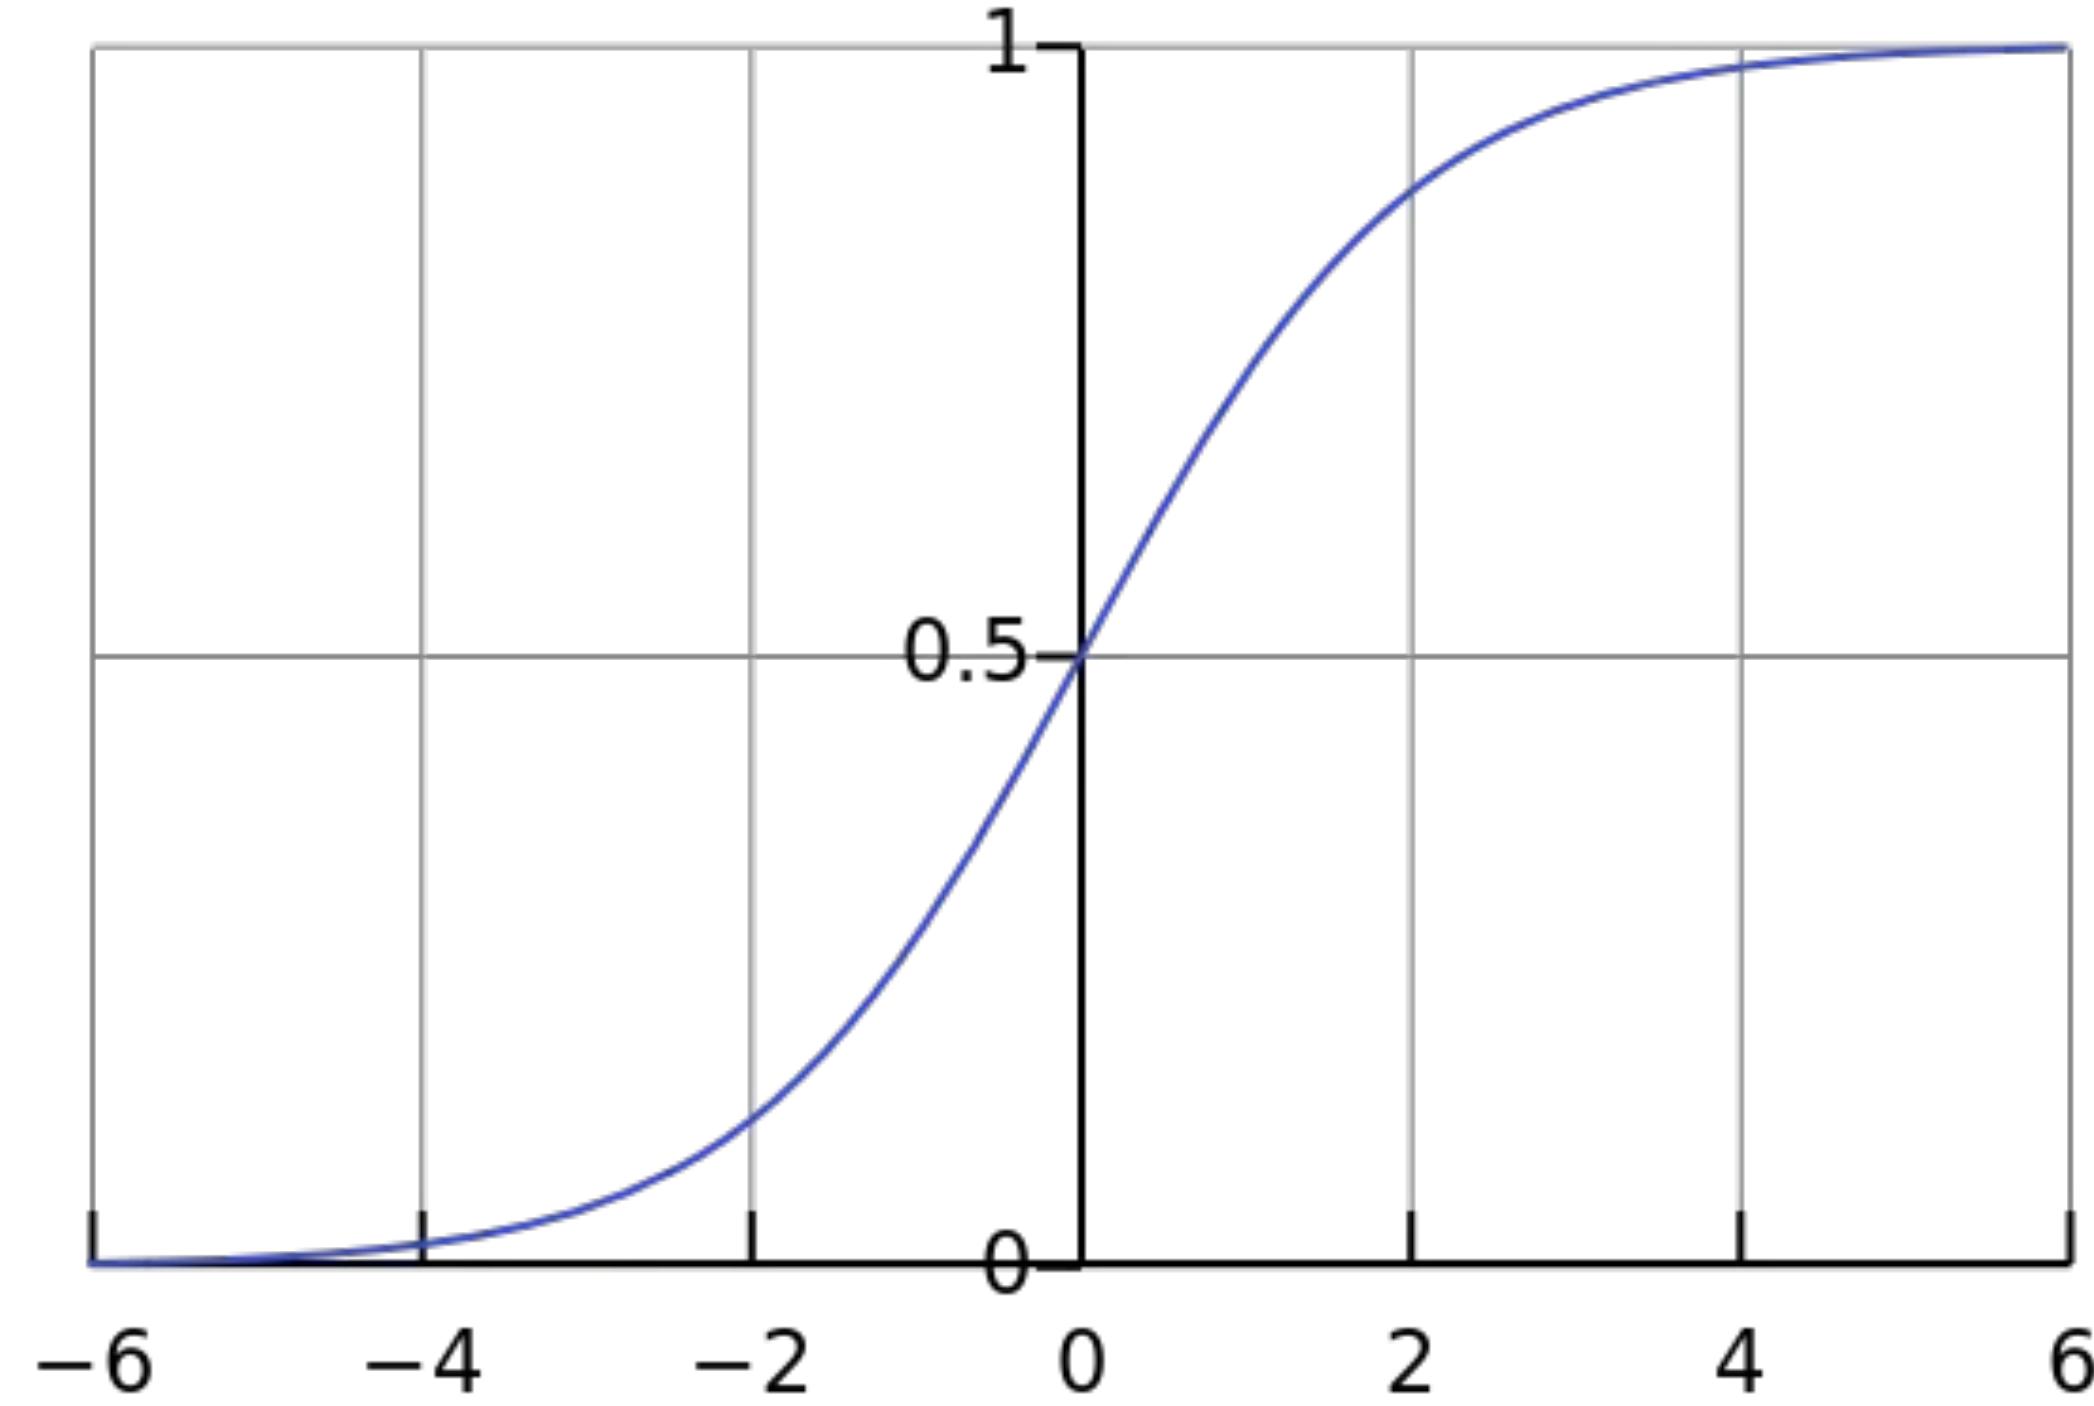


Machine learning mechanics

- Activation Functions

Sigmoid function

$$A = \frac{1}{1+e^{-x}}$$



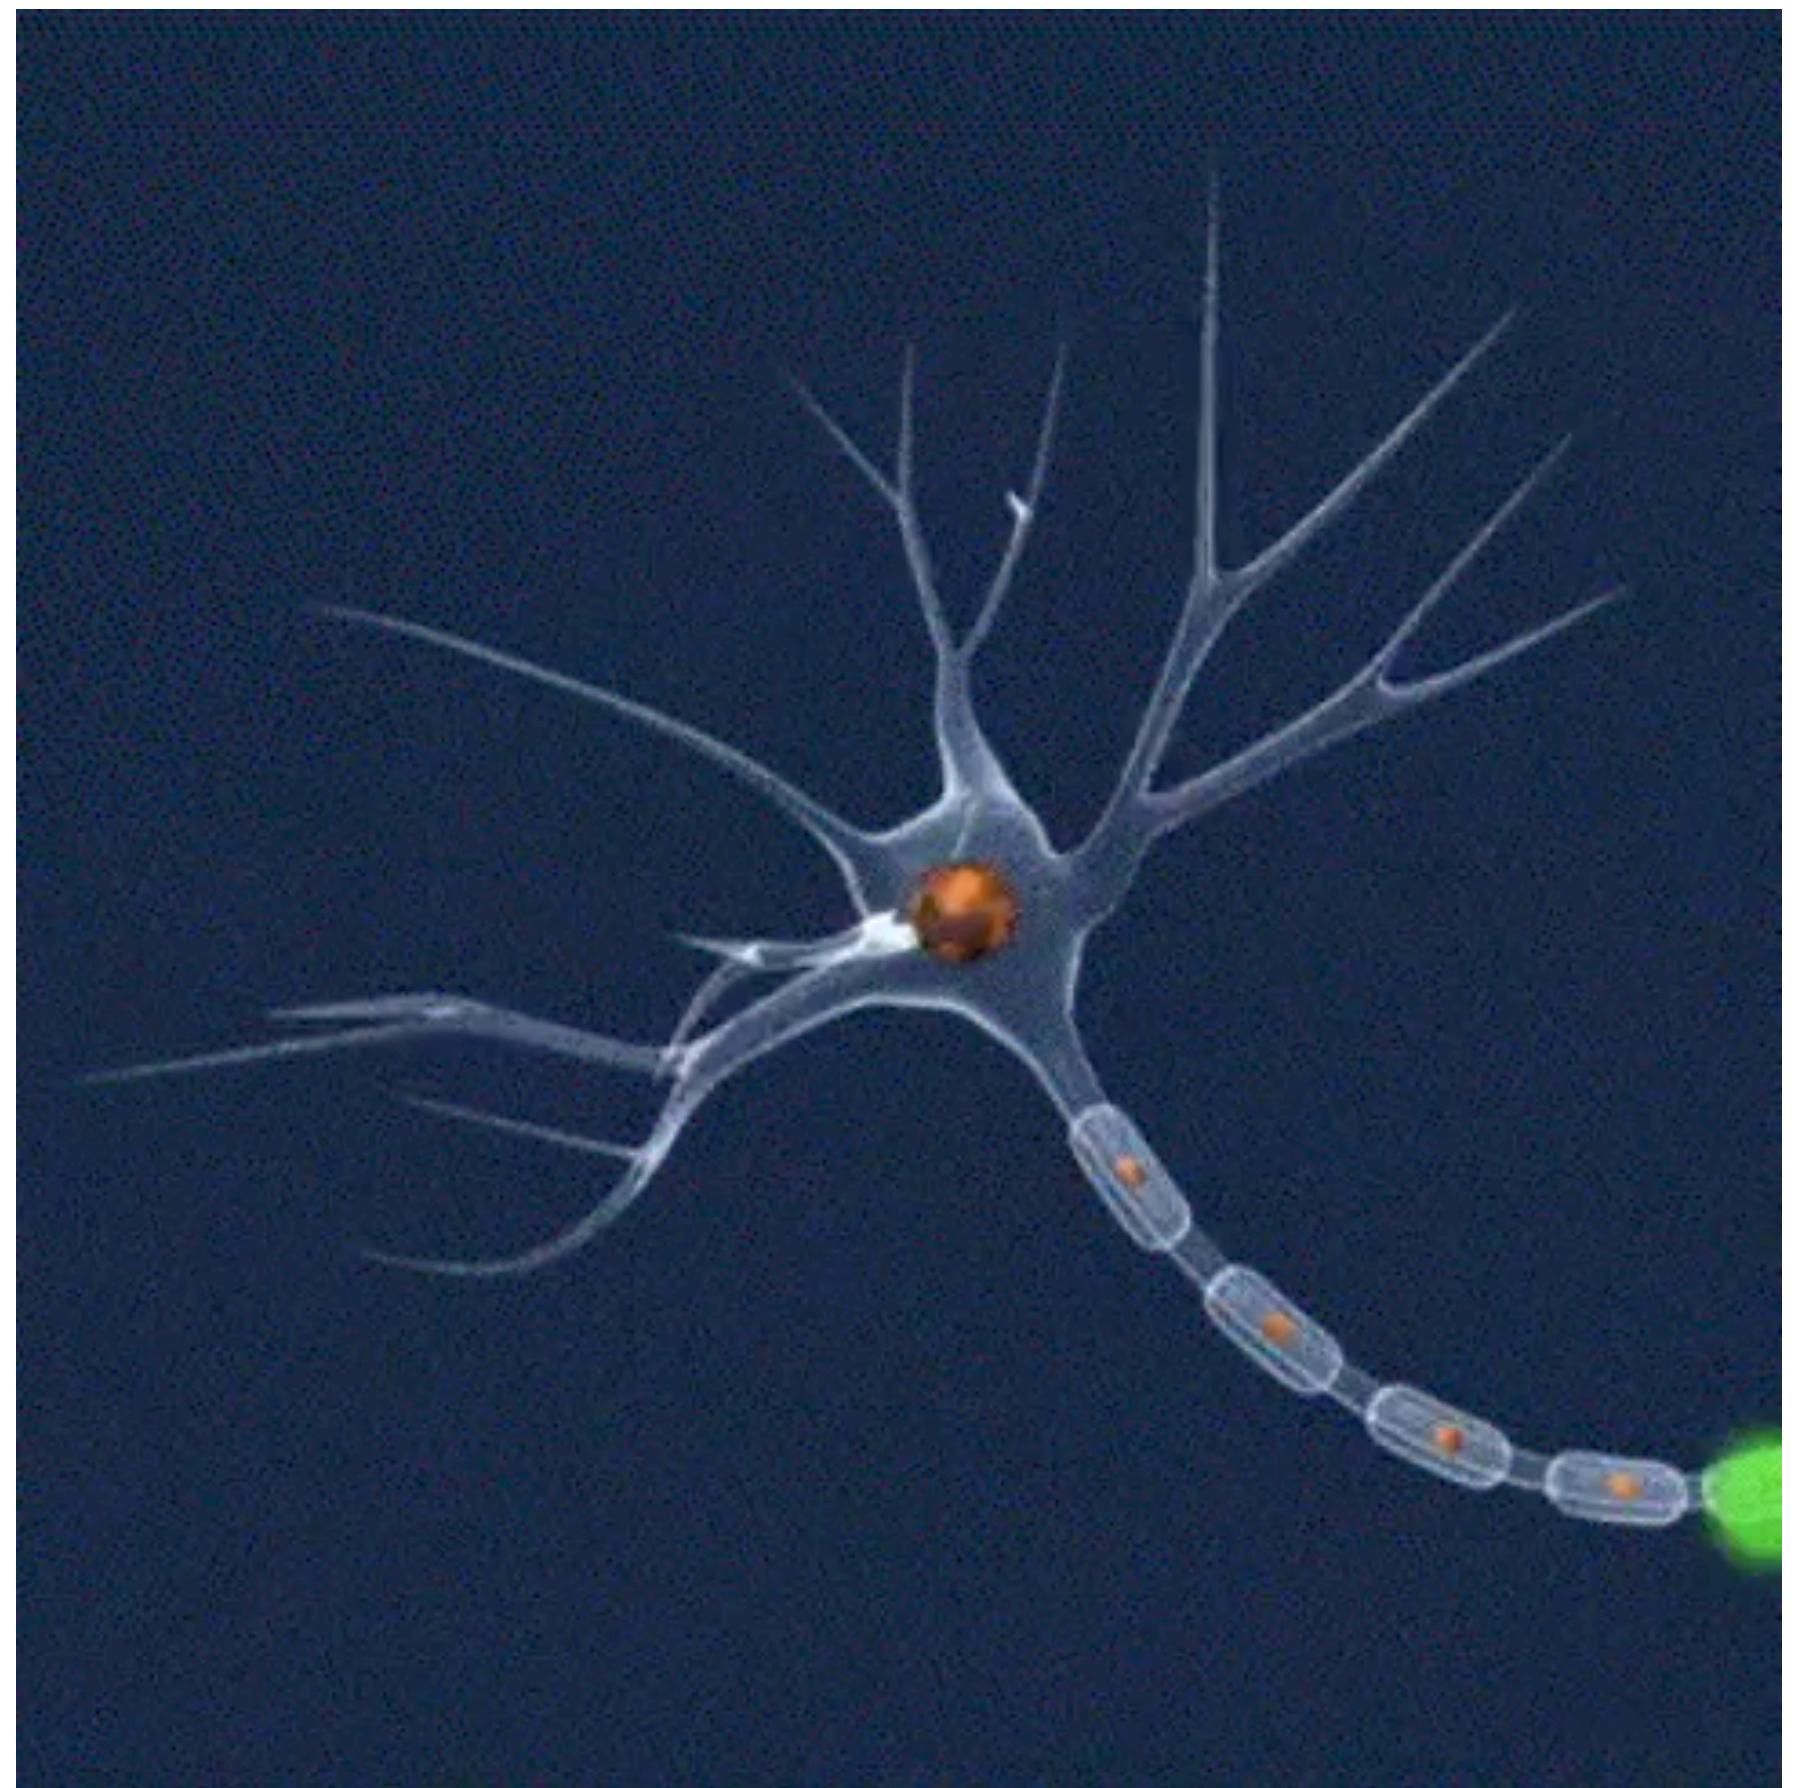
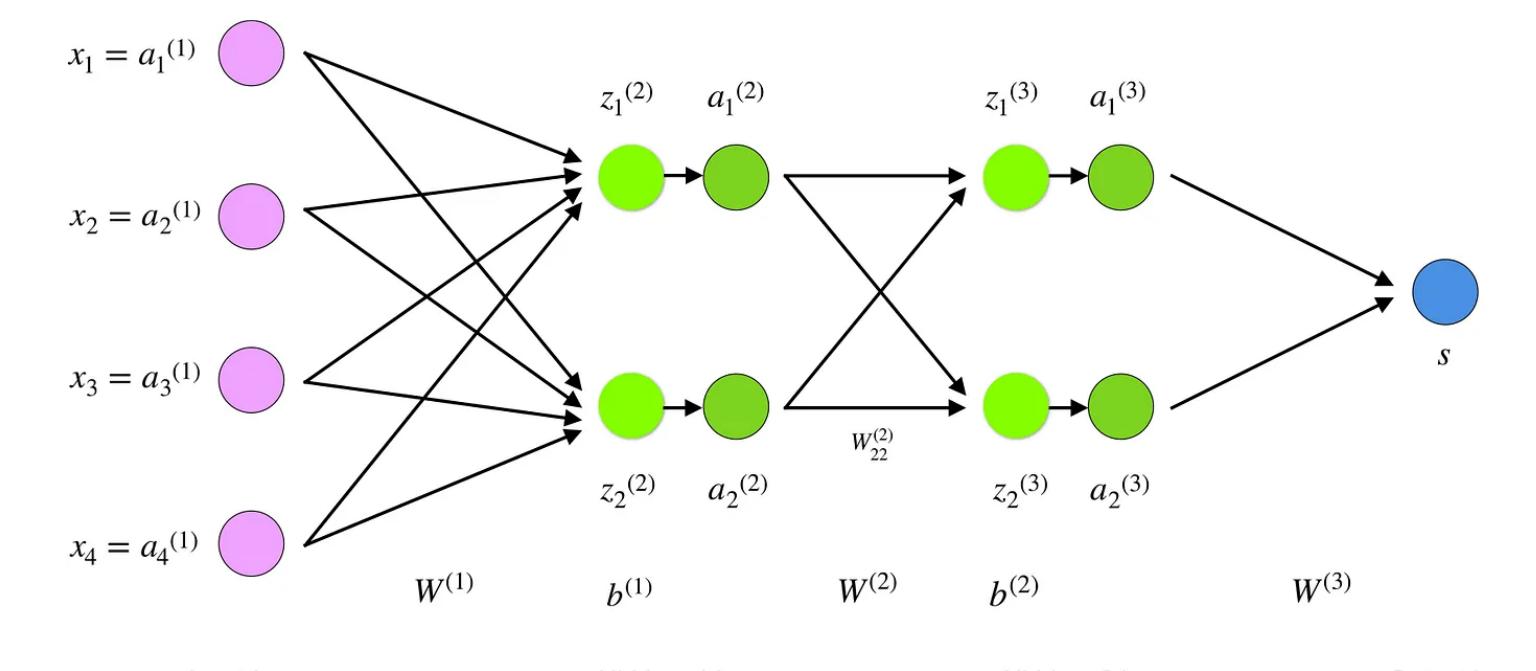
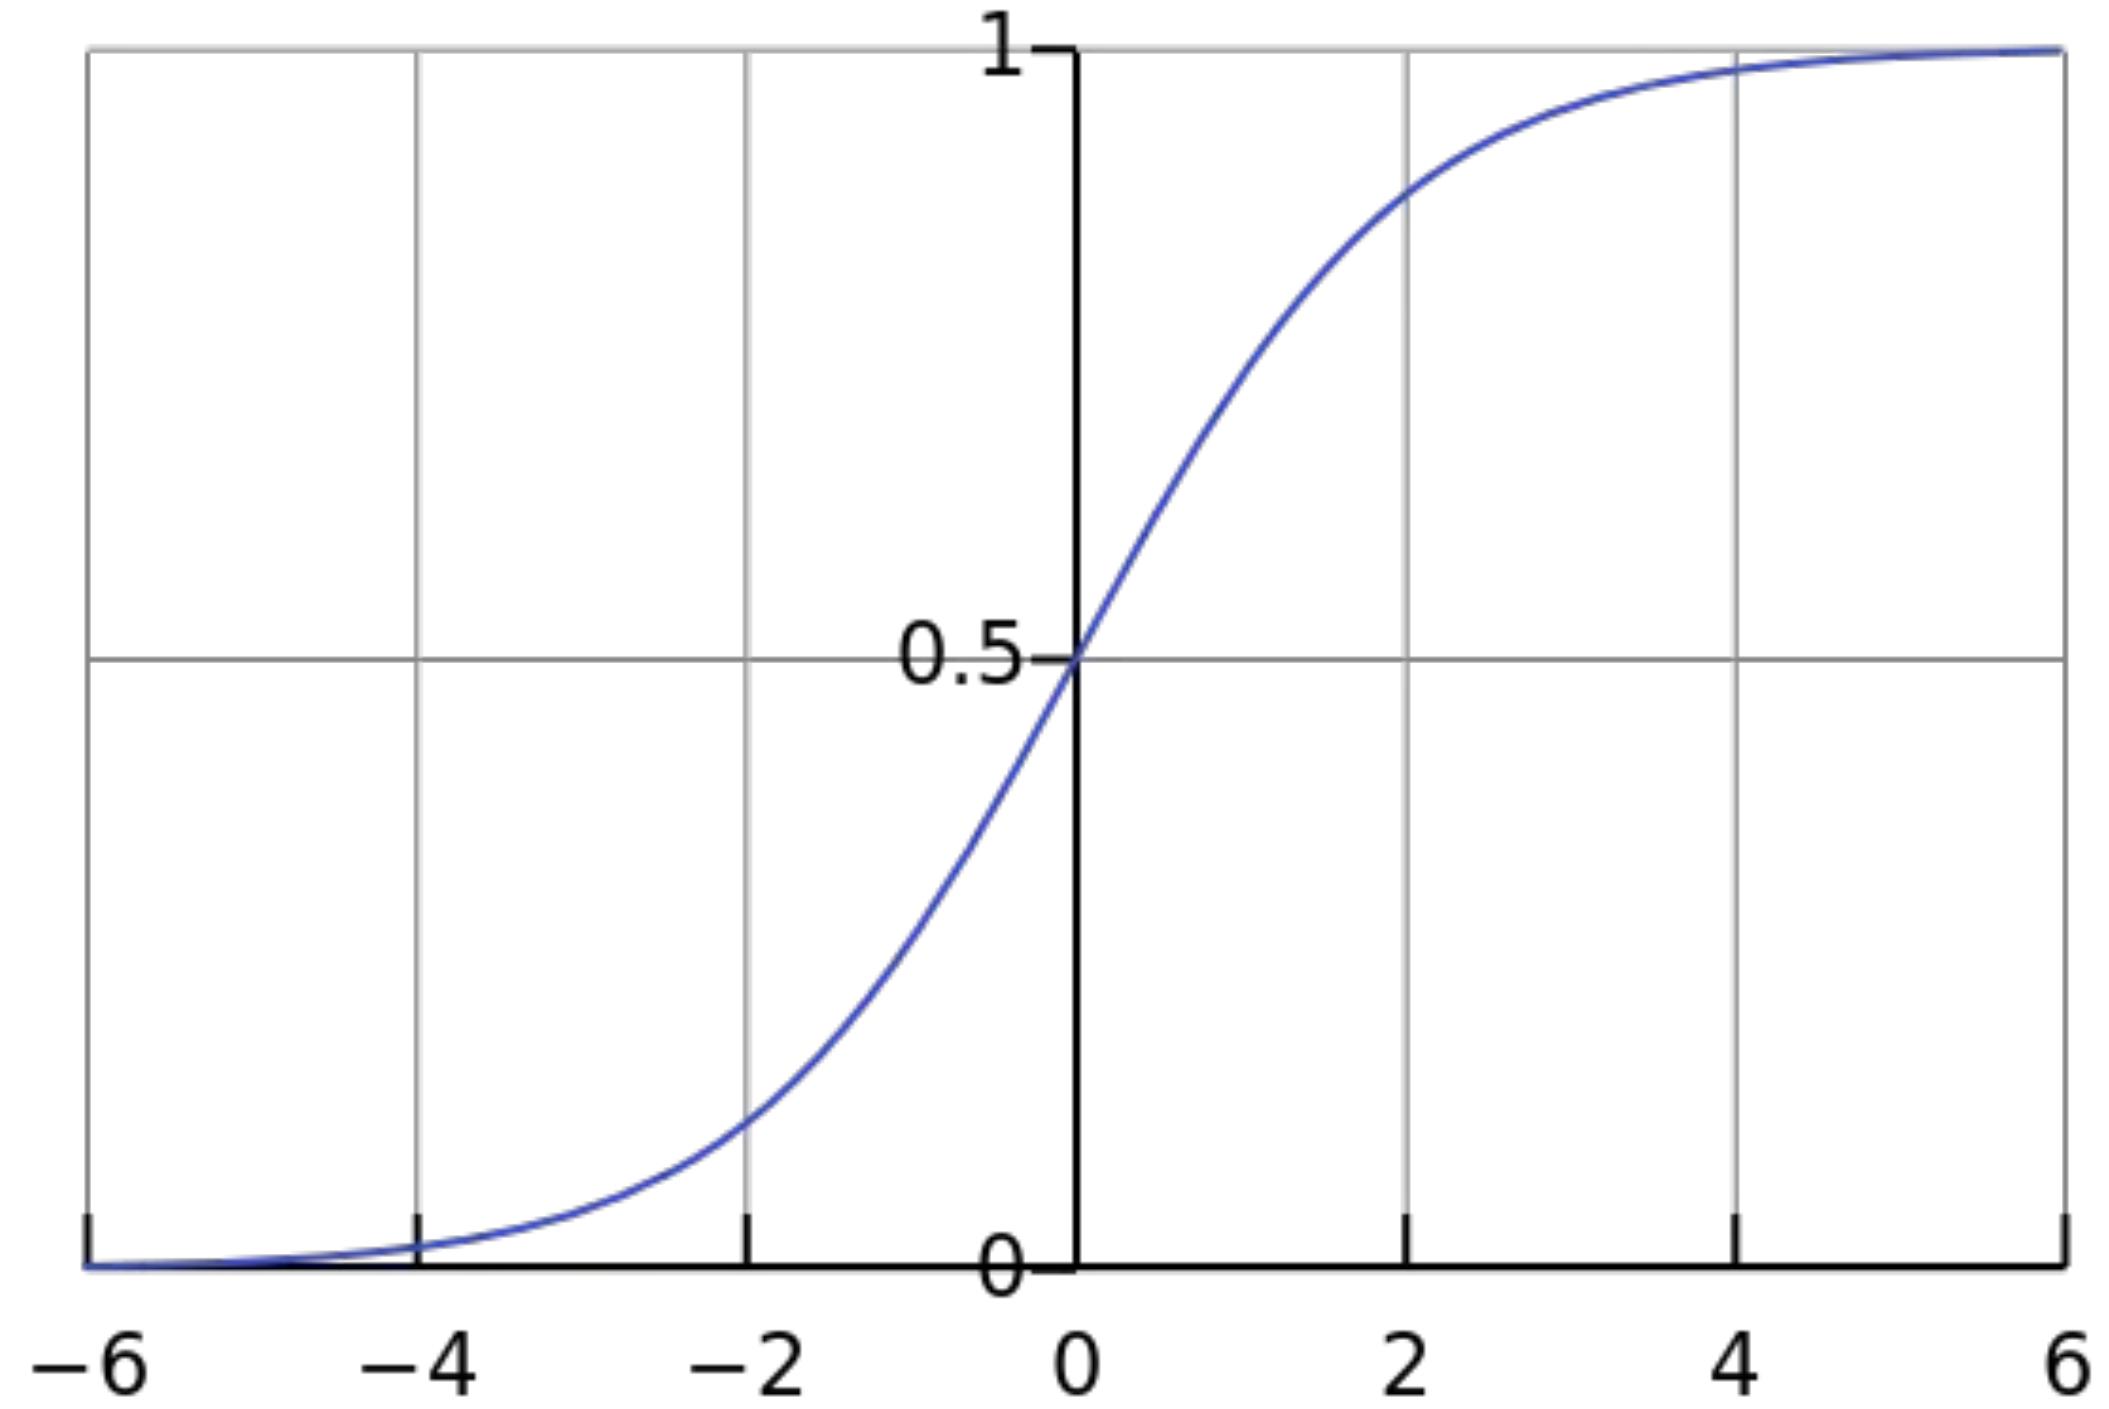
- Pros:
 - Smooth step
 - Smooth gradient
 - Not Binary
 - $Y \in [0; 1]$
- Cons:
 - Small gradient at the top and bottom
 - Computationally intensive

Machine learning mechanics

- Activation Functions

Sigmoid function

$$A = \frac{1}{1+e^{-x}}$$

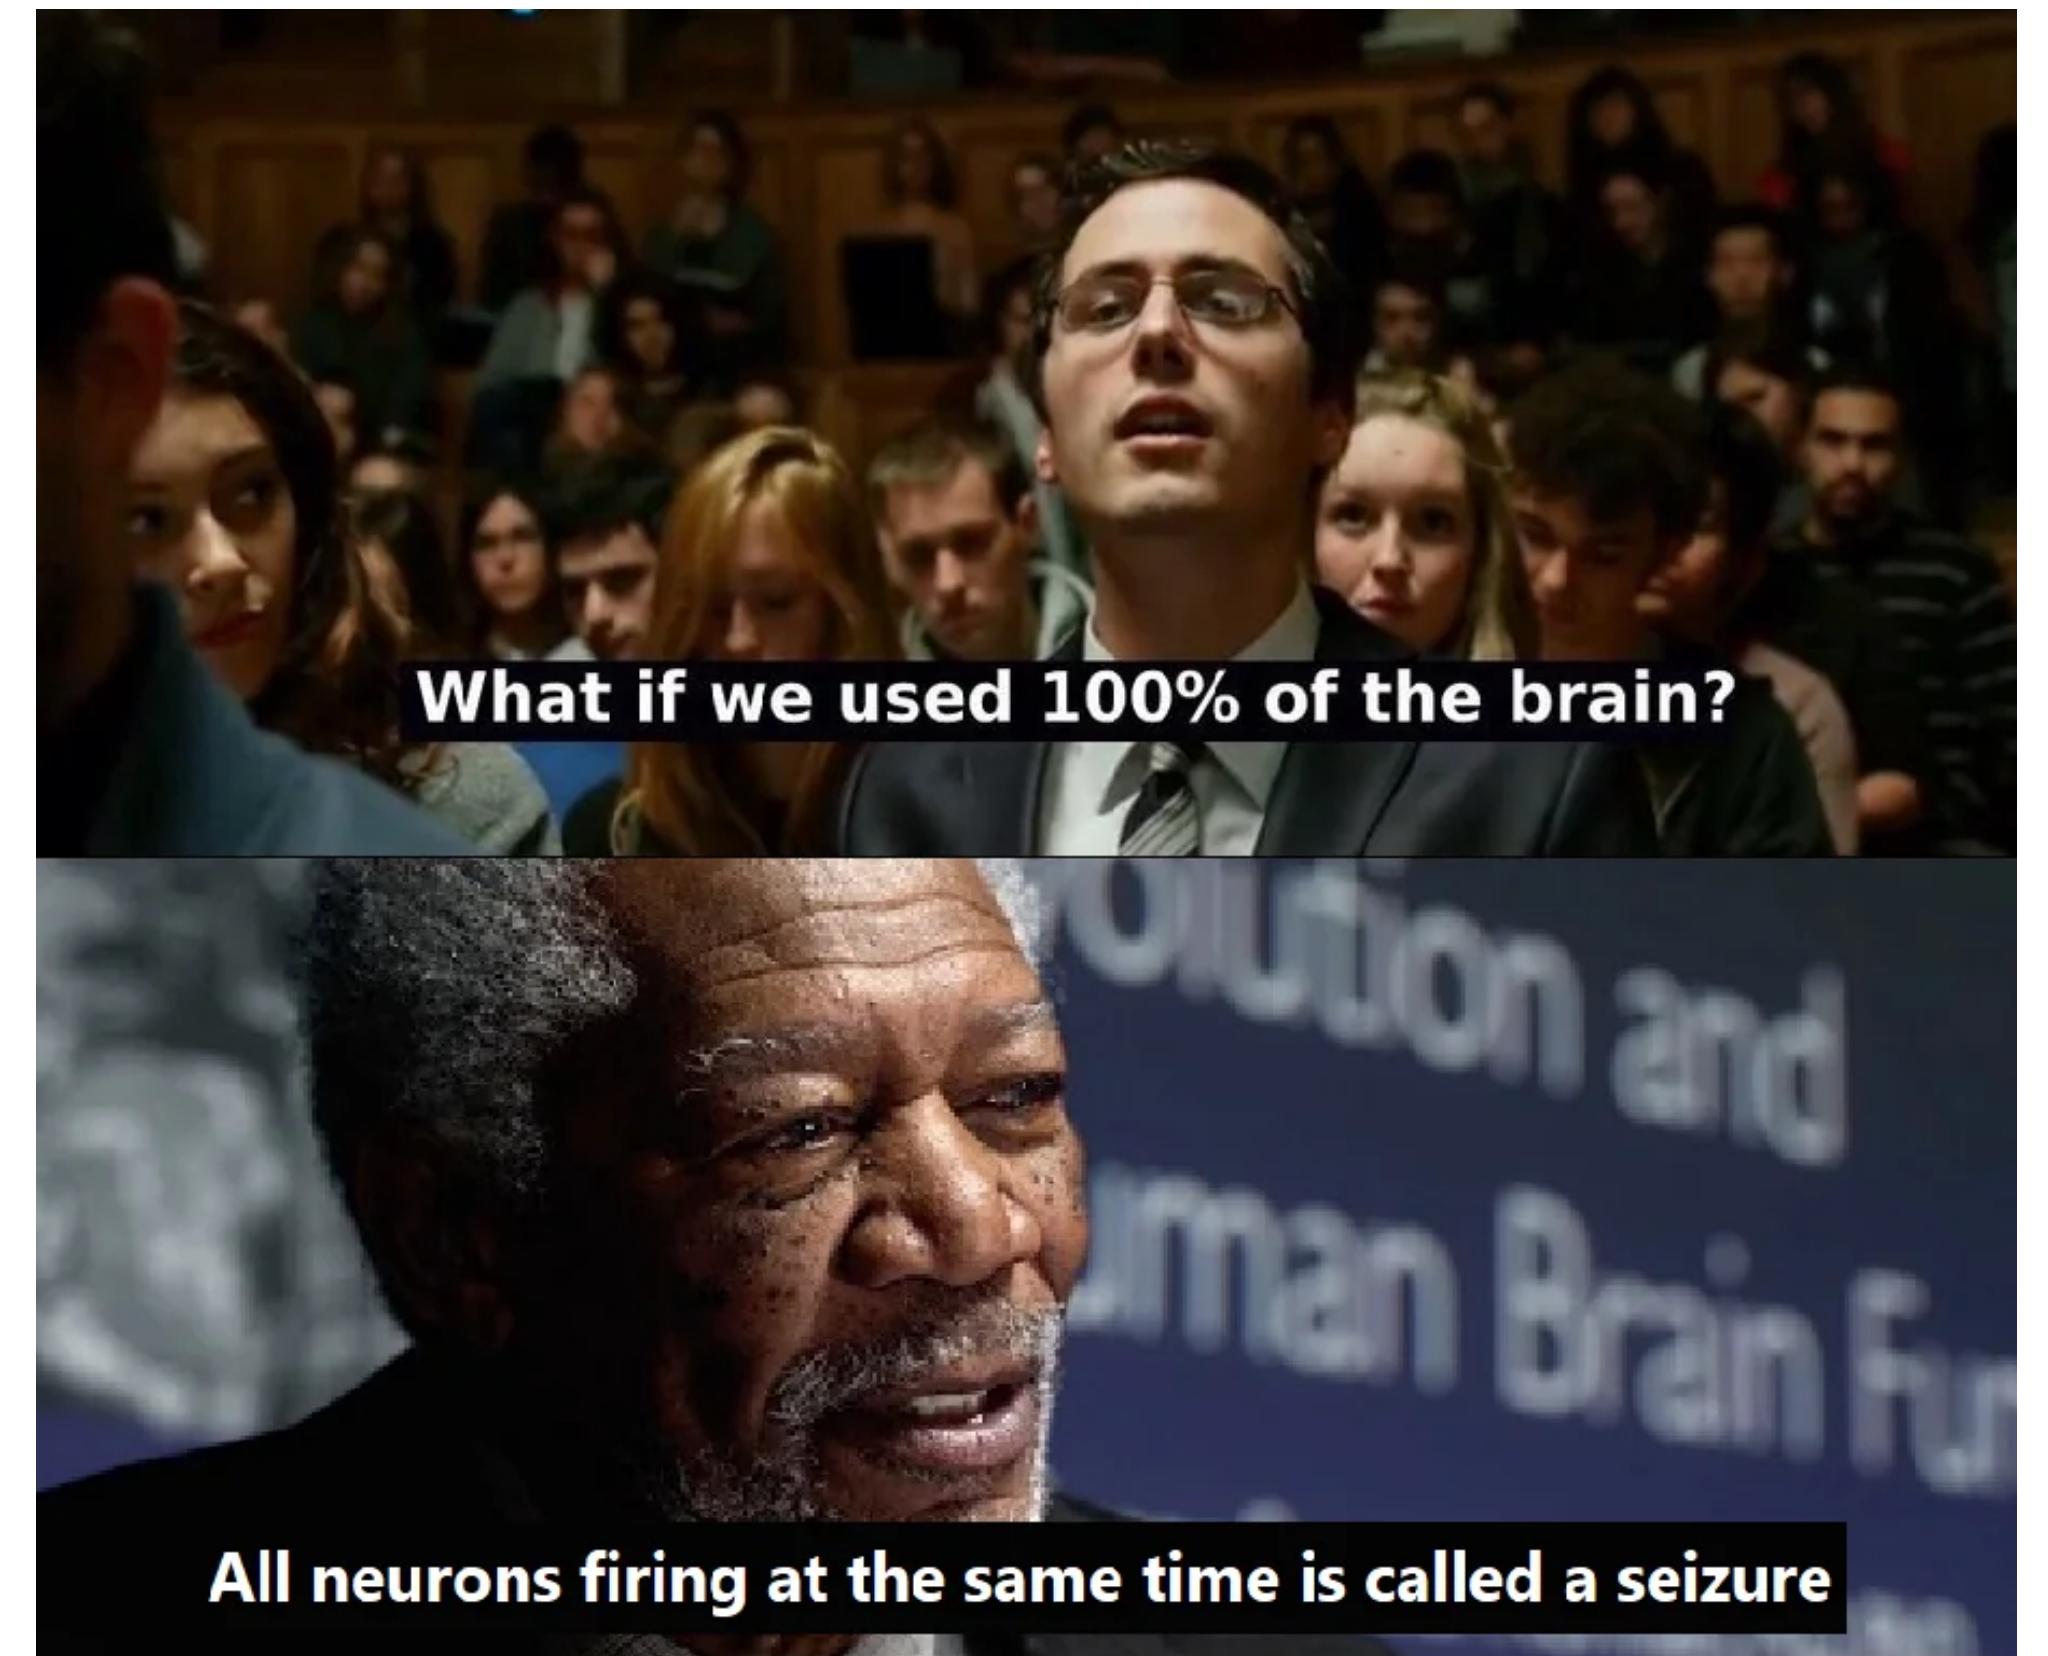
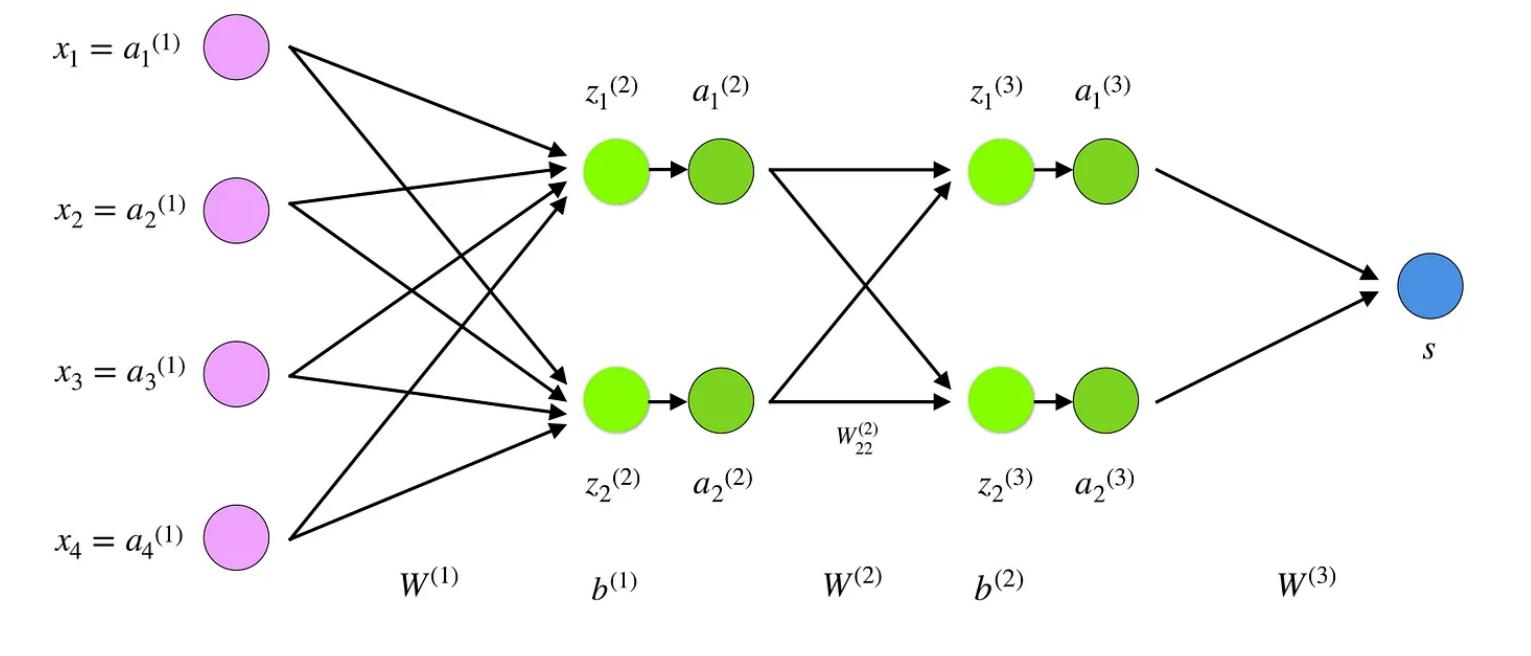
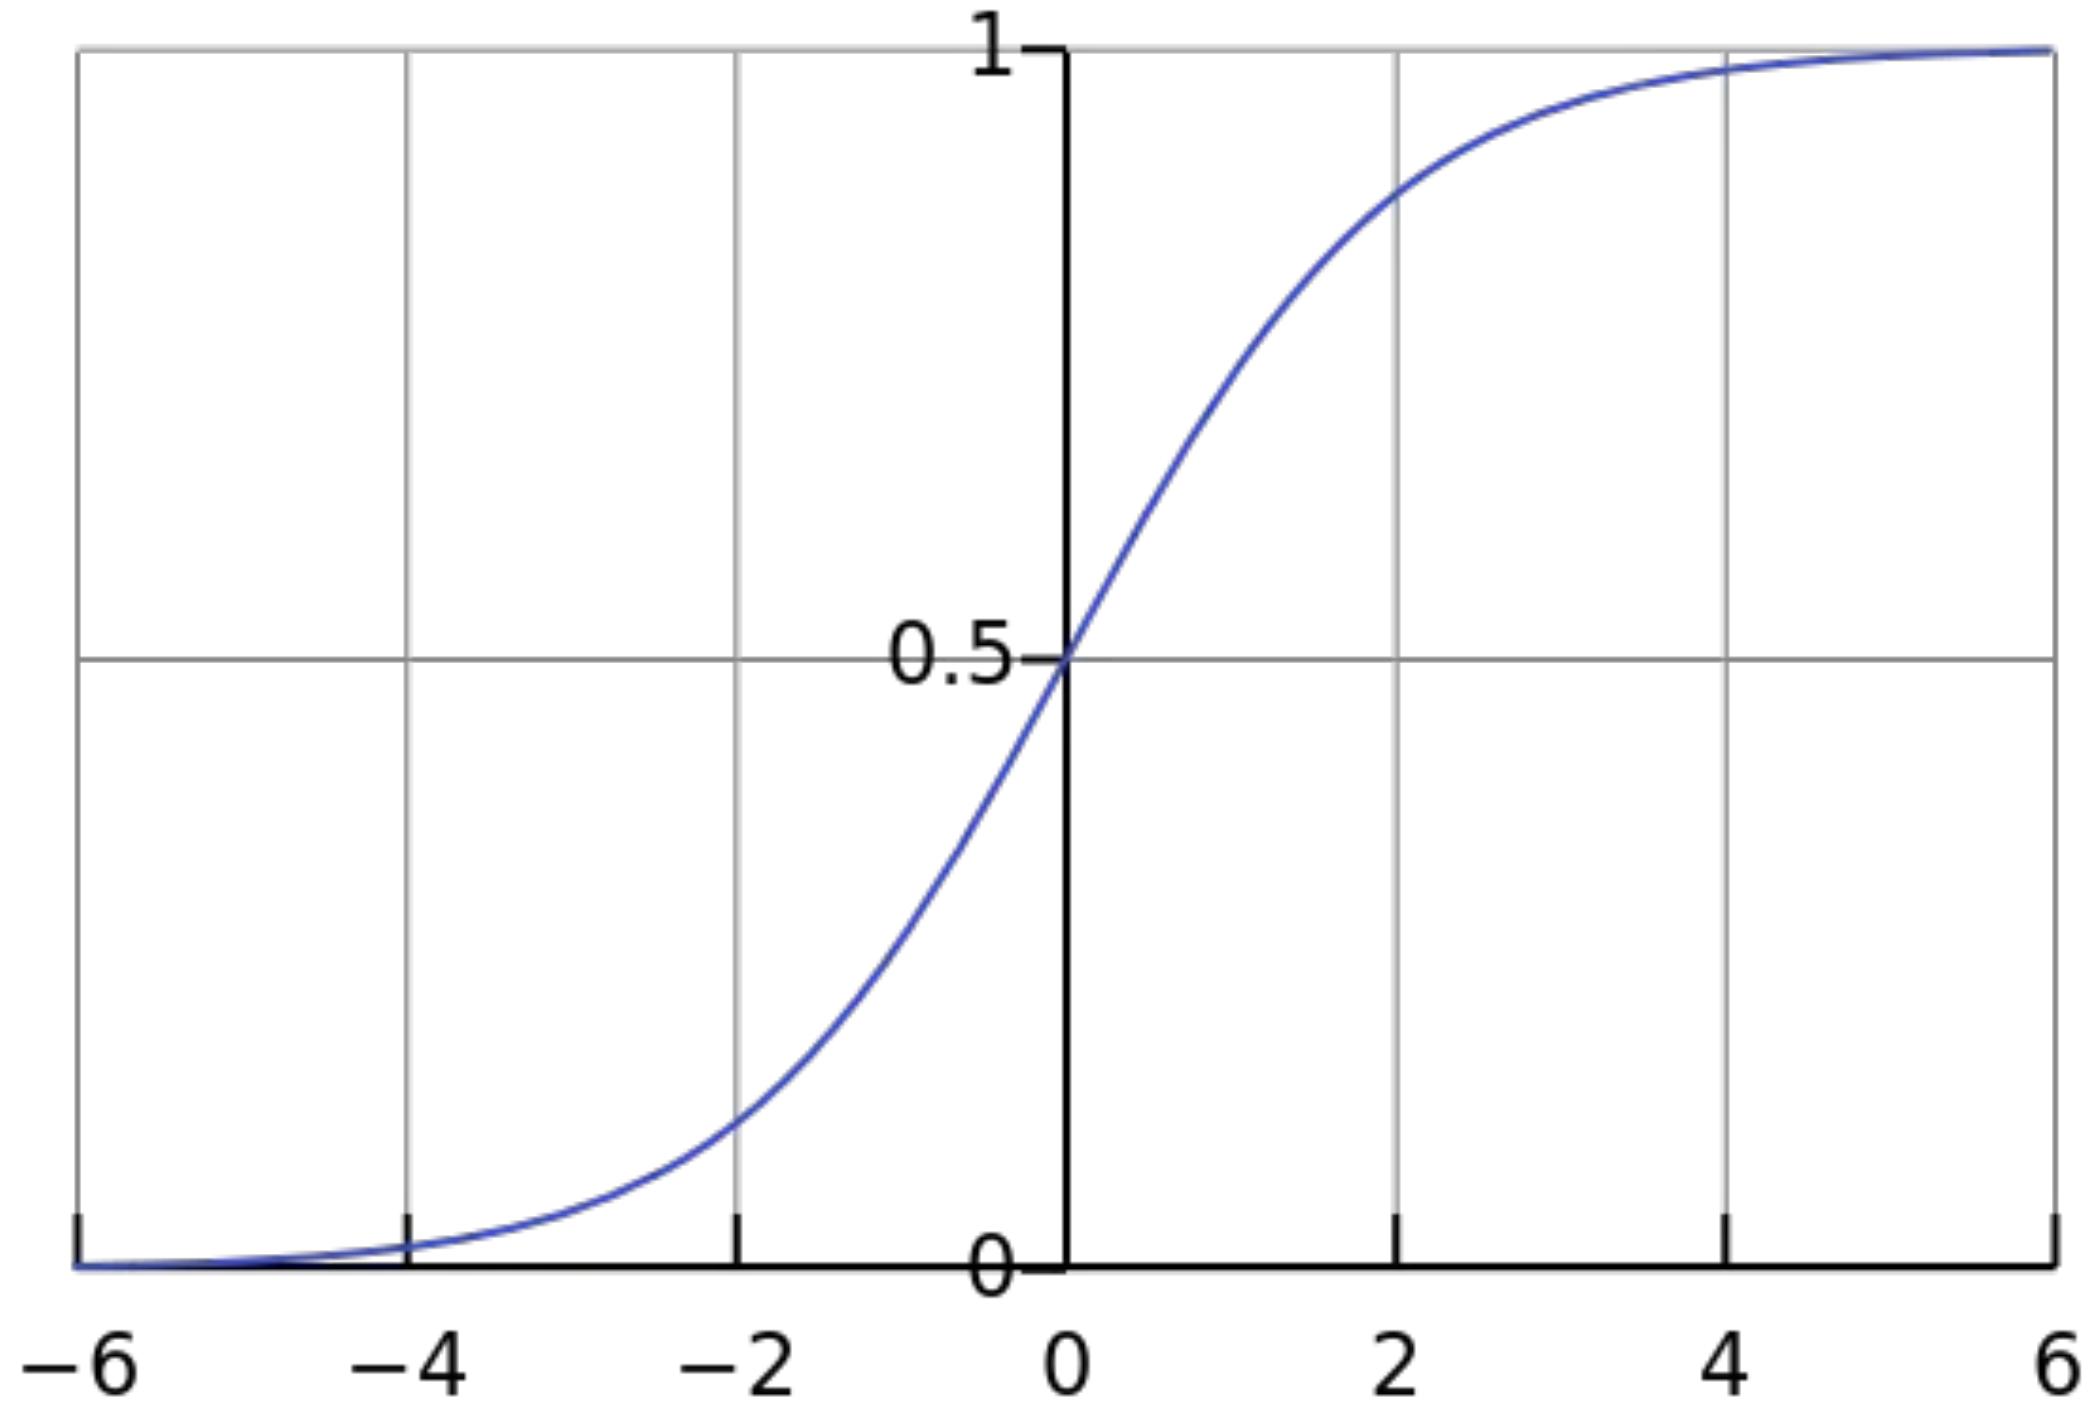


Machine learning mechanics

- Activation Functions

Sigmoid function

$$A = \frac{1}{1+e^{-x}}$$

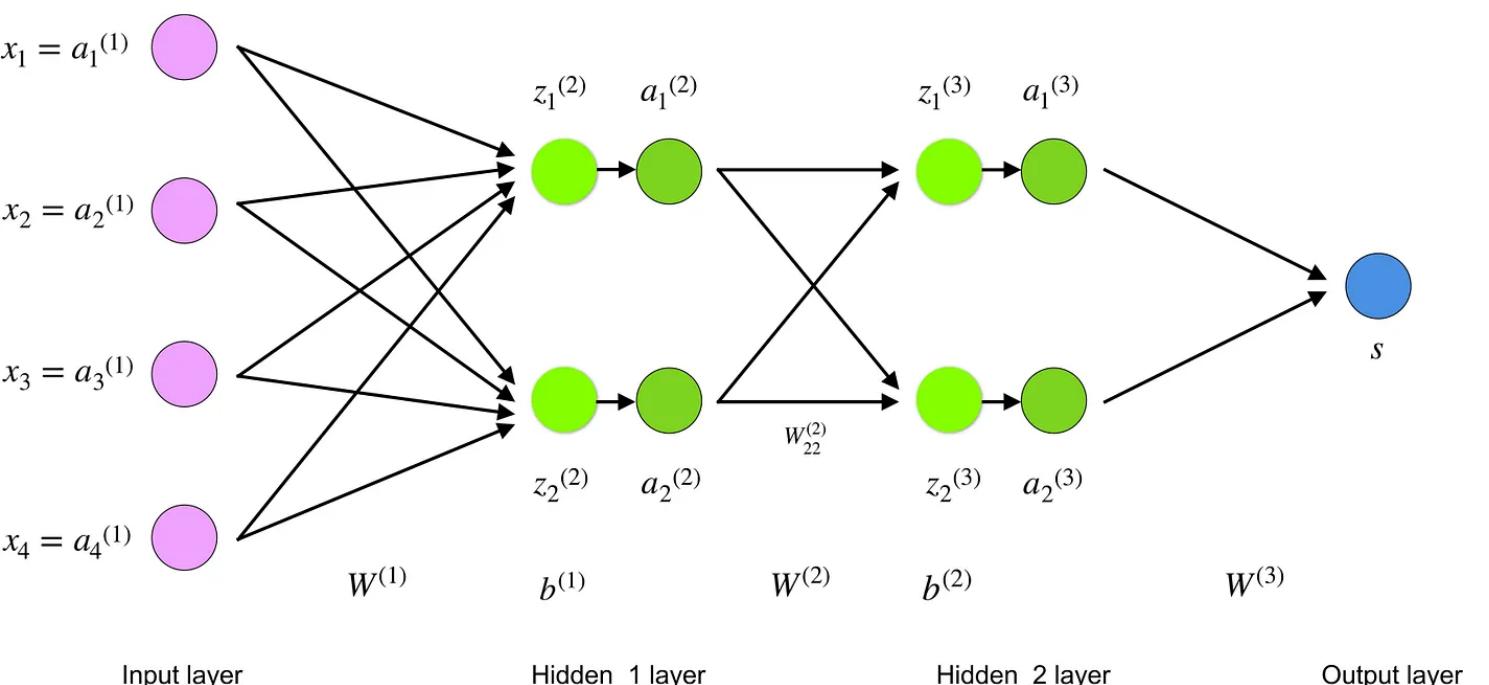
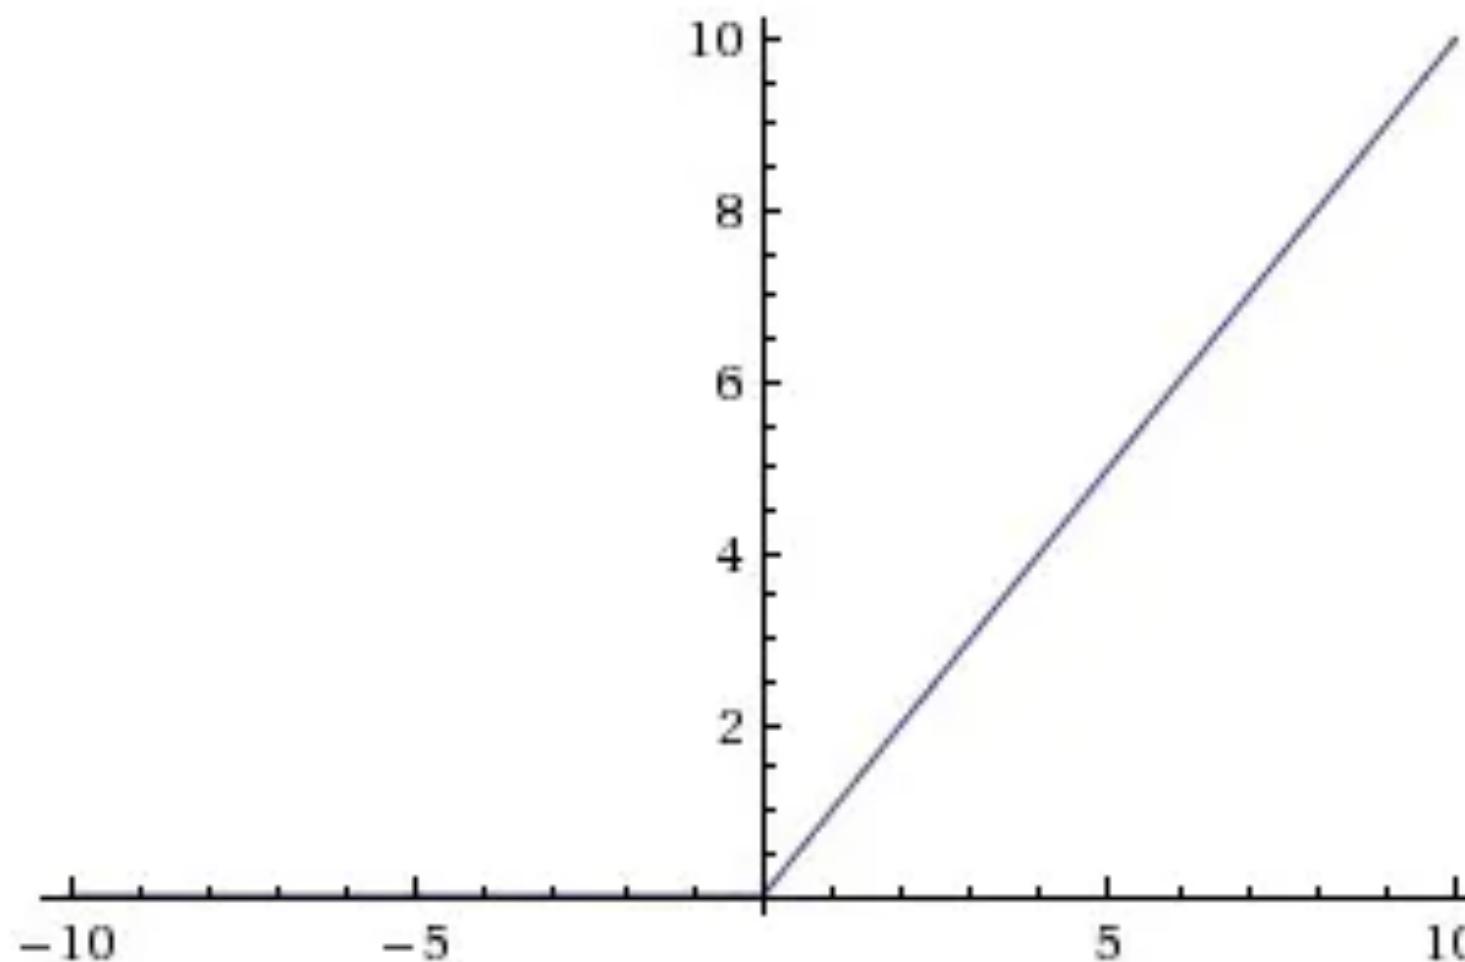


Machine learning mechanics

- Activation Functions

- ReLu function

$$f(x) = \max(0, x)$$

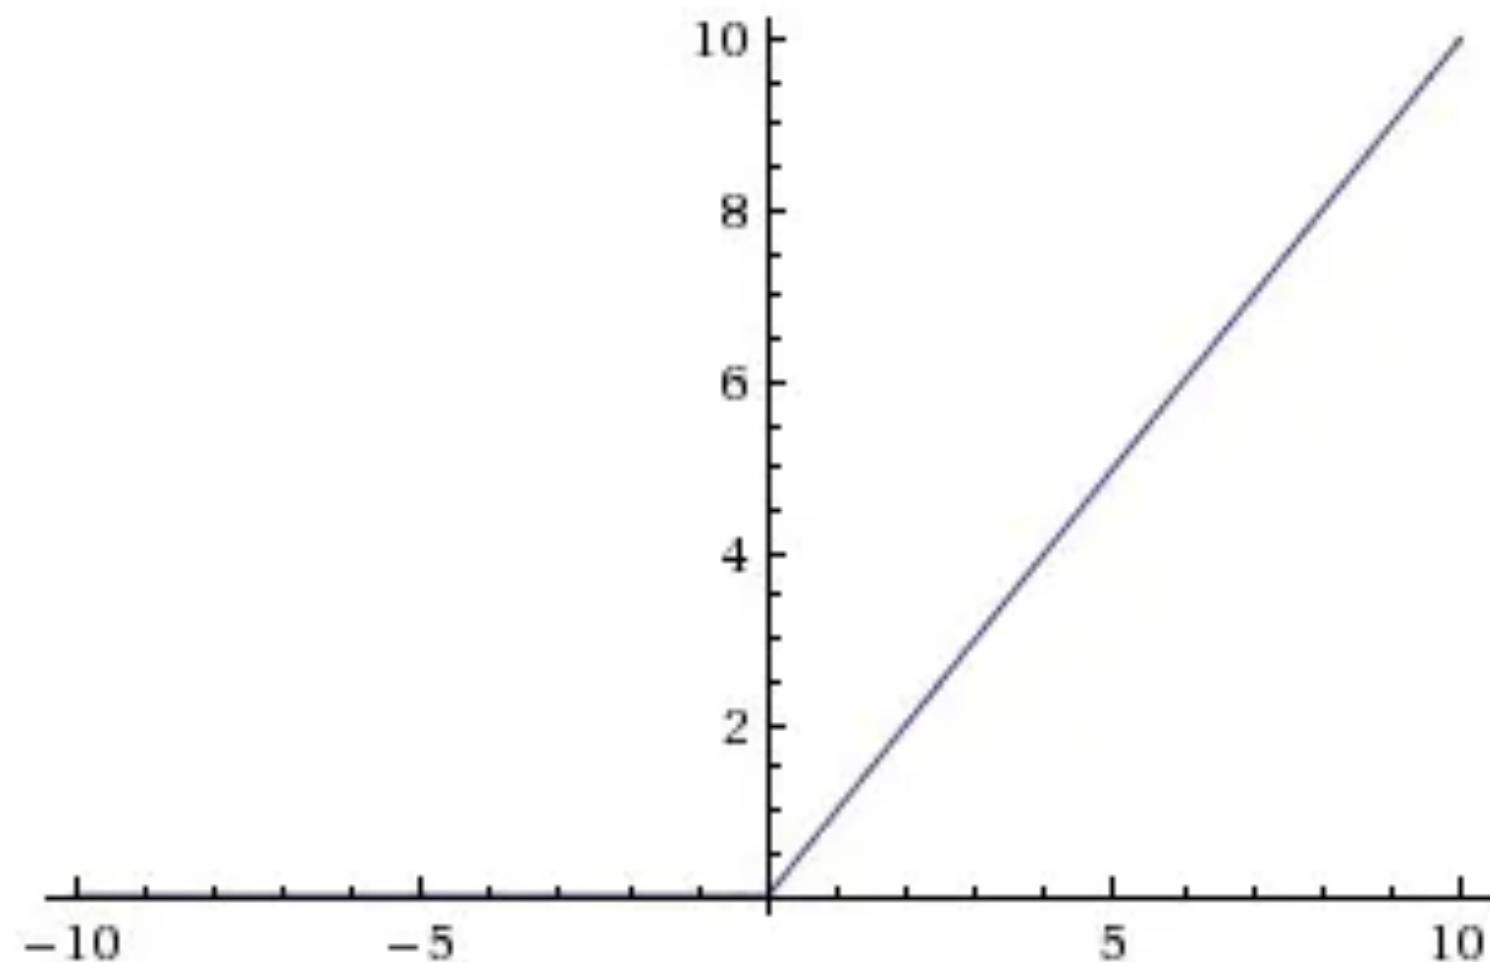


Machine learning mechanics

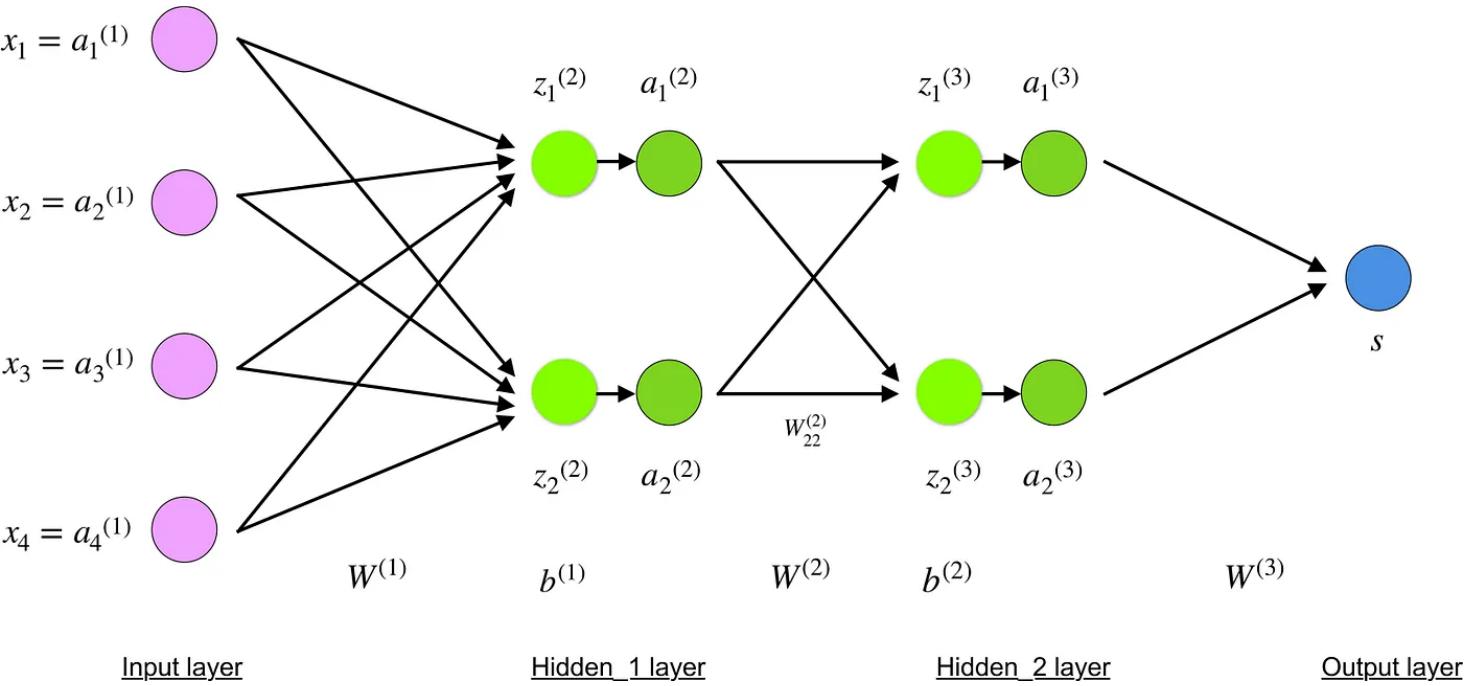
- Activation Functions

Rectified Linear Unit (ReLU)

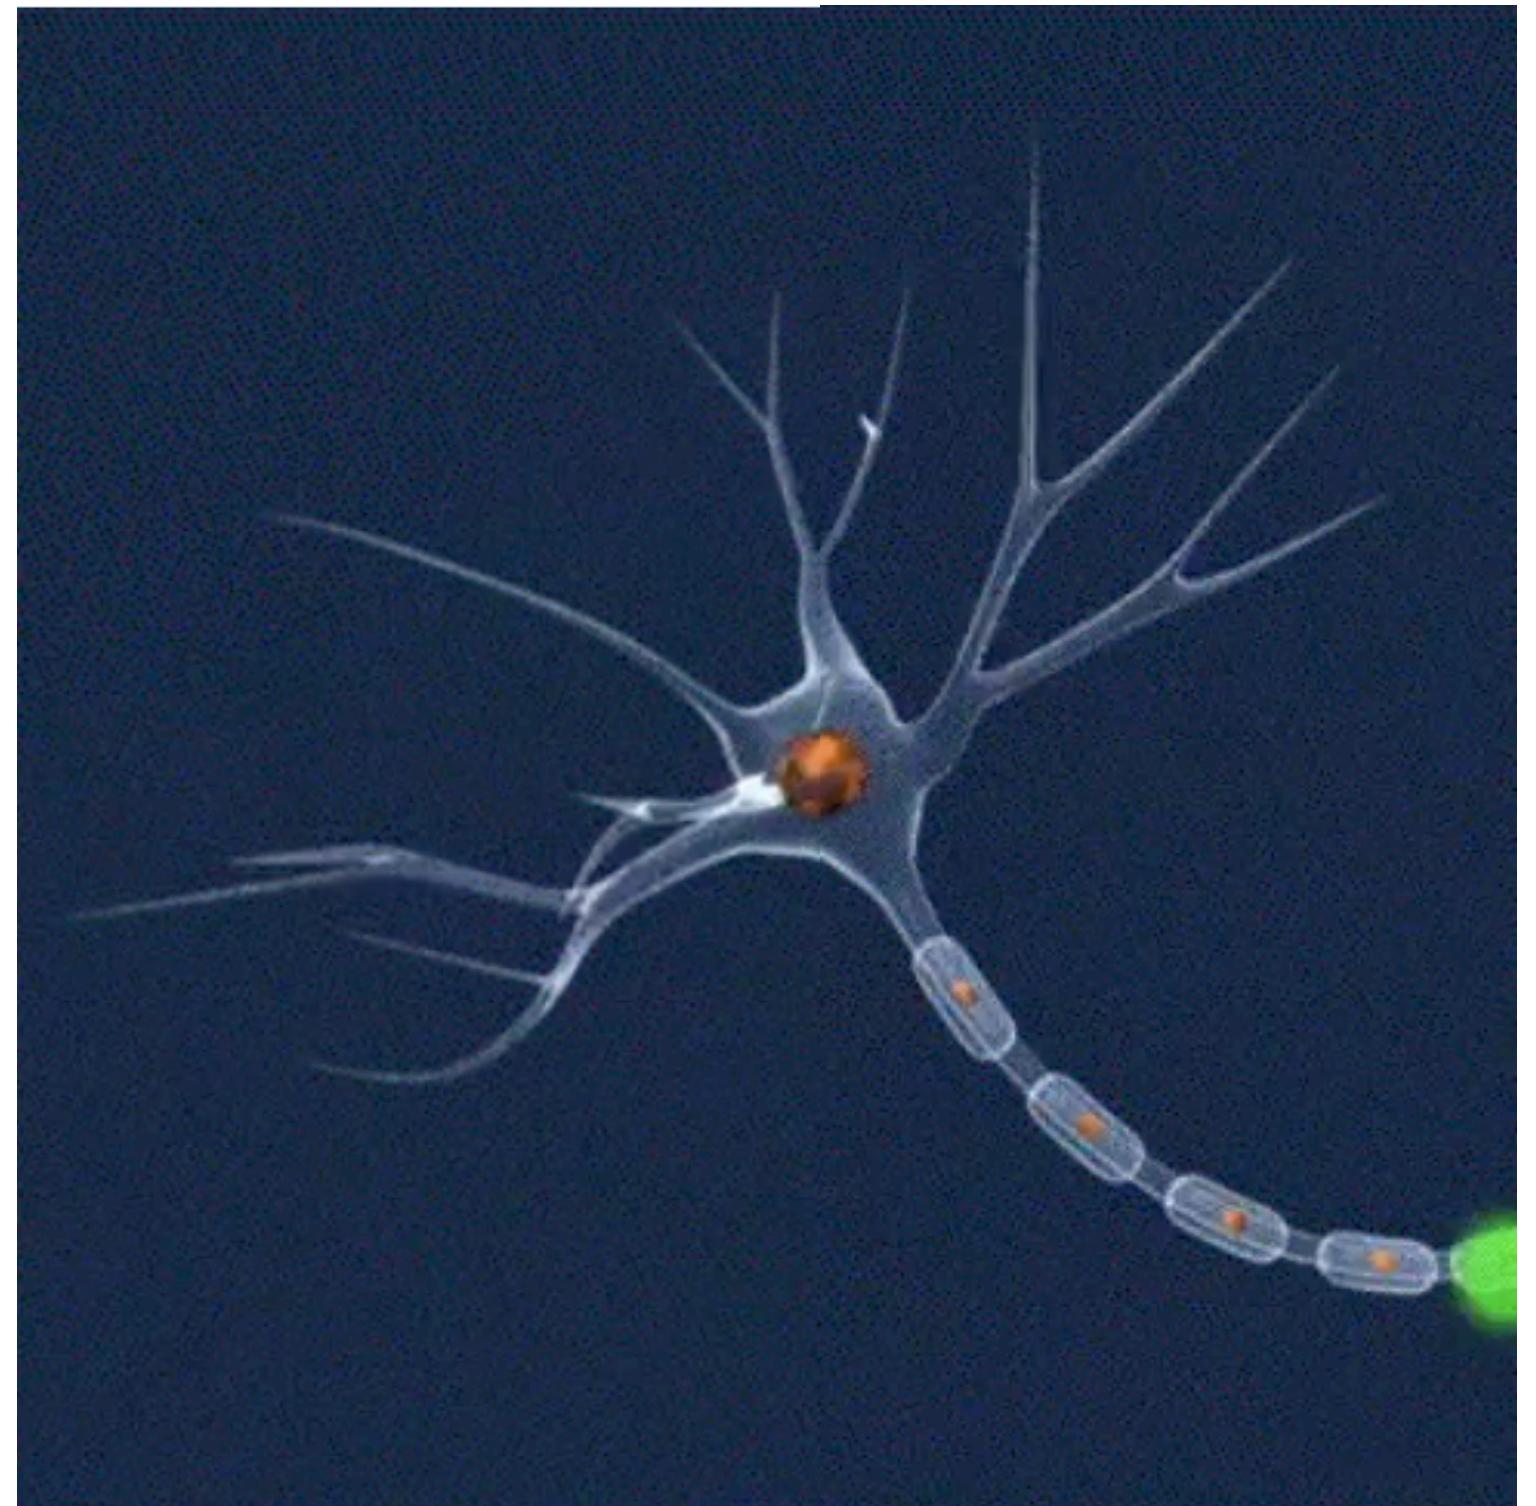
$$f(x) = \max(0, x)$$



$$Y \in [0; \infty]$$



Sparsity of the activation

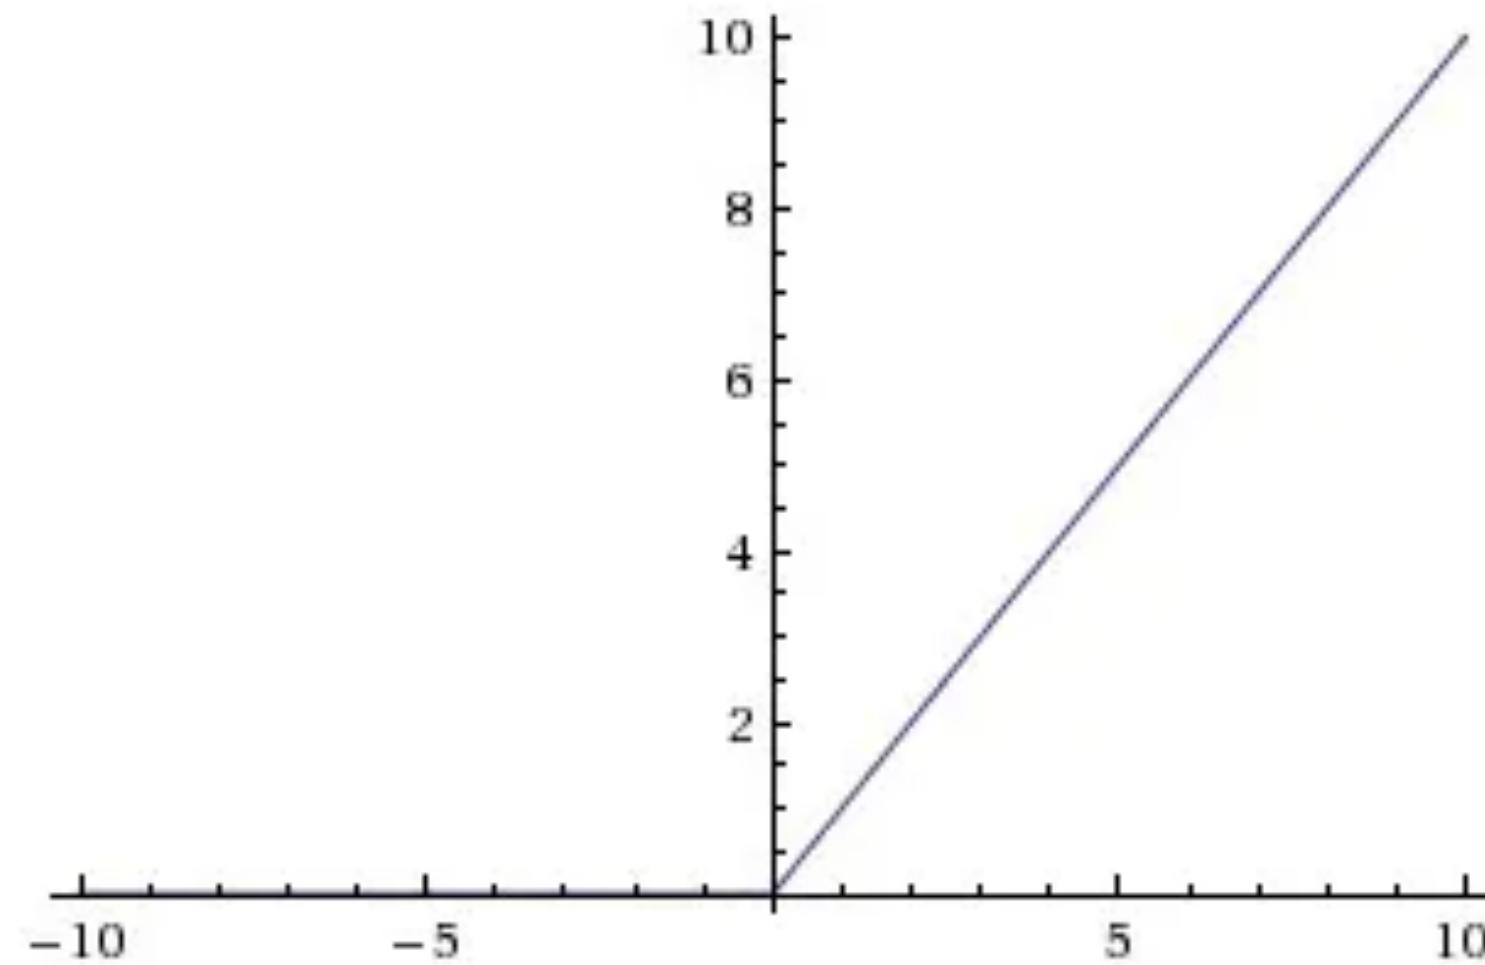


Machine learning mechanics

- Activation Functions

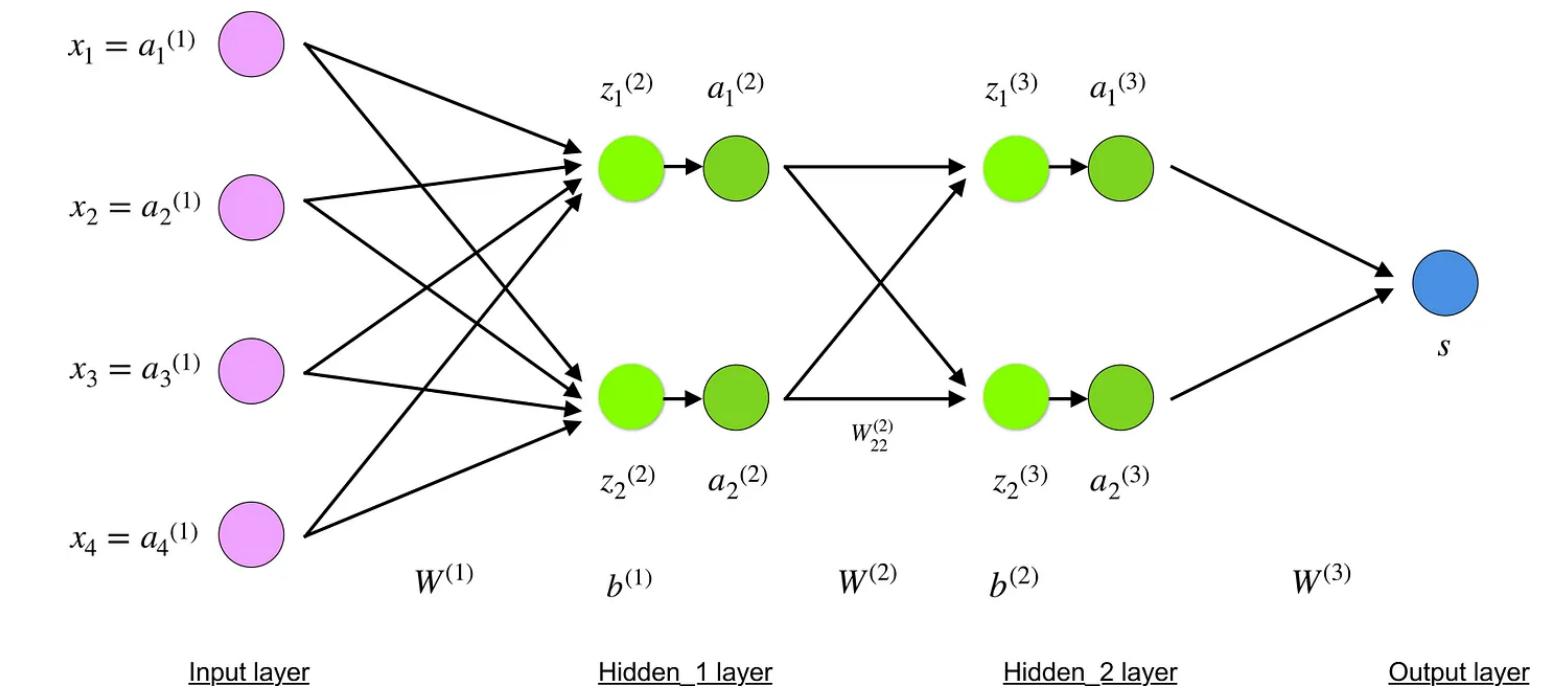
Rectified Linear Unit (ReLU)

$$f(x) = \max(0, x)$$



$$Y \in [0; \infty]$$

- Pros:
 - Less computationally intensive
 - Not Binary
 - Constant gradient
- Cons:
 - Dying ReLU
 - Can Blow up

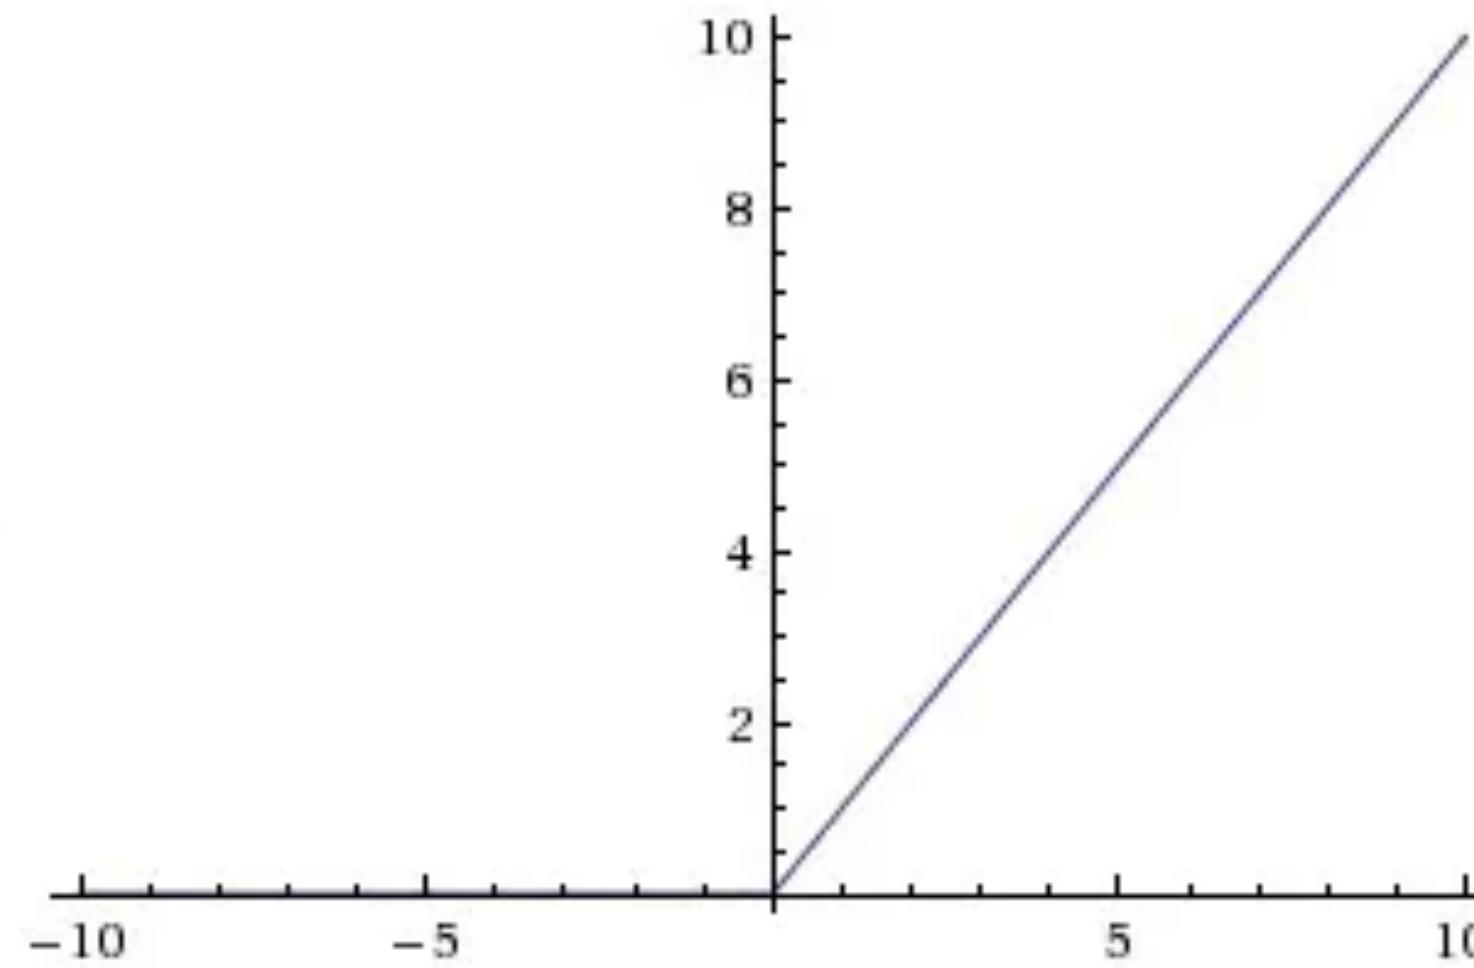


Machine learning mechanics

Activation Functions

Rectified Linear Unit (ReLU)

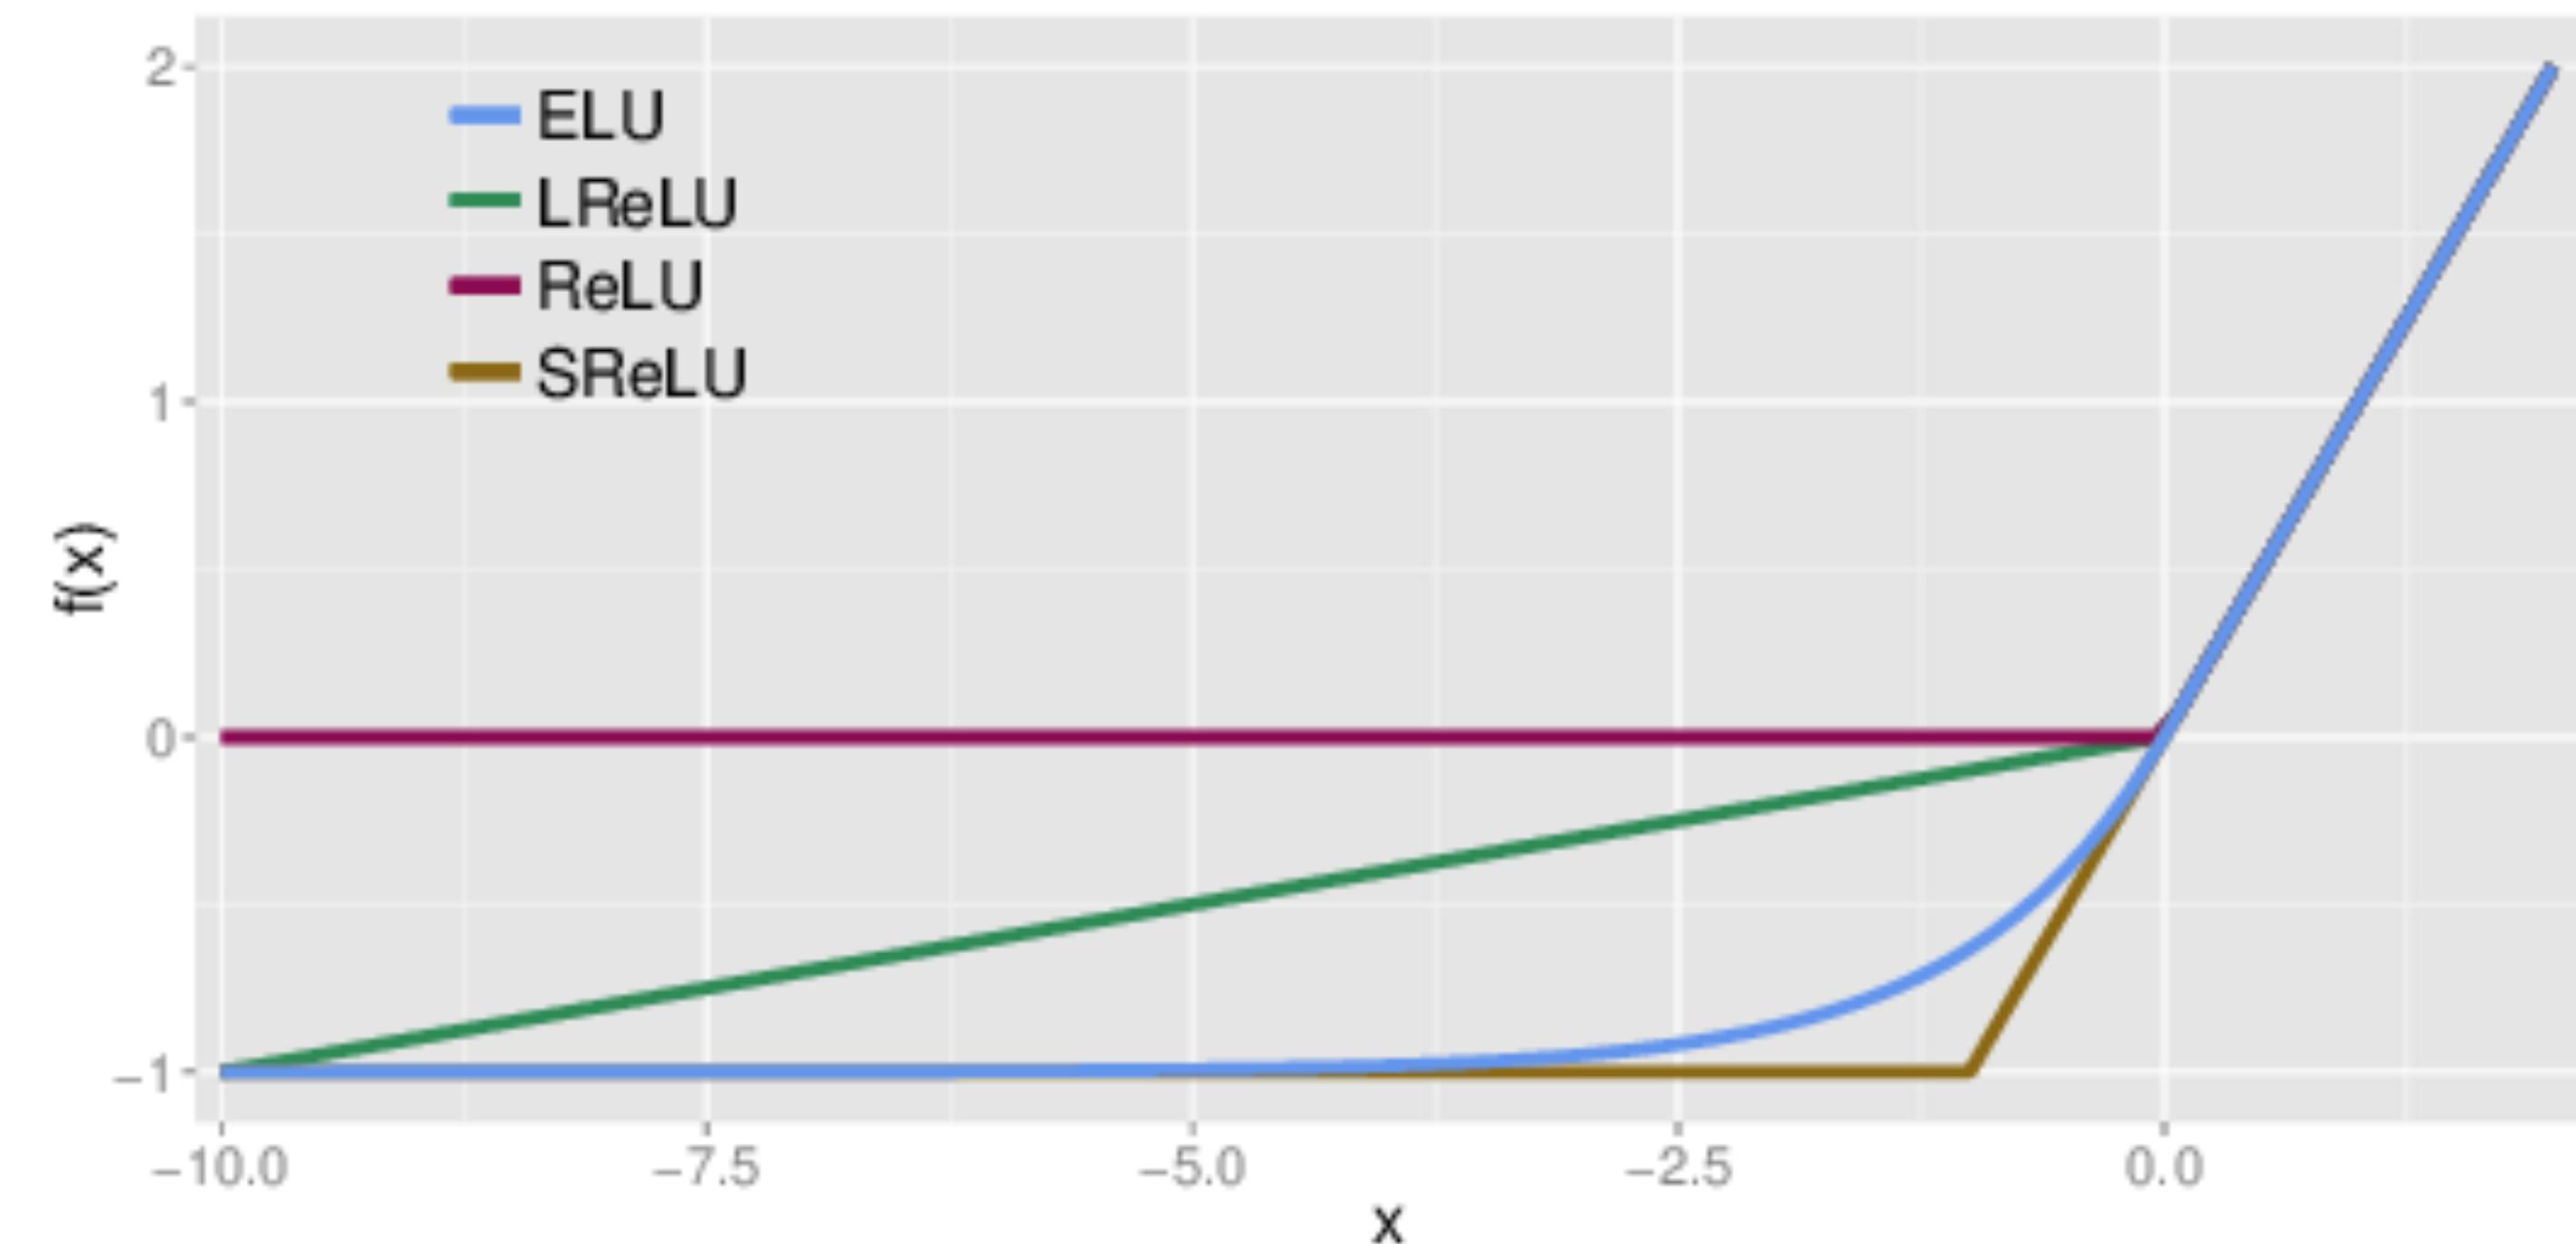
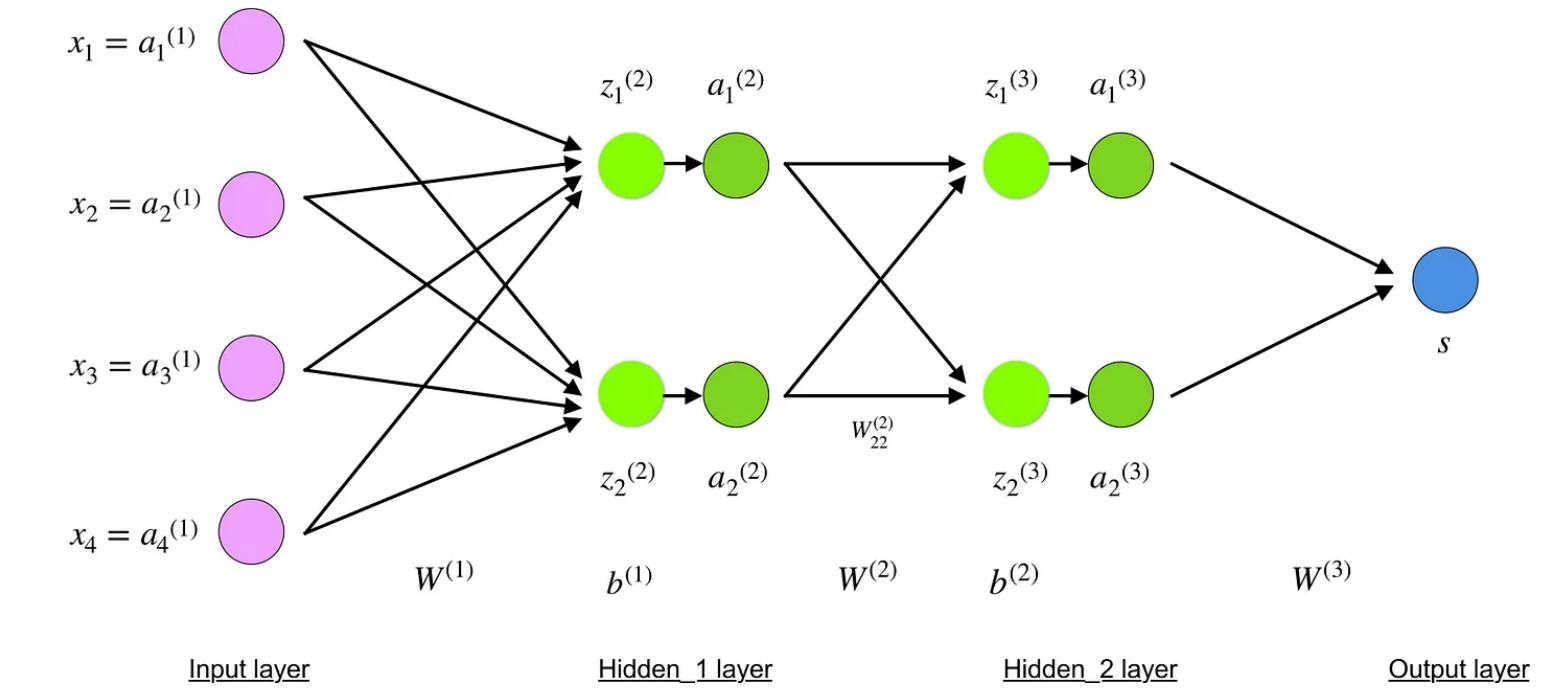
$$f(x) = \max(0, x)$$



$$Y \in [0; \infty]$$

<https://medium.com/the-theory/understanding-activation-functions-in-neural-networks-9491262884eo>

<https://www.nichd.nih.gov/newsroom/releases/031813-backwards-neurons>



Which one to choose?

- Depends on the size of the model
- A good and balanced combination of both 😊

Machine learning mechanics

Loss Functions

- Metric to characterize how far the model is from the label
 - Classification Losses:
 - Discrete (Ex: Character recognition, Animals, (n_2, I_{sat}))
 - Regression Losses:
 - Continuous (Ex: Someone's age, (n_2, I_{sat}))

Machine learning mechanics

- Loss Functions

- Classification Losses

Mean Square Error: $MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$ (better gradients, penalizes off values)

Mean Absolute Error: $MAE = \frac{1}{n} \sum_{i=1}^n | Y_i - \hat{Y}_i |$ (less penalizes off values)

Machine learning mechanics

- Loss Functions

- Regression Losses

Mean Square Error: $MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$ (better gradients, penalizes off values)

Mean Absolute Error: $MAE = \frac{1}{n} \sum_{i=1}^n |Y_i - \hat{Y}_i|$ (less penalizes off values)

Machine learning mechanics

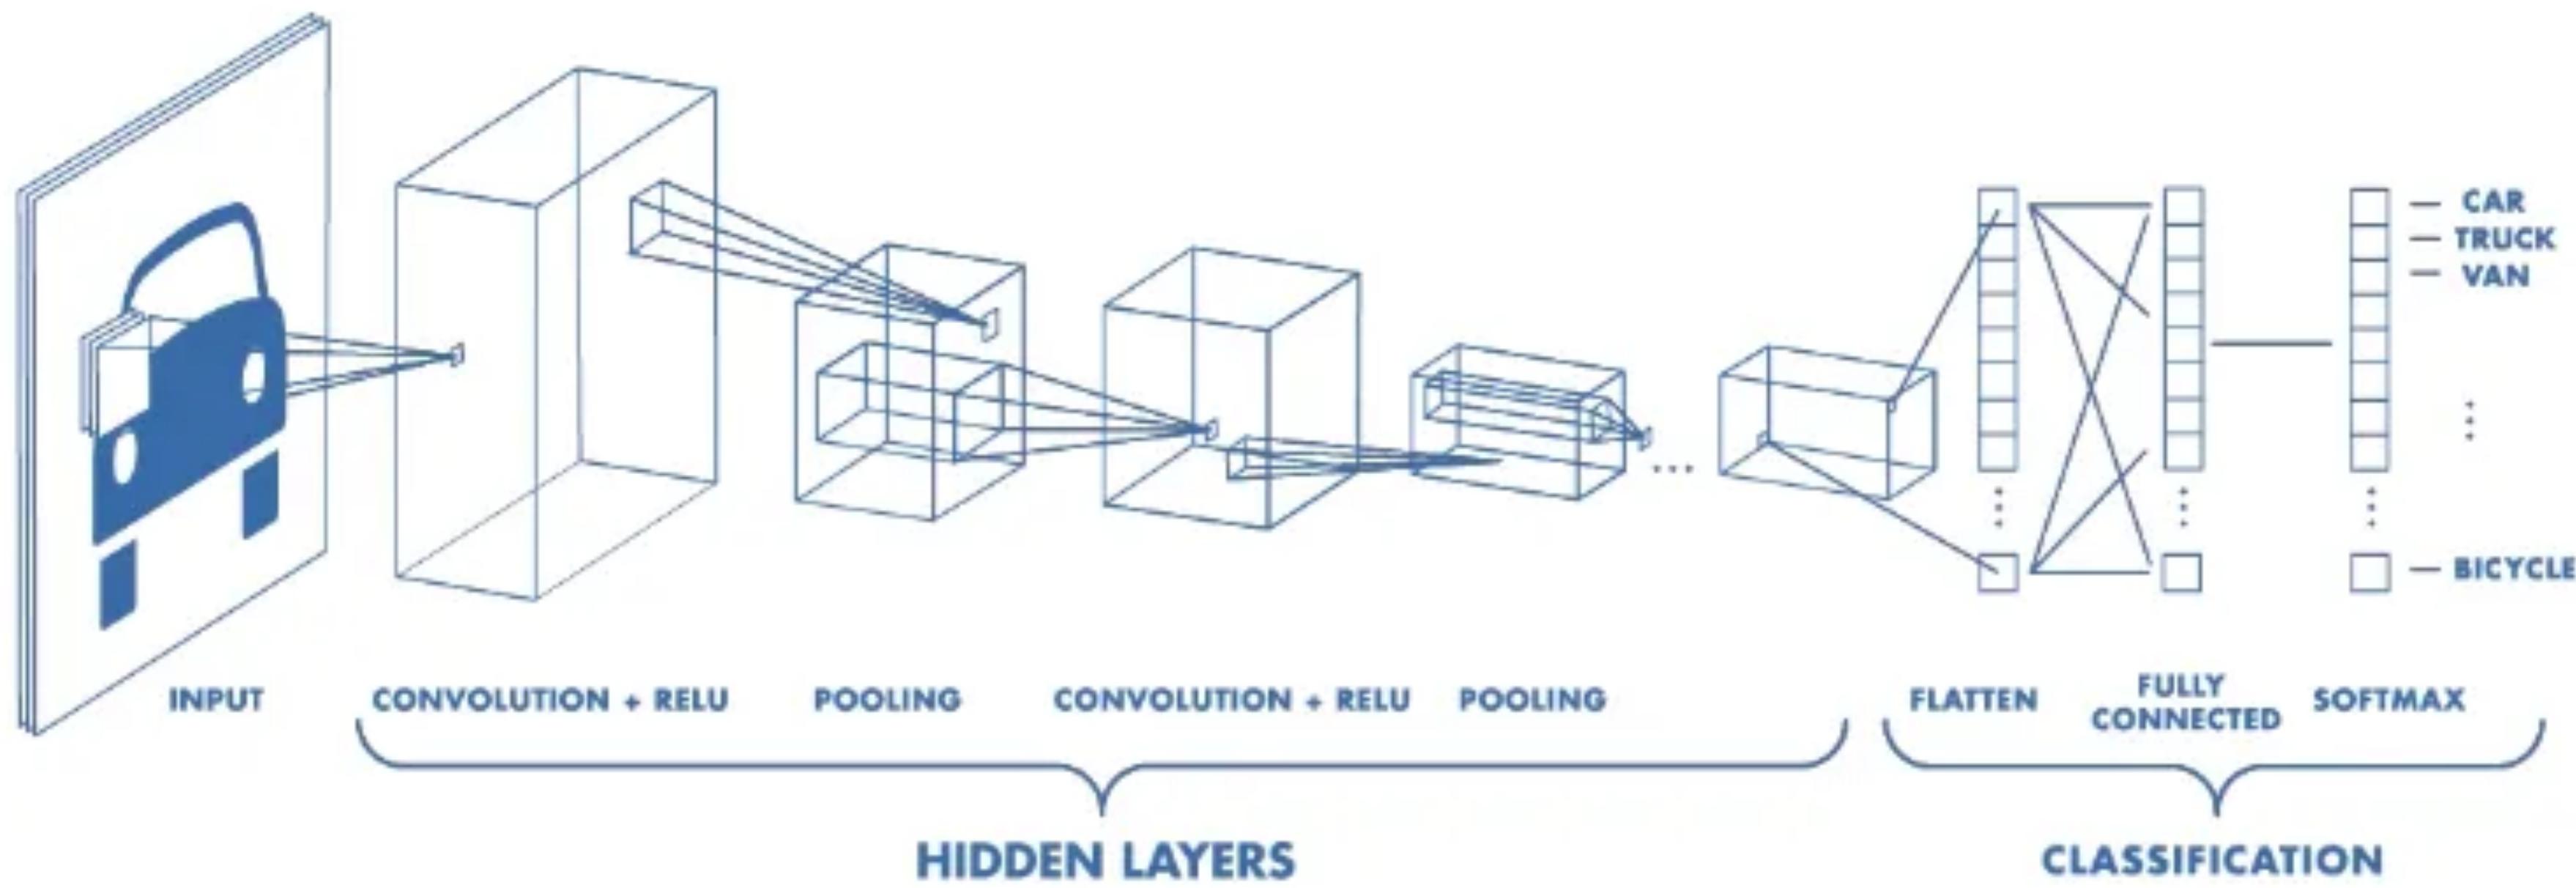
- Loss Functions

- Classification Losses

$$CrossEntropyLoss = -(y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$$

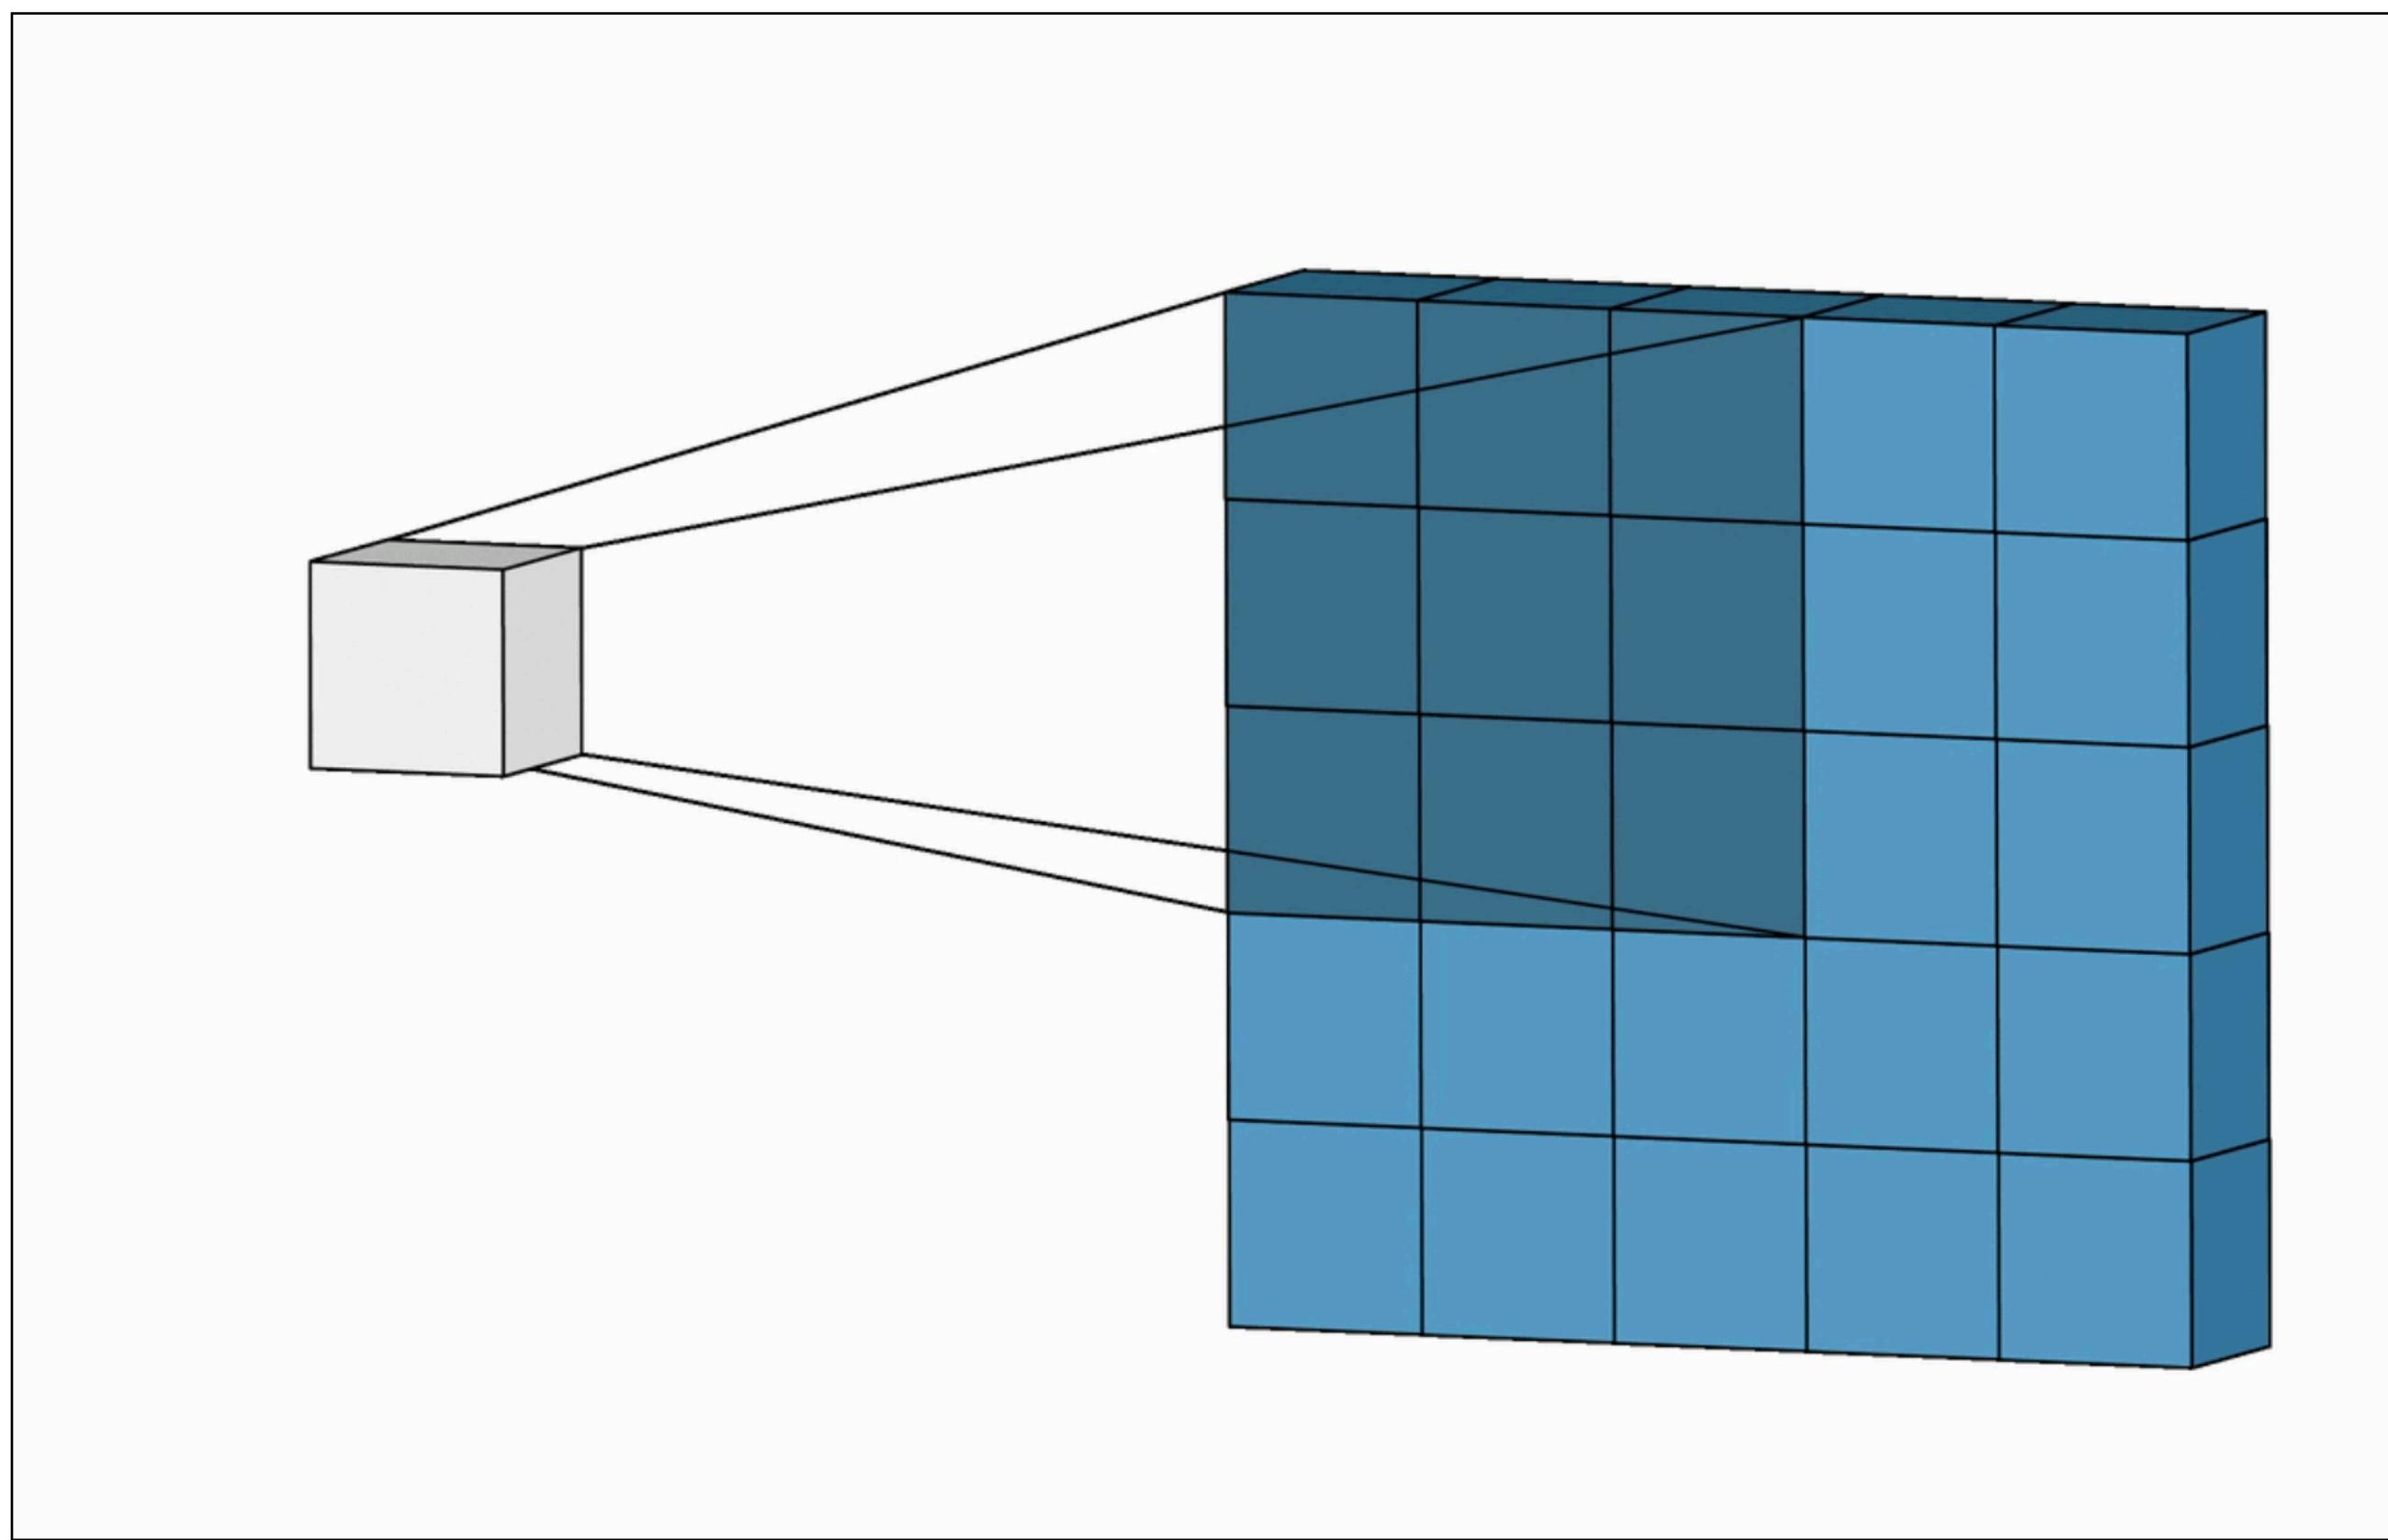
Machine learning mechanics

Convolutional Neural Networks



Machine learning mechanics

- Convolution

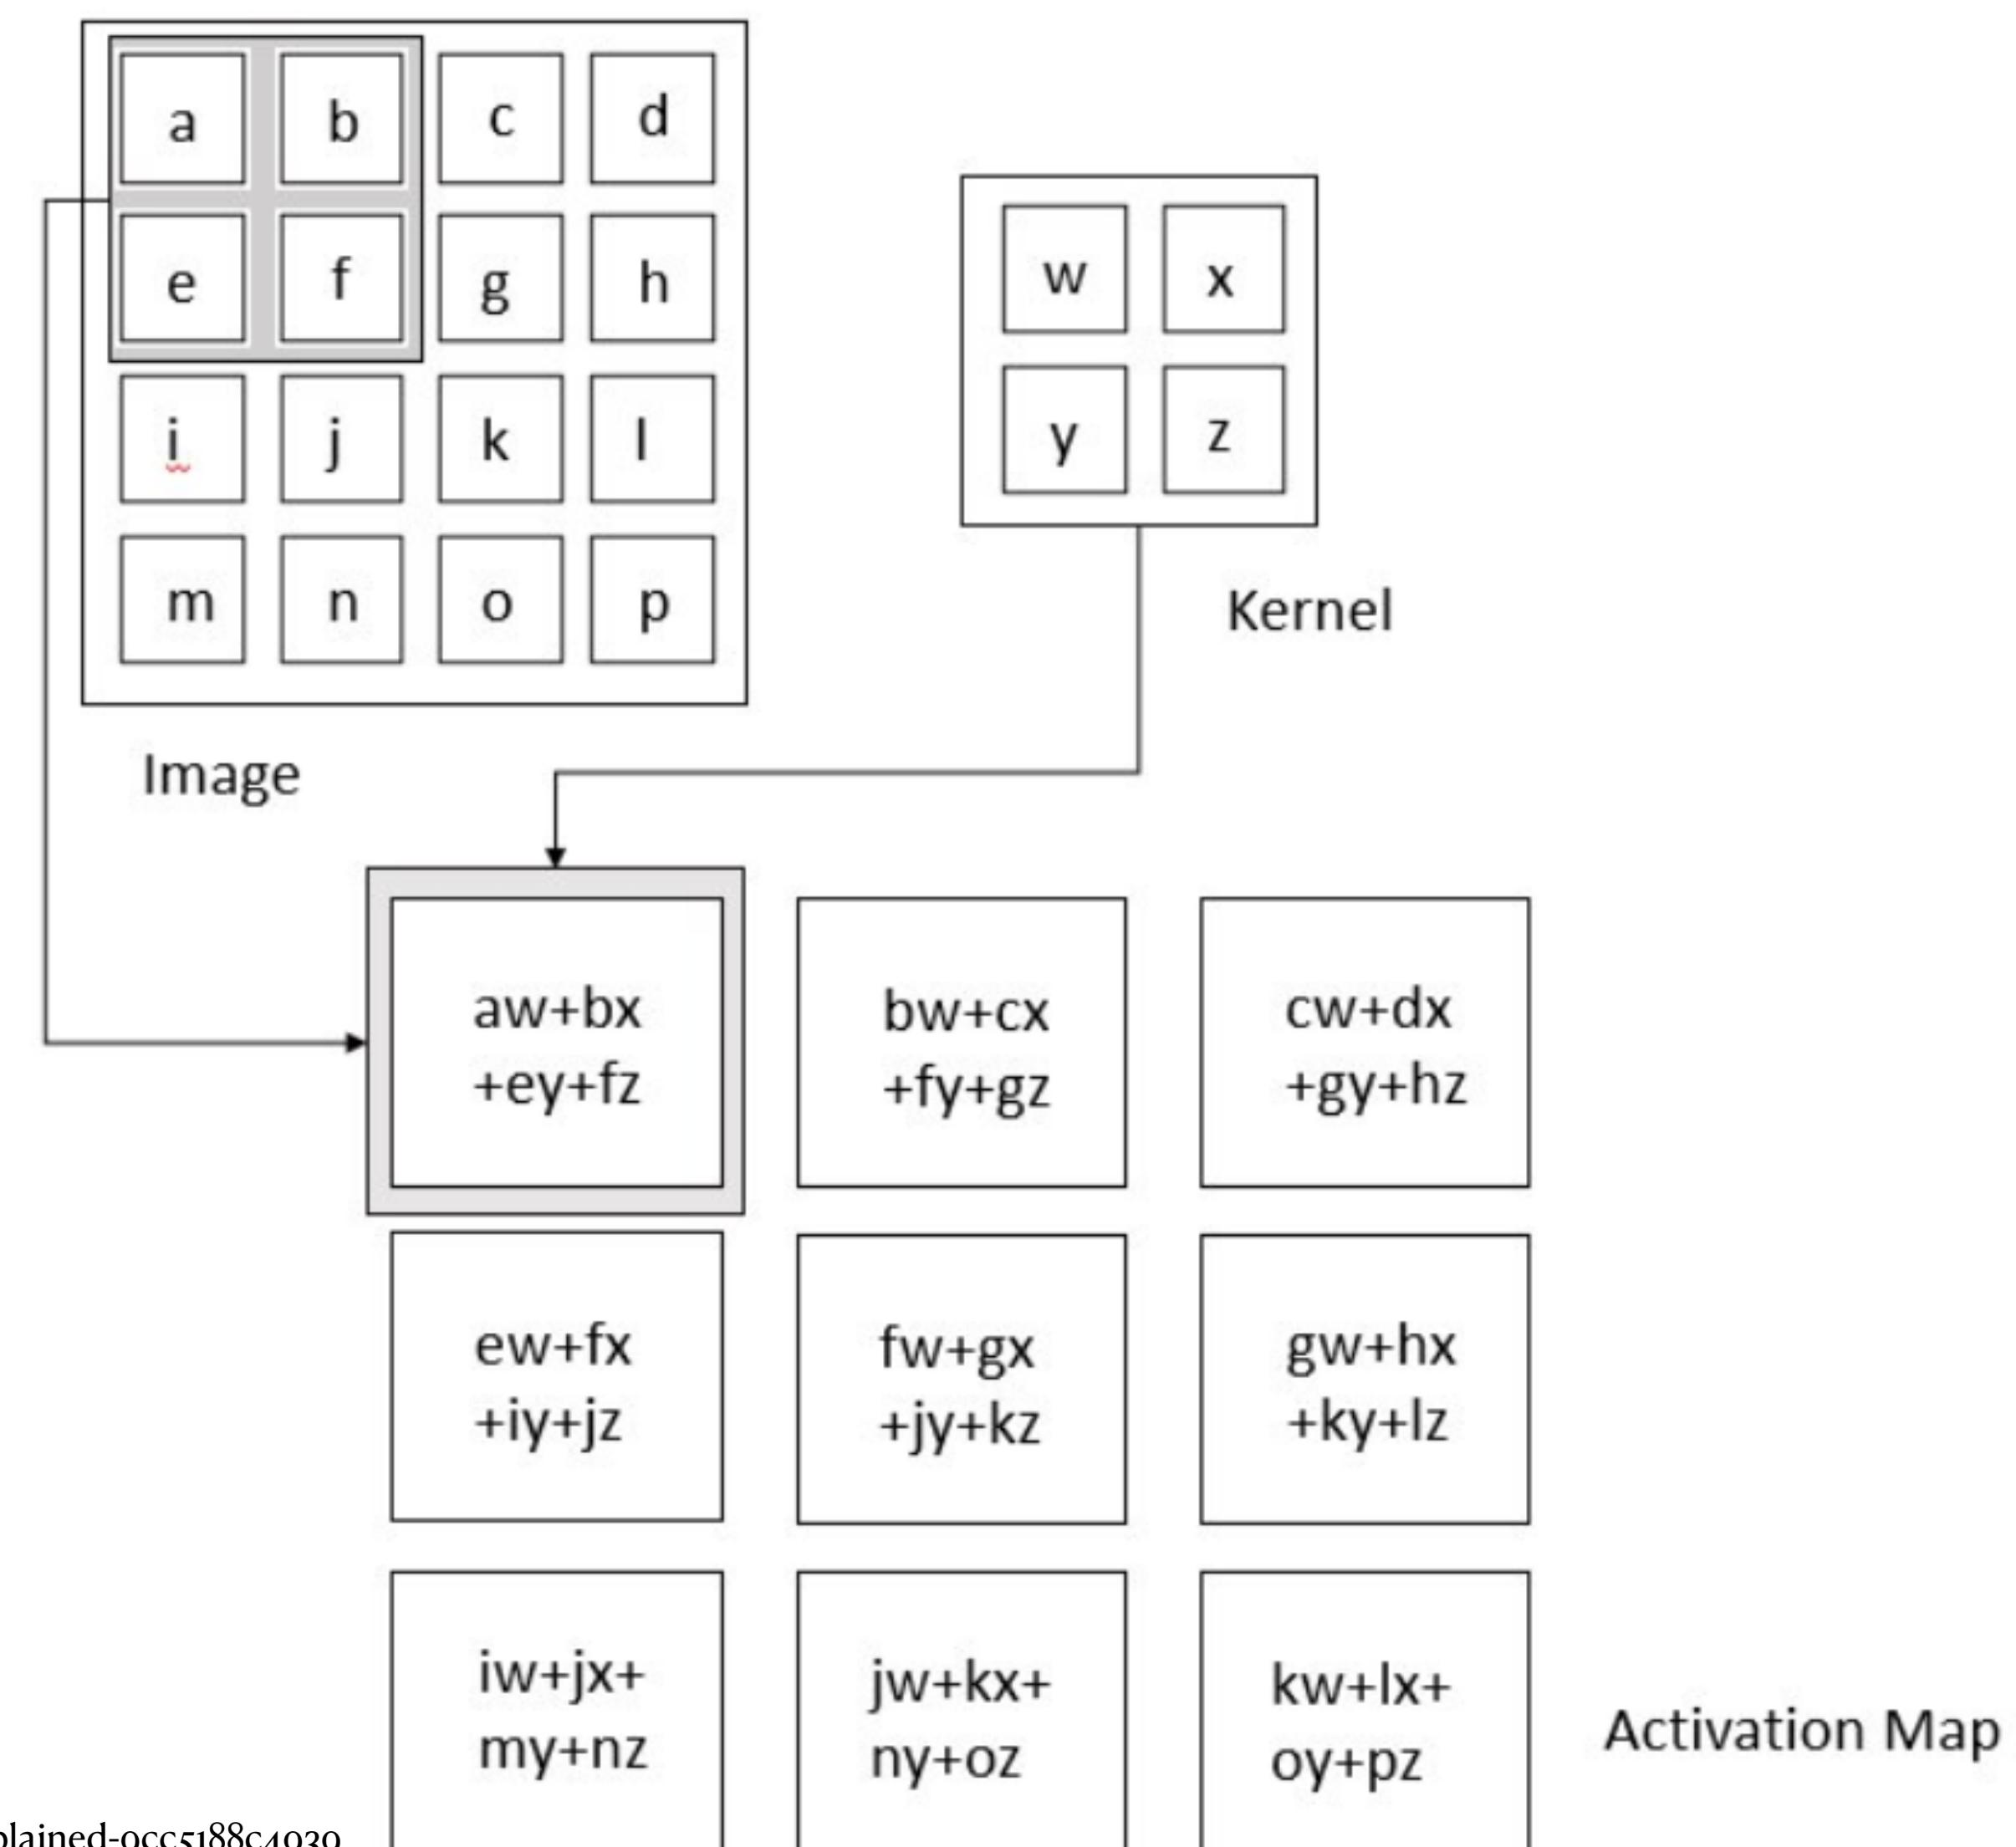


Machine learning mechanics

-

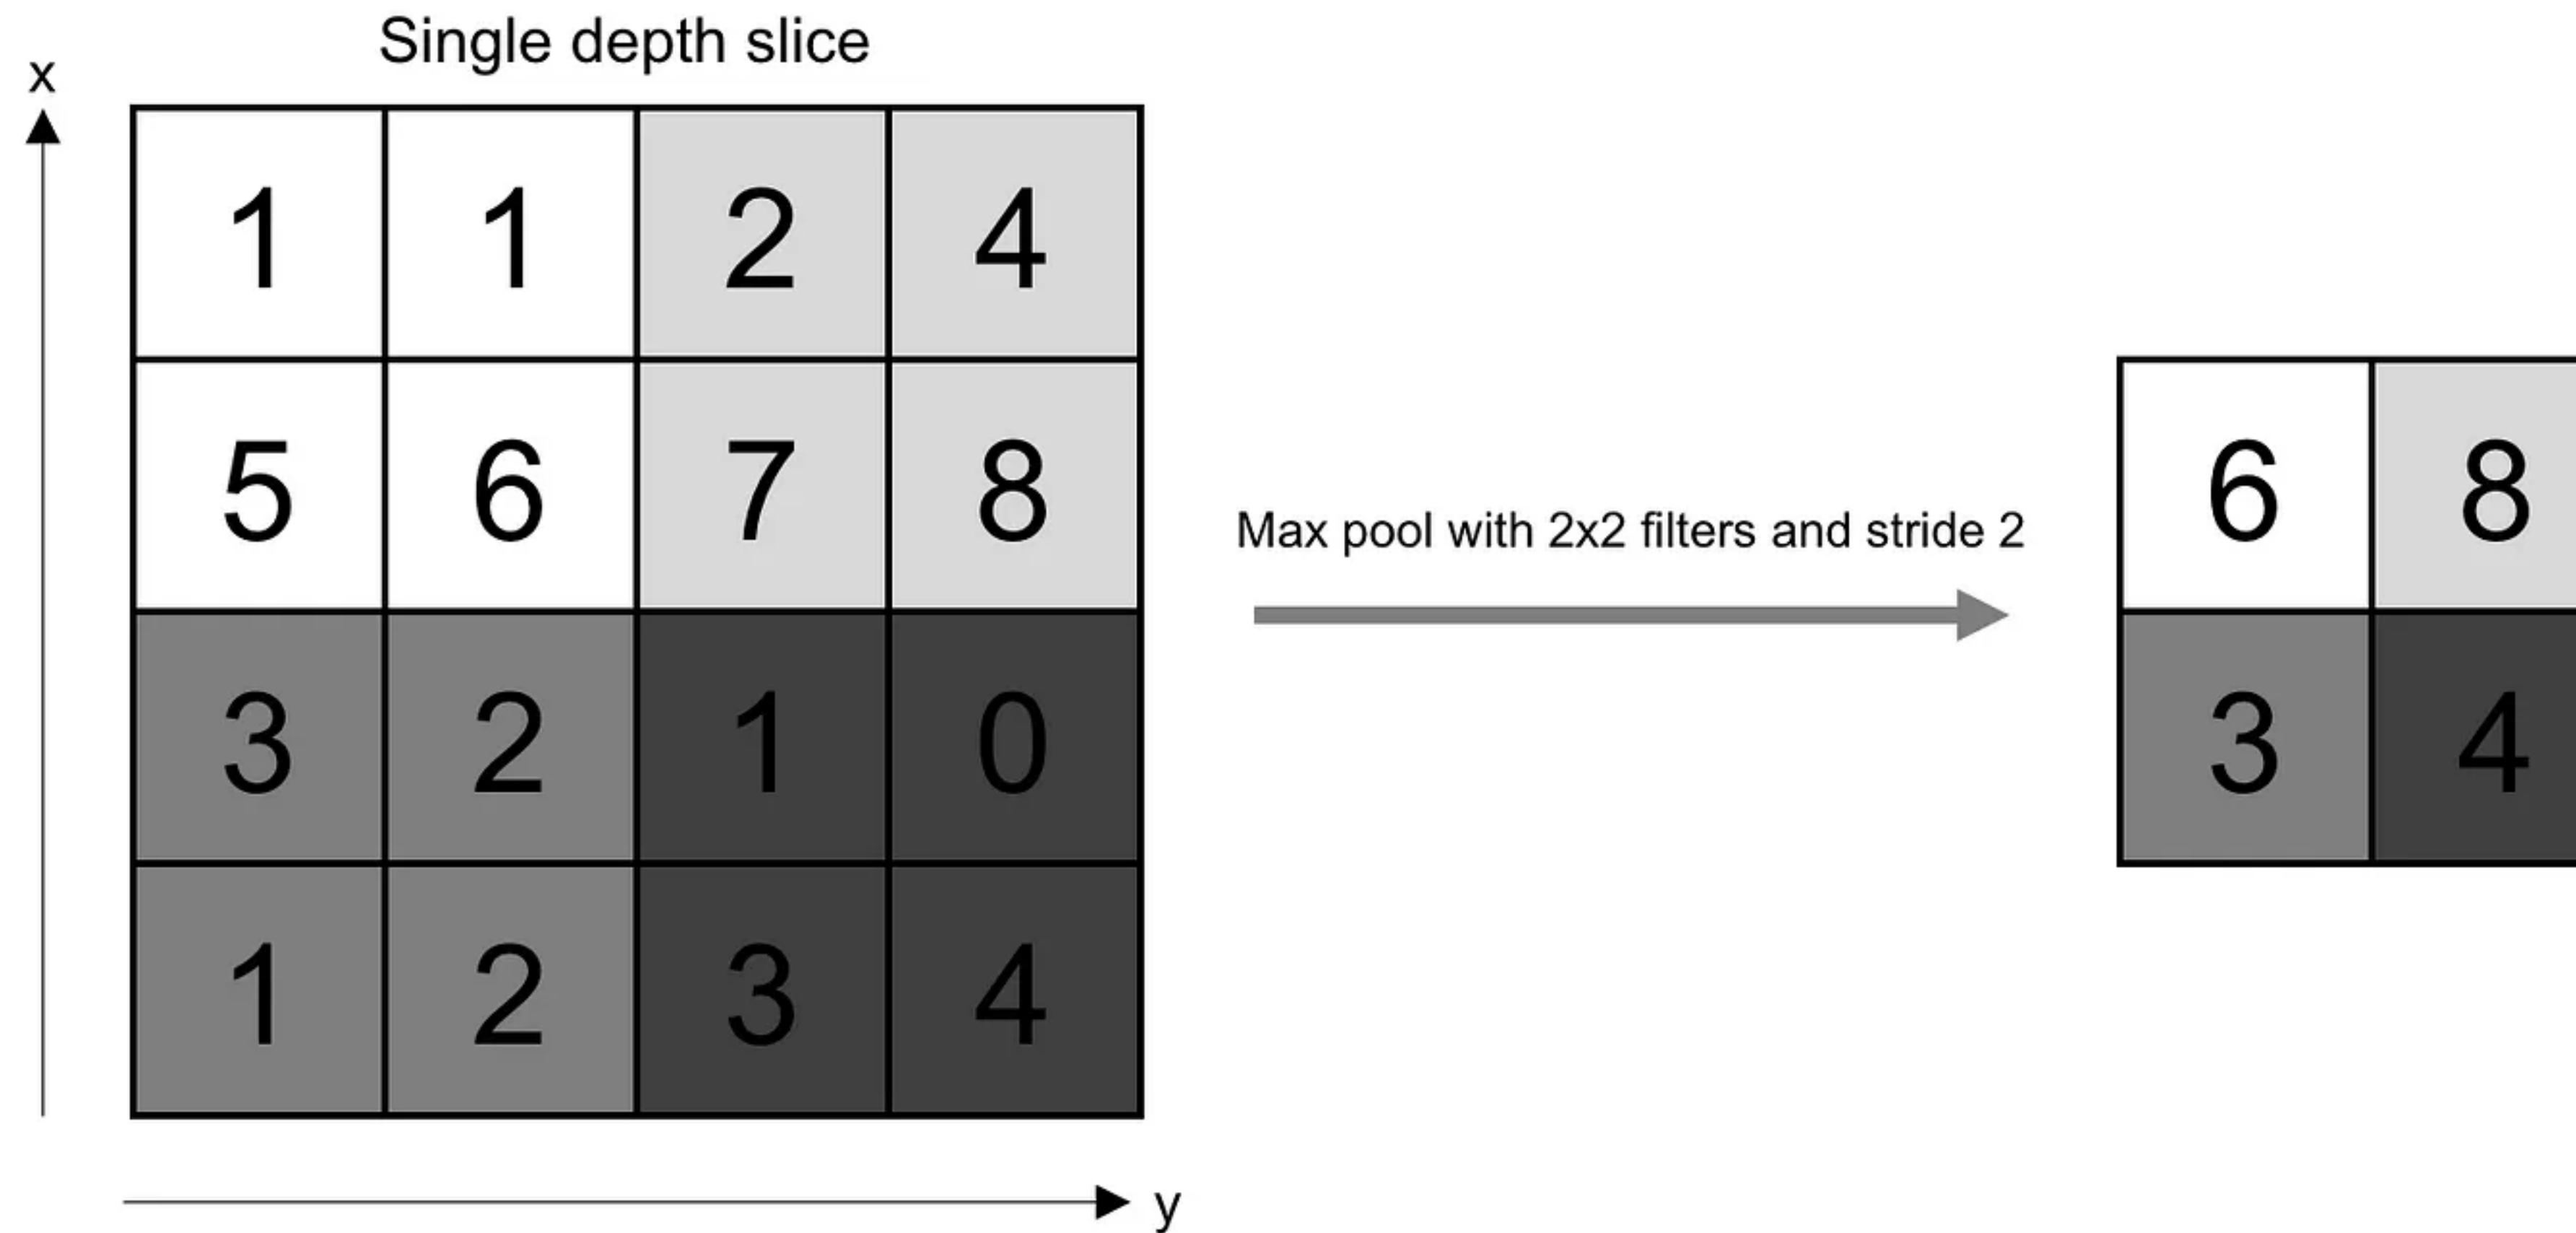
Convolution

$$W_{out} = \frac{W - F + 2P}{S} + 1$$



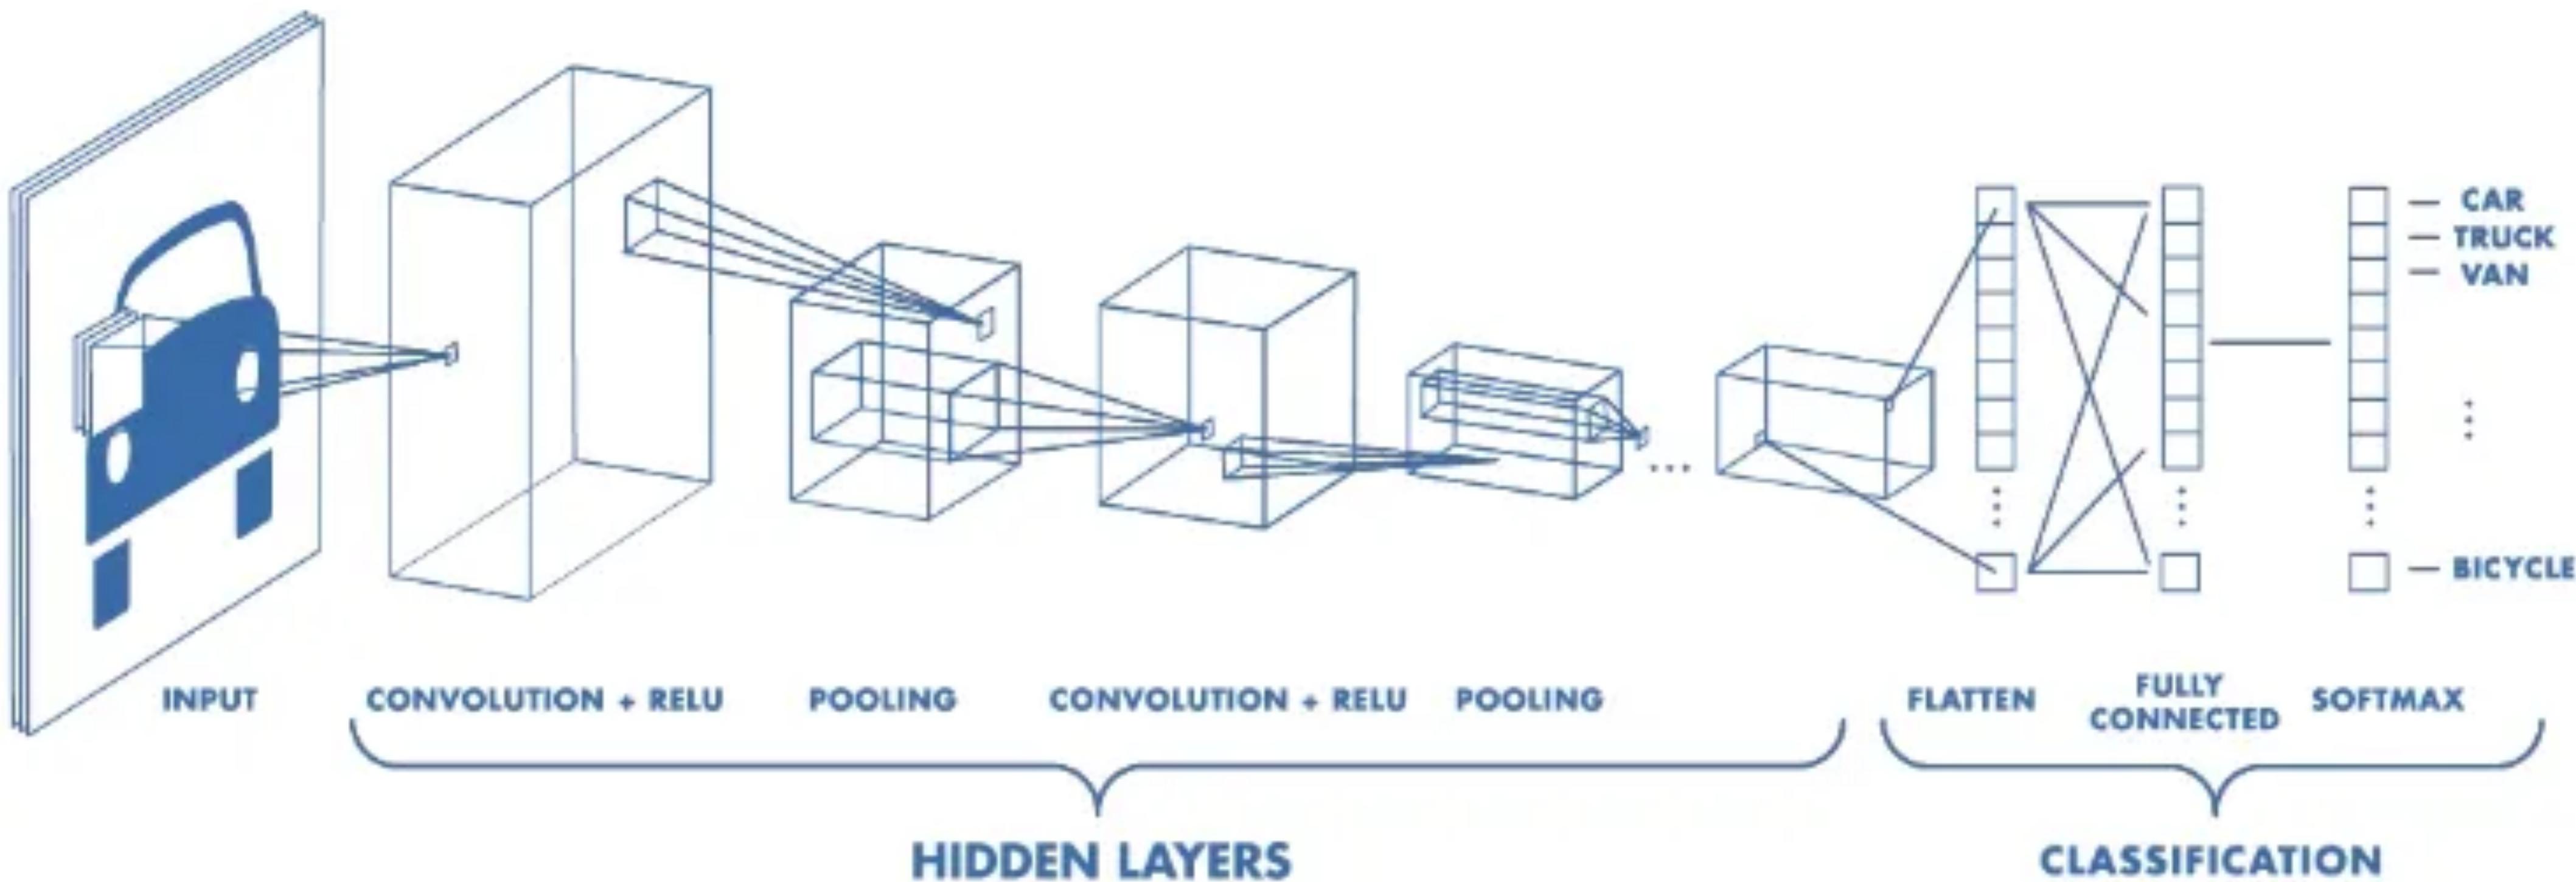
Machine learning mechanics

- Pooling Layer



Machine learning mechanics

Convolutional Neural Networks



Why CNN?

- Deals with signals and images as a whole
- Parameter sharing
- Better generalization

Thank You